


Imitation Learning

Saurabh Gupta

Convert into a Supervised Training Problem


Most General Case




Behavior Cloning

Train Time



Assume an expert e can solve this MDP.




1. Ask the expert e to solve this MDP.
2. Collect labeled dataset D from expert.
3. Train a function $\pi(o_t)$ that mimics $\pi_e(o_t)$ on D .


$$O_t$$


Train with back-propagation




Test Time



Supervision from Human Expert

Train Time




Assume an expert e can solve this MDP.



1. Ask the expert e to solve this MDP.
2. Collect labeled dataset D from expert.
3. Train a function $\pi(o_t)$ that mimics $\pi_e(o_t)$ on D .




o_t



=

Train with back-propagation


Test Time



Supervision from ~~Human~~ Algorithmic Expert

Train Time

1. Instrument the environment such that it becomes a known MDP.



Fully Observed System

Known or Learned Transition Function

Known Reward Function


$$o_t = s_t$$

$$s_{t+1} \sim \hat{T}(s_t, a_t)$$

$$R(s_{t+1}, s_t, a_t)$$

2. Train a function $\hat{\pi}(o_t)$ that mimics $\pi(o_t, T, R)$

Test Time



Fully Observed System

Known or Learned Transition Function

Known Reward Function

$$o_t = s_t$$

$$s_{t+1} \sim \hat{T}(s_t, a_t)$$

$$R(s_{t+1}, s_t, a_t)$$

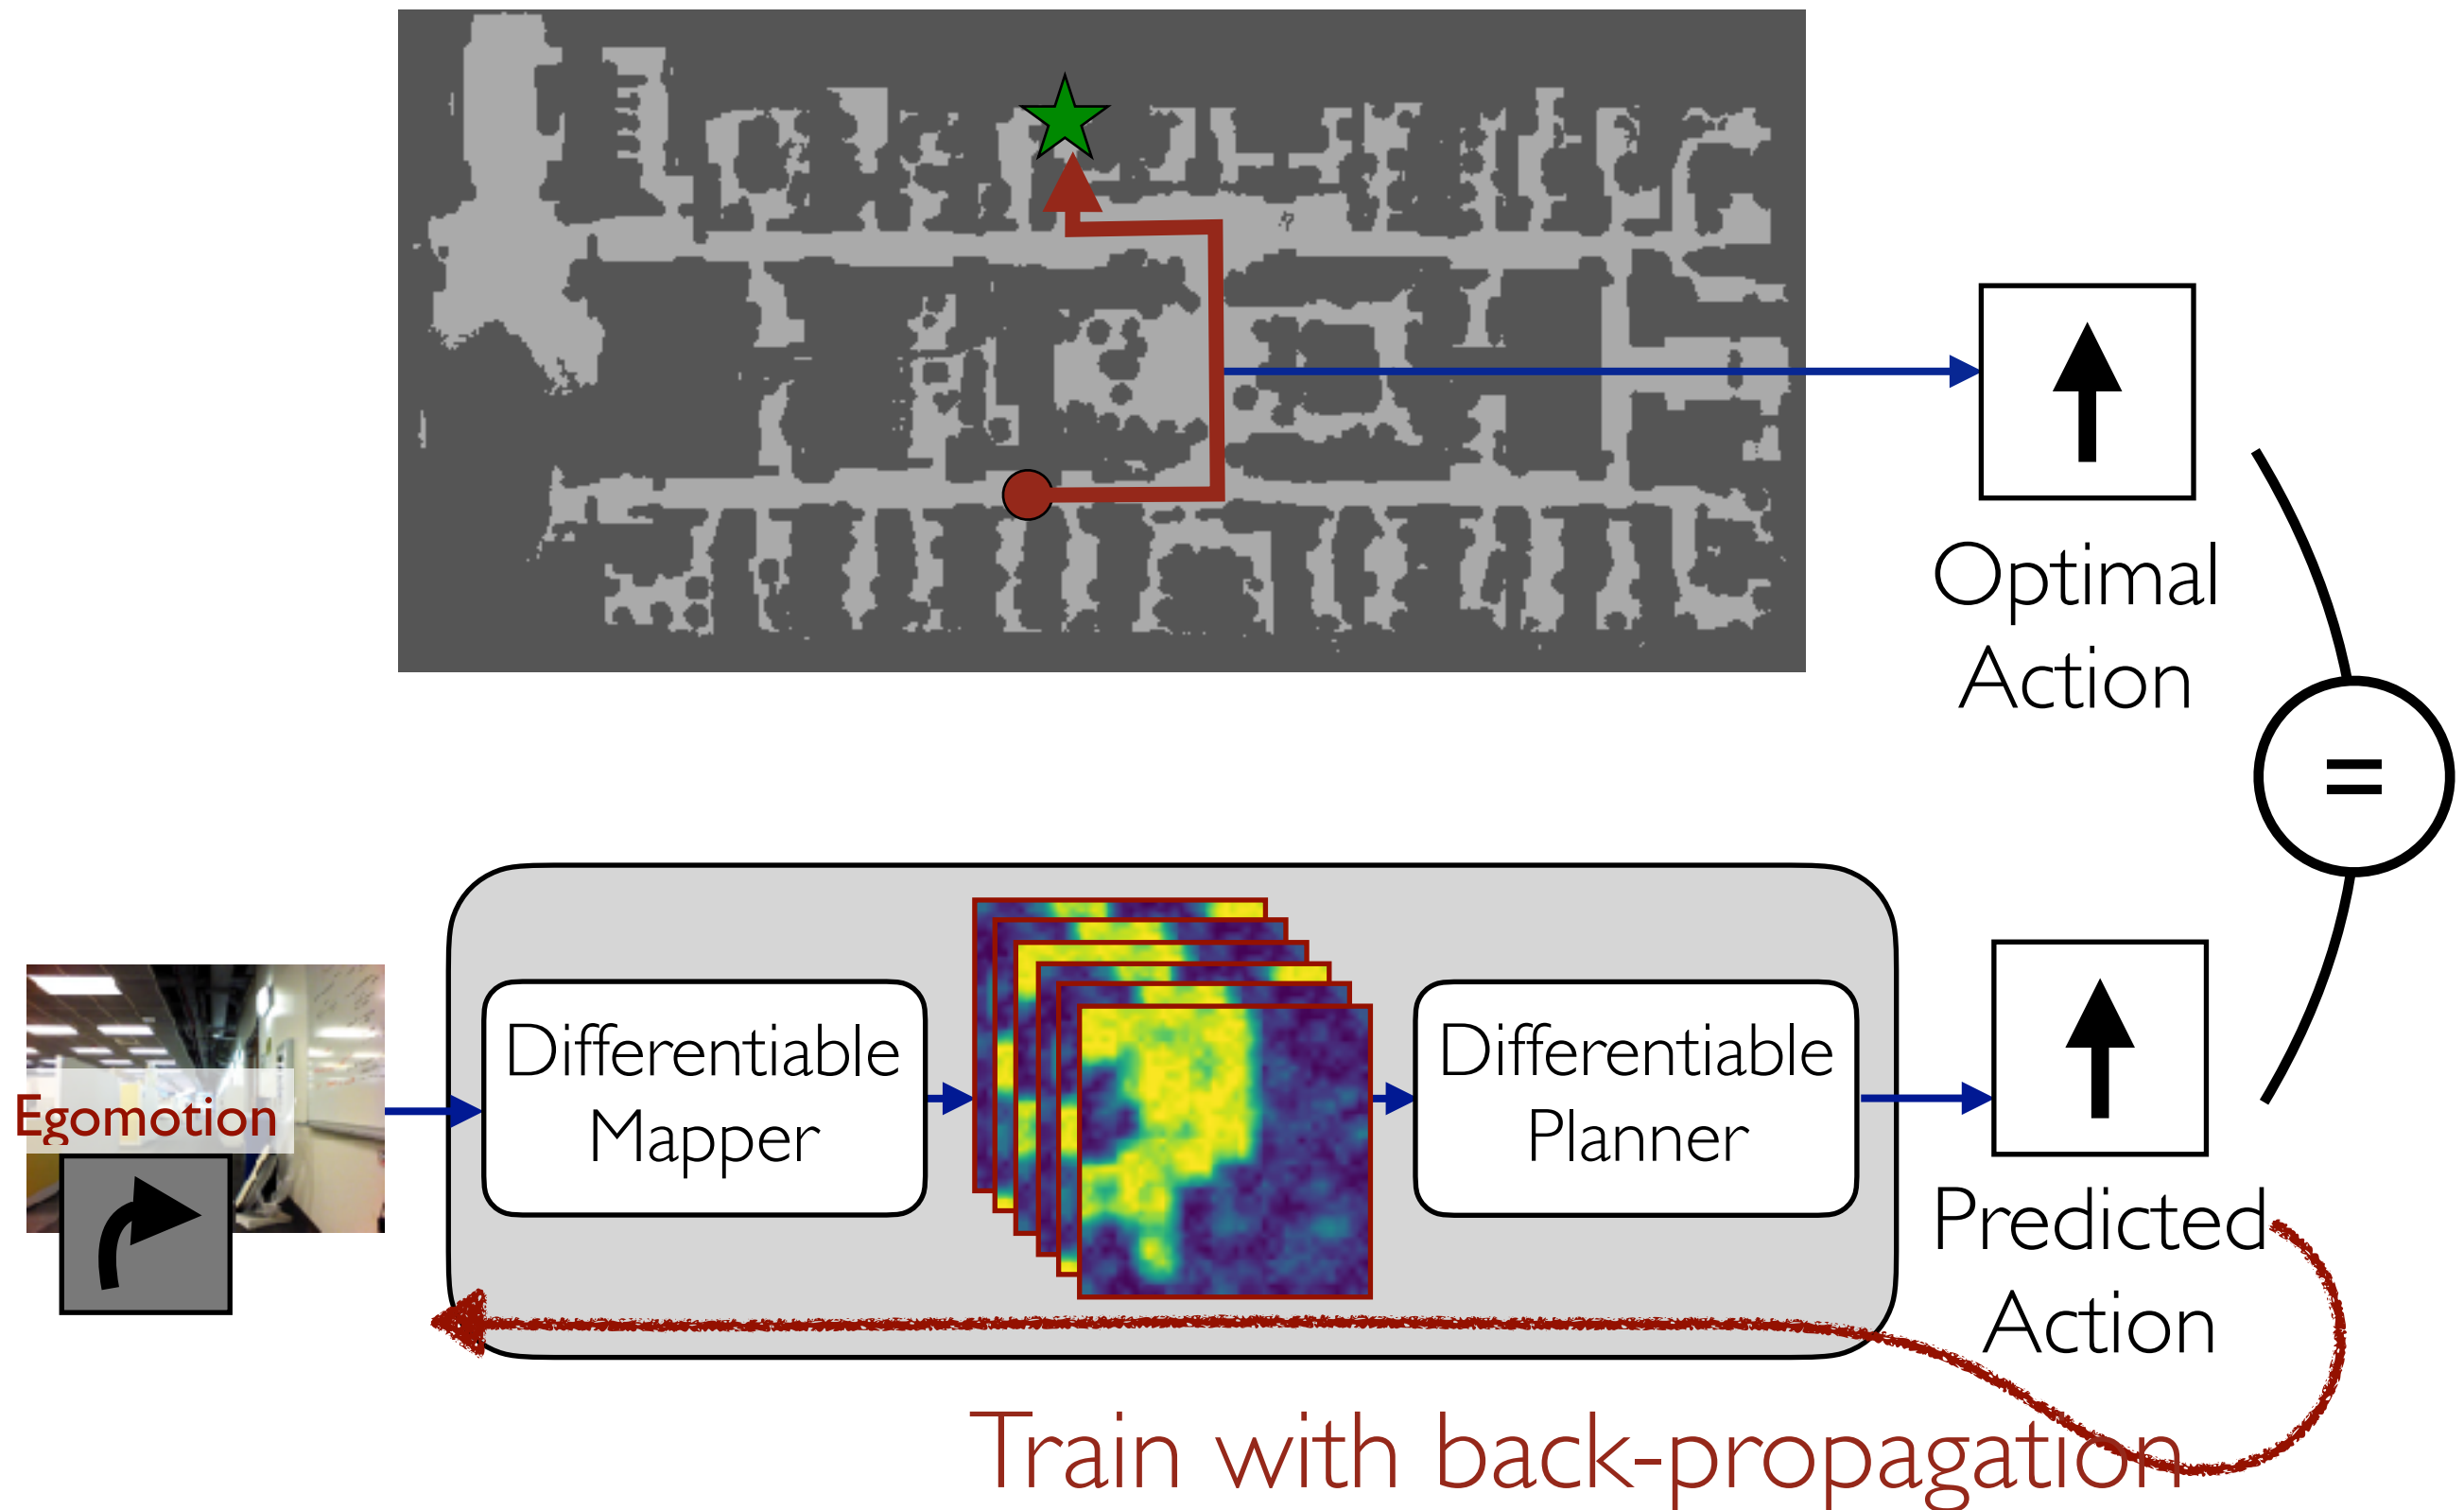
Supervision from ~~Human~~ Algorithmic Expert


Deep Sensorimotor Learning

rll.berkeley.edu/deeplearningrobotics

Department of Electrical Engineering and Computer Sciences
University of California, Berkeley

Supervision from ~~Human~~ Algorithmic Expert



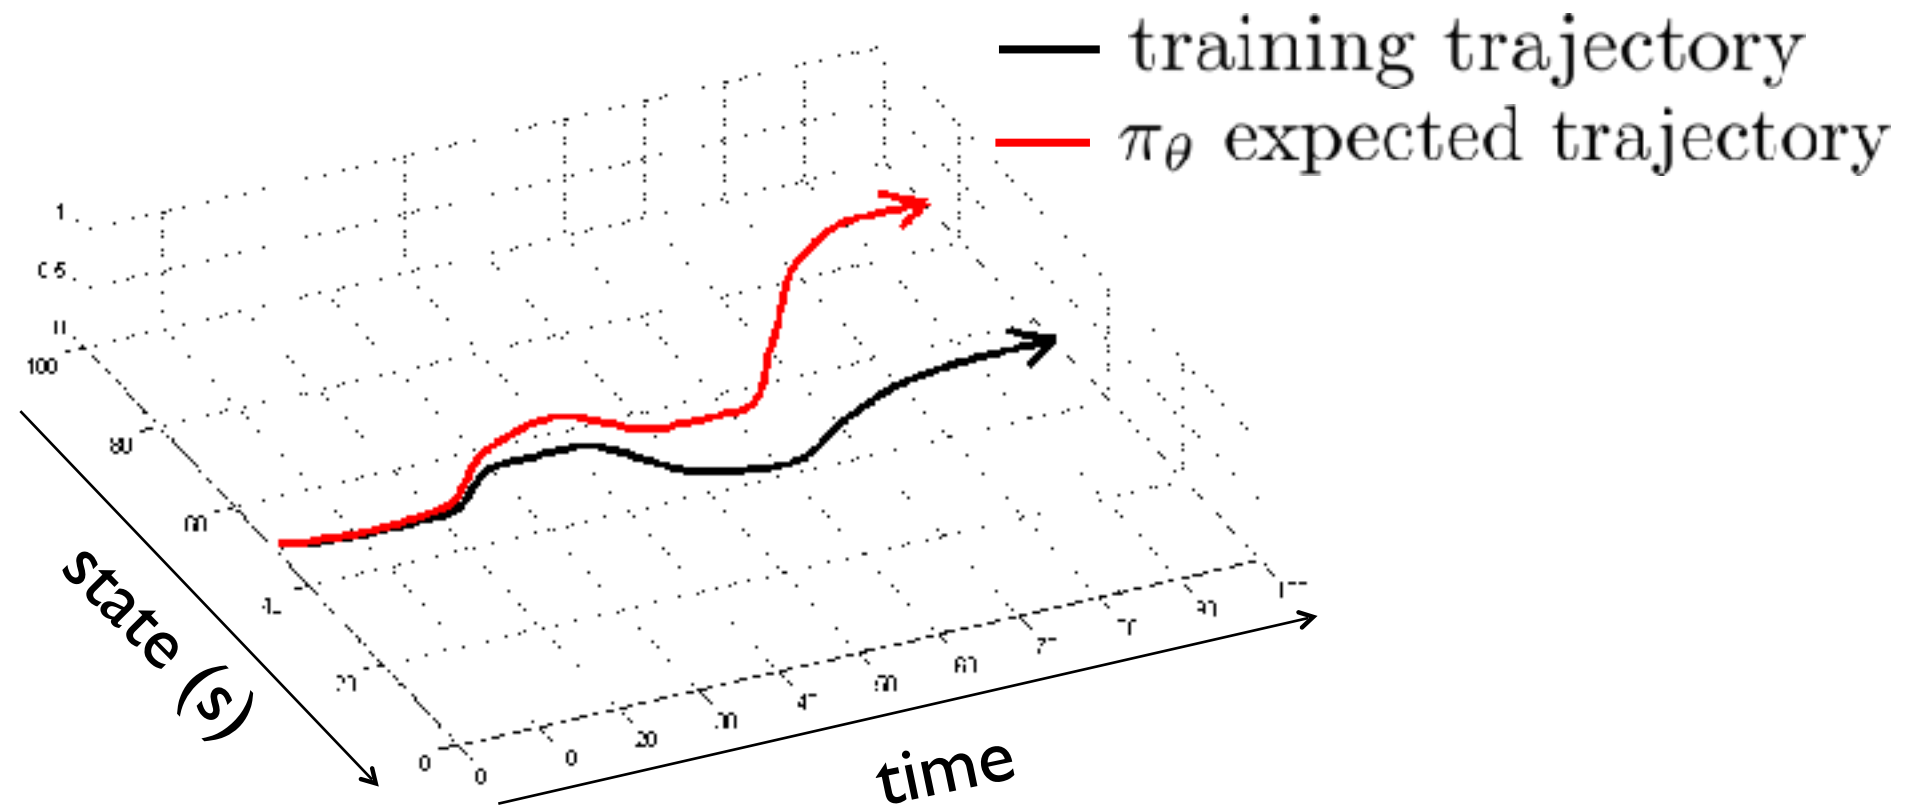


ALVINN: An Autonomous Land Vehicle in a Neural Network. Pomerleau. NeurIPS 1988.

Behavior Cloning

Does it always work?

No, data mis-match problem

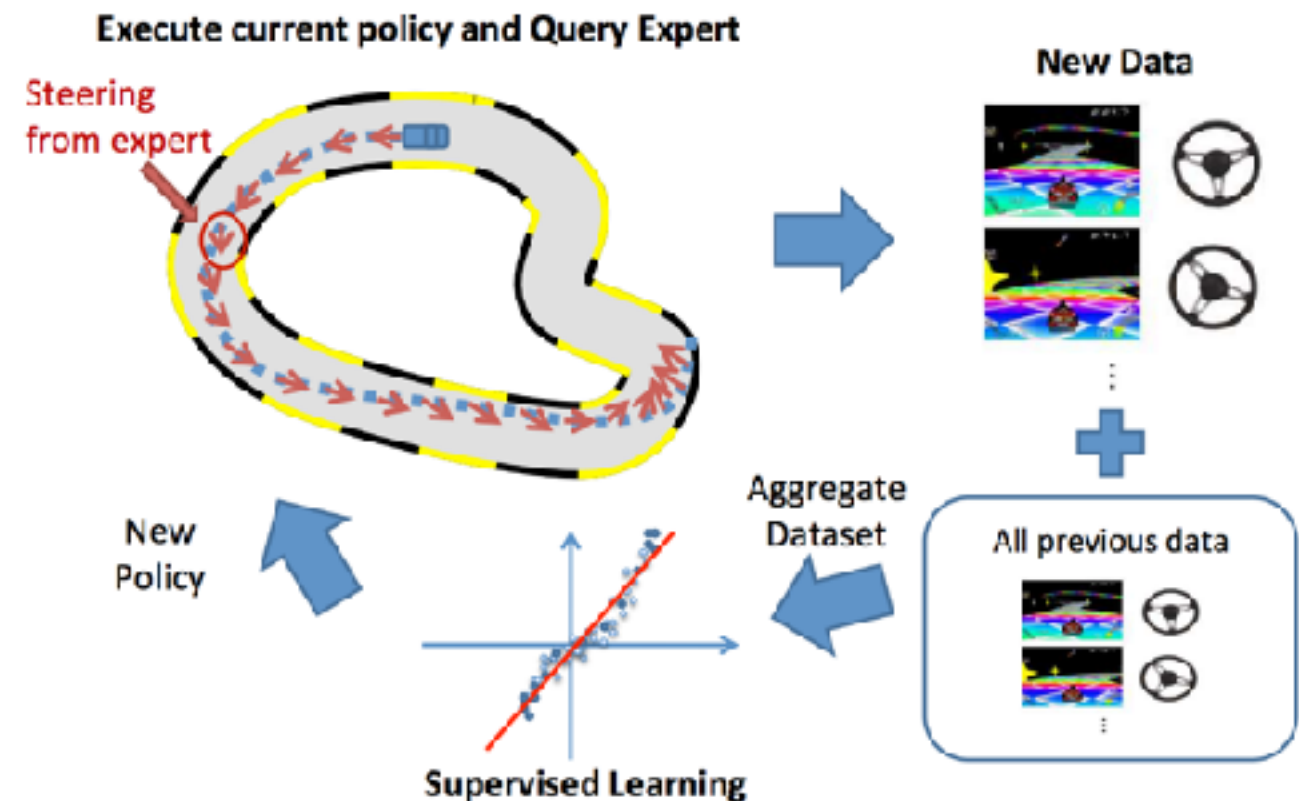


Fix Data Mis-Match Problem

*D*Agger: Dataset Aggregation

Collect labels on states visited by $\pi(o_t)$ instead of $\pi_e(o_t)$.

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_π with actions \mathbf{a}_t
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$



A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning


Stéphane Ross
Robotics Institute
Carnegie Mellon University
Pittsburgh, PA 15213, USA
stephaneross@cmu.edu

Geoffrey J. Gordon
Machine Learning Department
Carnegie Mellon University
Pittsburgh, PA 15213, USA
ggordon@cs.cmu.edu


J. Andrew Bagnell
Robotics Institute
Carnegie Mellon University
Pittsburgh, PA 15213, USA
dbagnell@ri.cmu.edu

Execute current policy and Query Expert

Steering
from expert



New Data




...



**Aggregate
Dataset**




**New
Policy**



Supervised Learning

All previous data



...

Forward Training and SMILe Algorithm

Initialize π_1^0, \dots, π_T^0 to query and execute π^* .
for $i = 1$ **to** T **do**
 Sample T -step trajectories by following π^{i-1} .
 Get dataset $\mathcal{D} = \{(s_i, \pi^*(s_i))\}$ of states, actions taken by expert at step i .
 Train classifier $\pi_i^i = \operatorname{argmin}_{\pi \in \Pi} \mathbb{E}_{s \sim \mathcal{D}}(e_\pi(s))$.
 $\pi_j^i = \pi_j^{i-1}$ for all $j \neq i$
end for
Return π_1^T, \dots, π_T^T

Algorithm 3.1: Forward Training Algorithm.

Initialize $\pi^0 \leftarrow \pi^*$ to query and execute expert.
for $i = 1$ **to** N **do**
 Execute π^{i-1} to get $\mathcal{D} = \{(s, \pi^*(s))\}$.
 Train classifier $\hat{\pi}^{*i} = \operatorname{argmin}_{\pi \in \Pi} \mathbb{E}_{s \sim \mathcal{D}}(e_\pi(s))$.
 $\pi^i = (1 - \alpha)^i \pi^* + \alpha \sum_{j=1}^i (1 - \alpha)^{j-1} \hat{\pi}^{*j}$.
end for
Remove expert queries: $\tilde{\pi}^N = \frac{\pi^N - (1 - \alpha)^N \pi^*}{1 - (1 - \alpha)^N}$
Return $\tilde{\pi}^N$

Algorithm 4.1: The SMILe Algorithm.

DAgger

```
Initialize  $\mathcal{D} \leftarrow \emptyset$ .  
Initialize  $\hat{\pi}_1$  to any policy in  $\Pi$ .  
for  $i = 1$  to  $N$  do  
    Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .  
    Sample  $T$ -step trajectories using  $\pi_i$ .  
    Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$   
    and actions given by expert.  
    Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .  
    Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .  
end for  
Return best  $\hat{\pi}_i$  on validation.
```

Algorithm 3.1: DAGGER Algorithm.

Super Tux Kart




Figure 2: Average falls/lap as a function of training data.

Super Mario Bros.




Figure 4: Average distance/stage as a function of data.

Structured Prediction




Figure 5: Character accuracy as a function of iteration.

Thank you