

Self-supervision

Saurabh Gupta

Solving a RL Problem

Better Optimization

**Solve a Related but
Supervision-rich Problem**

Model-free RL
with sparse
rewards

Better Reward Signals

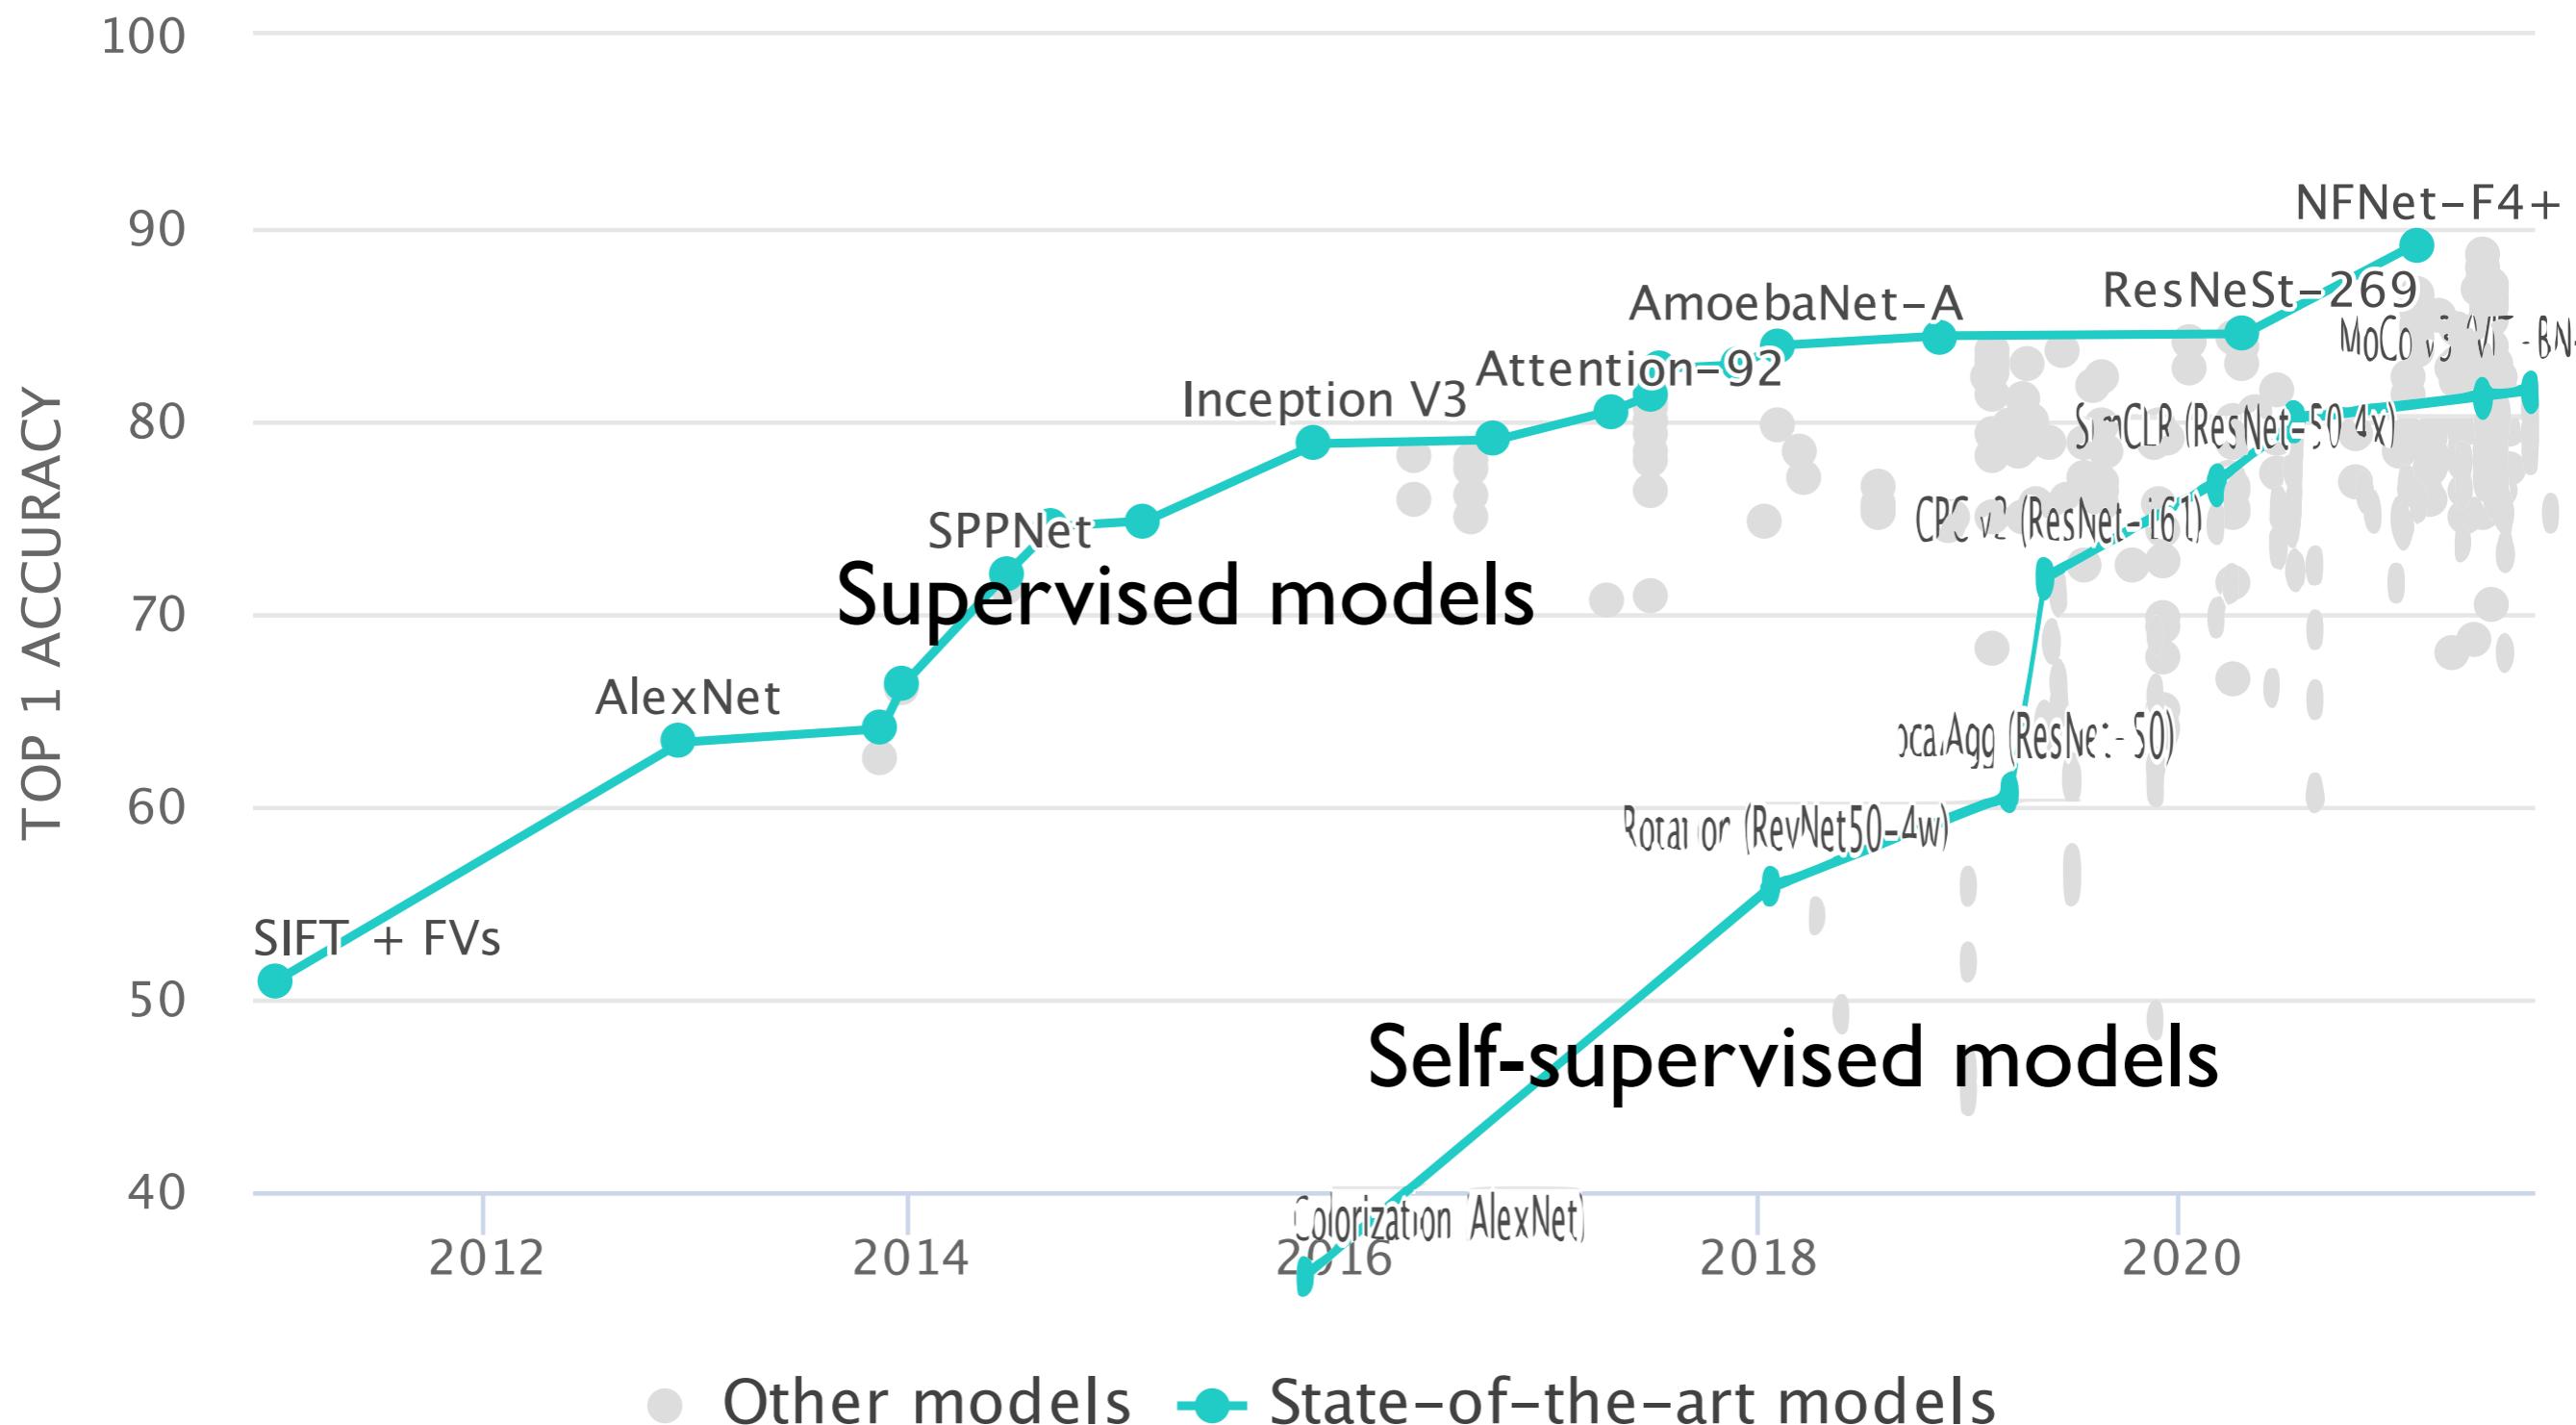
Convert into a
Supervised Training
Problem

Build Models and Plan
with Them

Known reward,
known model.
Model-based RL

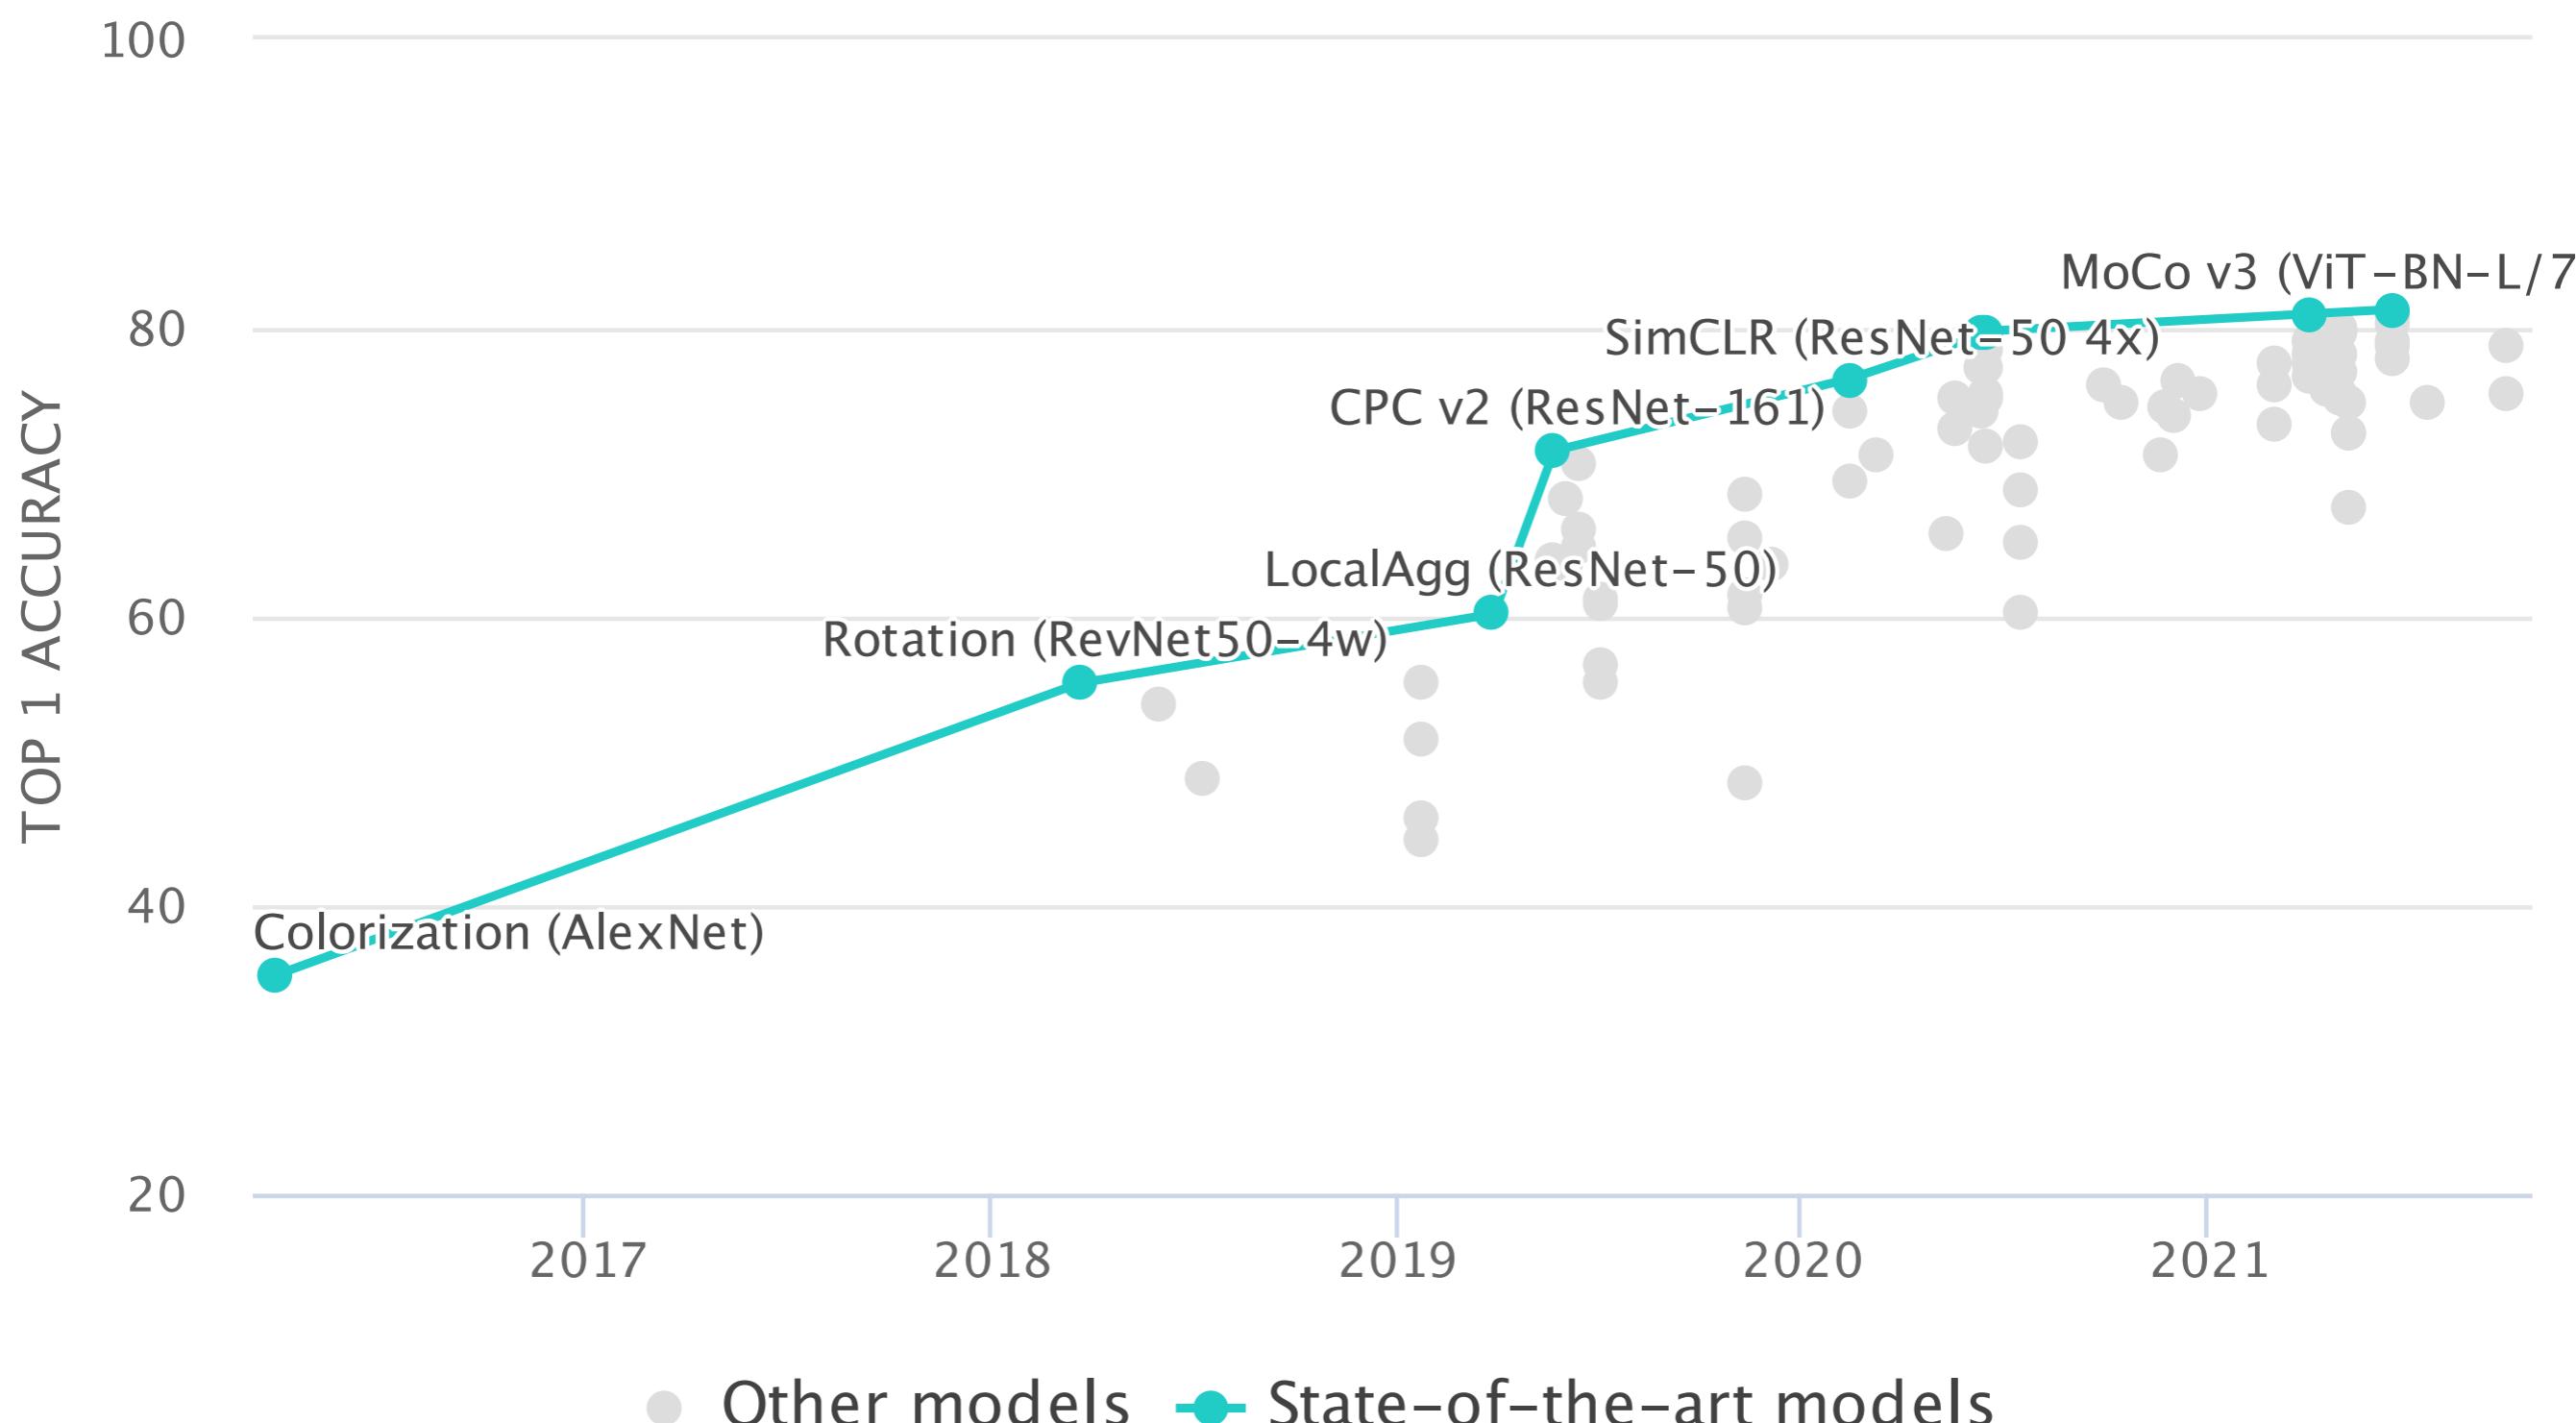


In computer vision, *Solve a related but easy to supervise problem*



In computer vision,

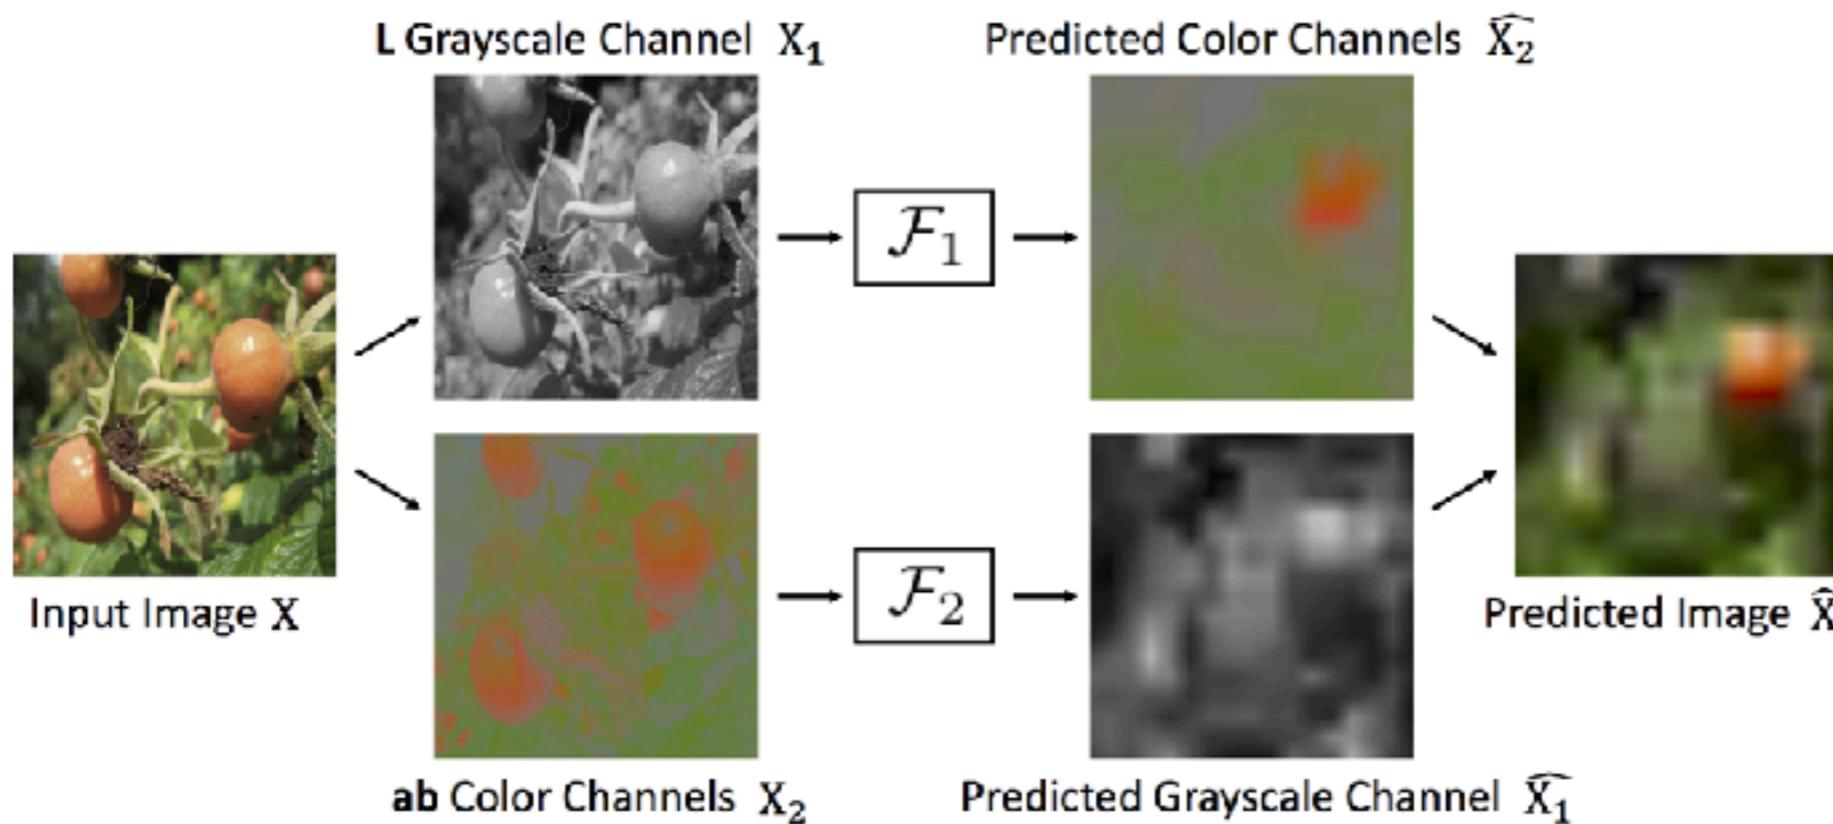
Solve a related but easy to supervise problem



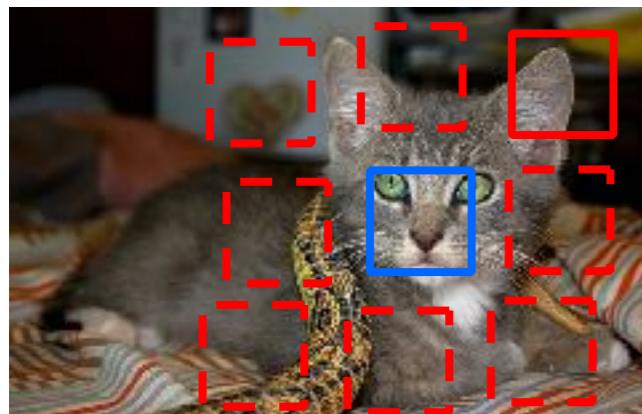
In computer vision,

Solve a related but easy to supervise problem

- Collecting large-scale labeled datasets is a bottleneck
- Can we train models without relying on semantic supervision?
- Kind of. Supervise using *pretext* tasks.
 - *Pretext task*: easy to supervise, yet lead to good learning
- Eg: colorization:



In computer vision,
Solve a related but easy to supervise problem



Images



Videos



Across modalities

Solve a Related but Supervision-rich Problem

Design setups that can allow rapid data gathering

Supersizing Self-supervision: Learning to Grasp from 50K Tries and 700 Robot Hours

Lerrel Pinto and Abhinav Gupta

The Robotics Institute, Carnegie Mellon University

(lerrelp, abhinavg)@cs.cmu.edu

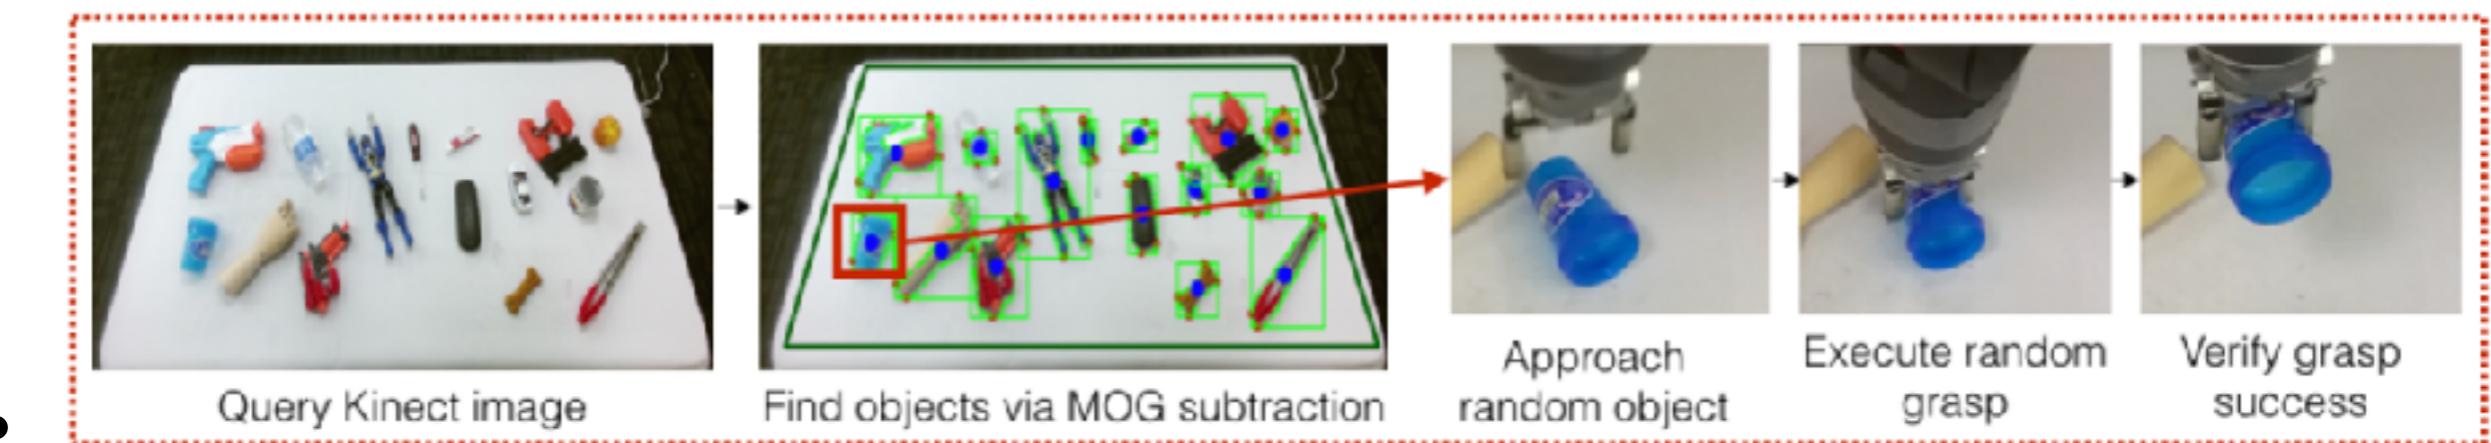
Problem Setup

- Grasp objects on a table-top
- Known robot, unknown objects
- Planar grasping, parameterized by (x, y, θ) .
- Observation space?
- Action space?



Data collection

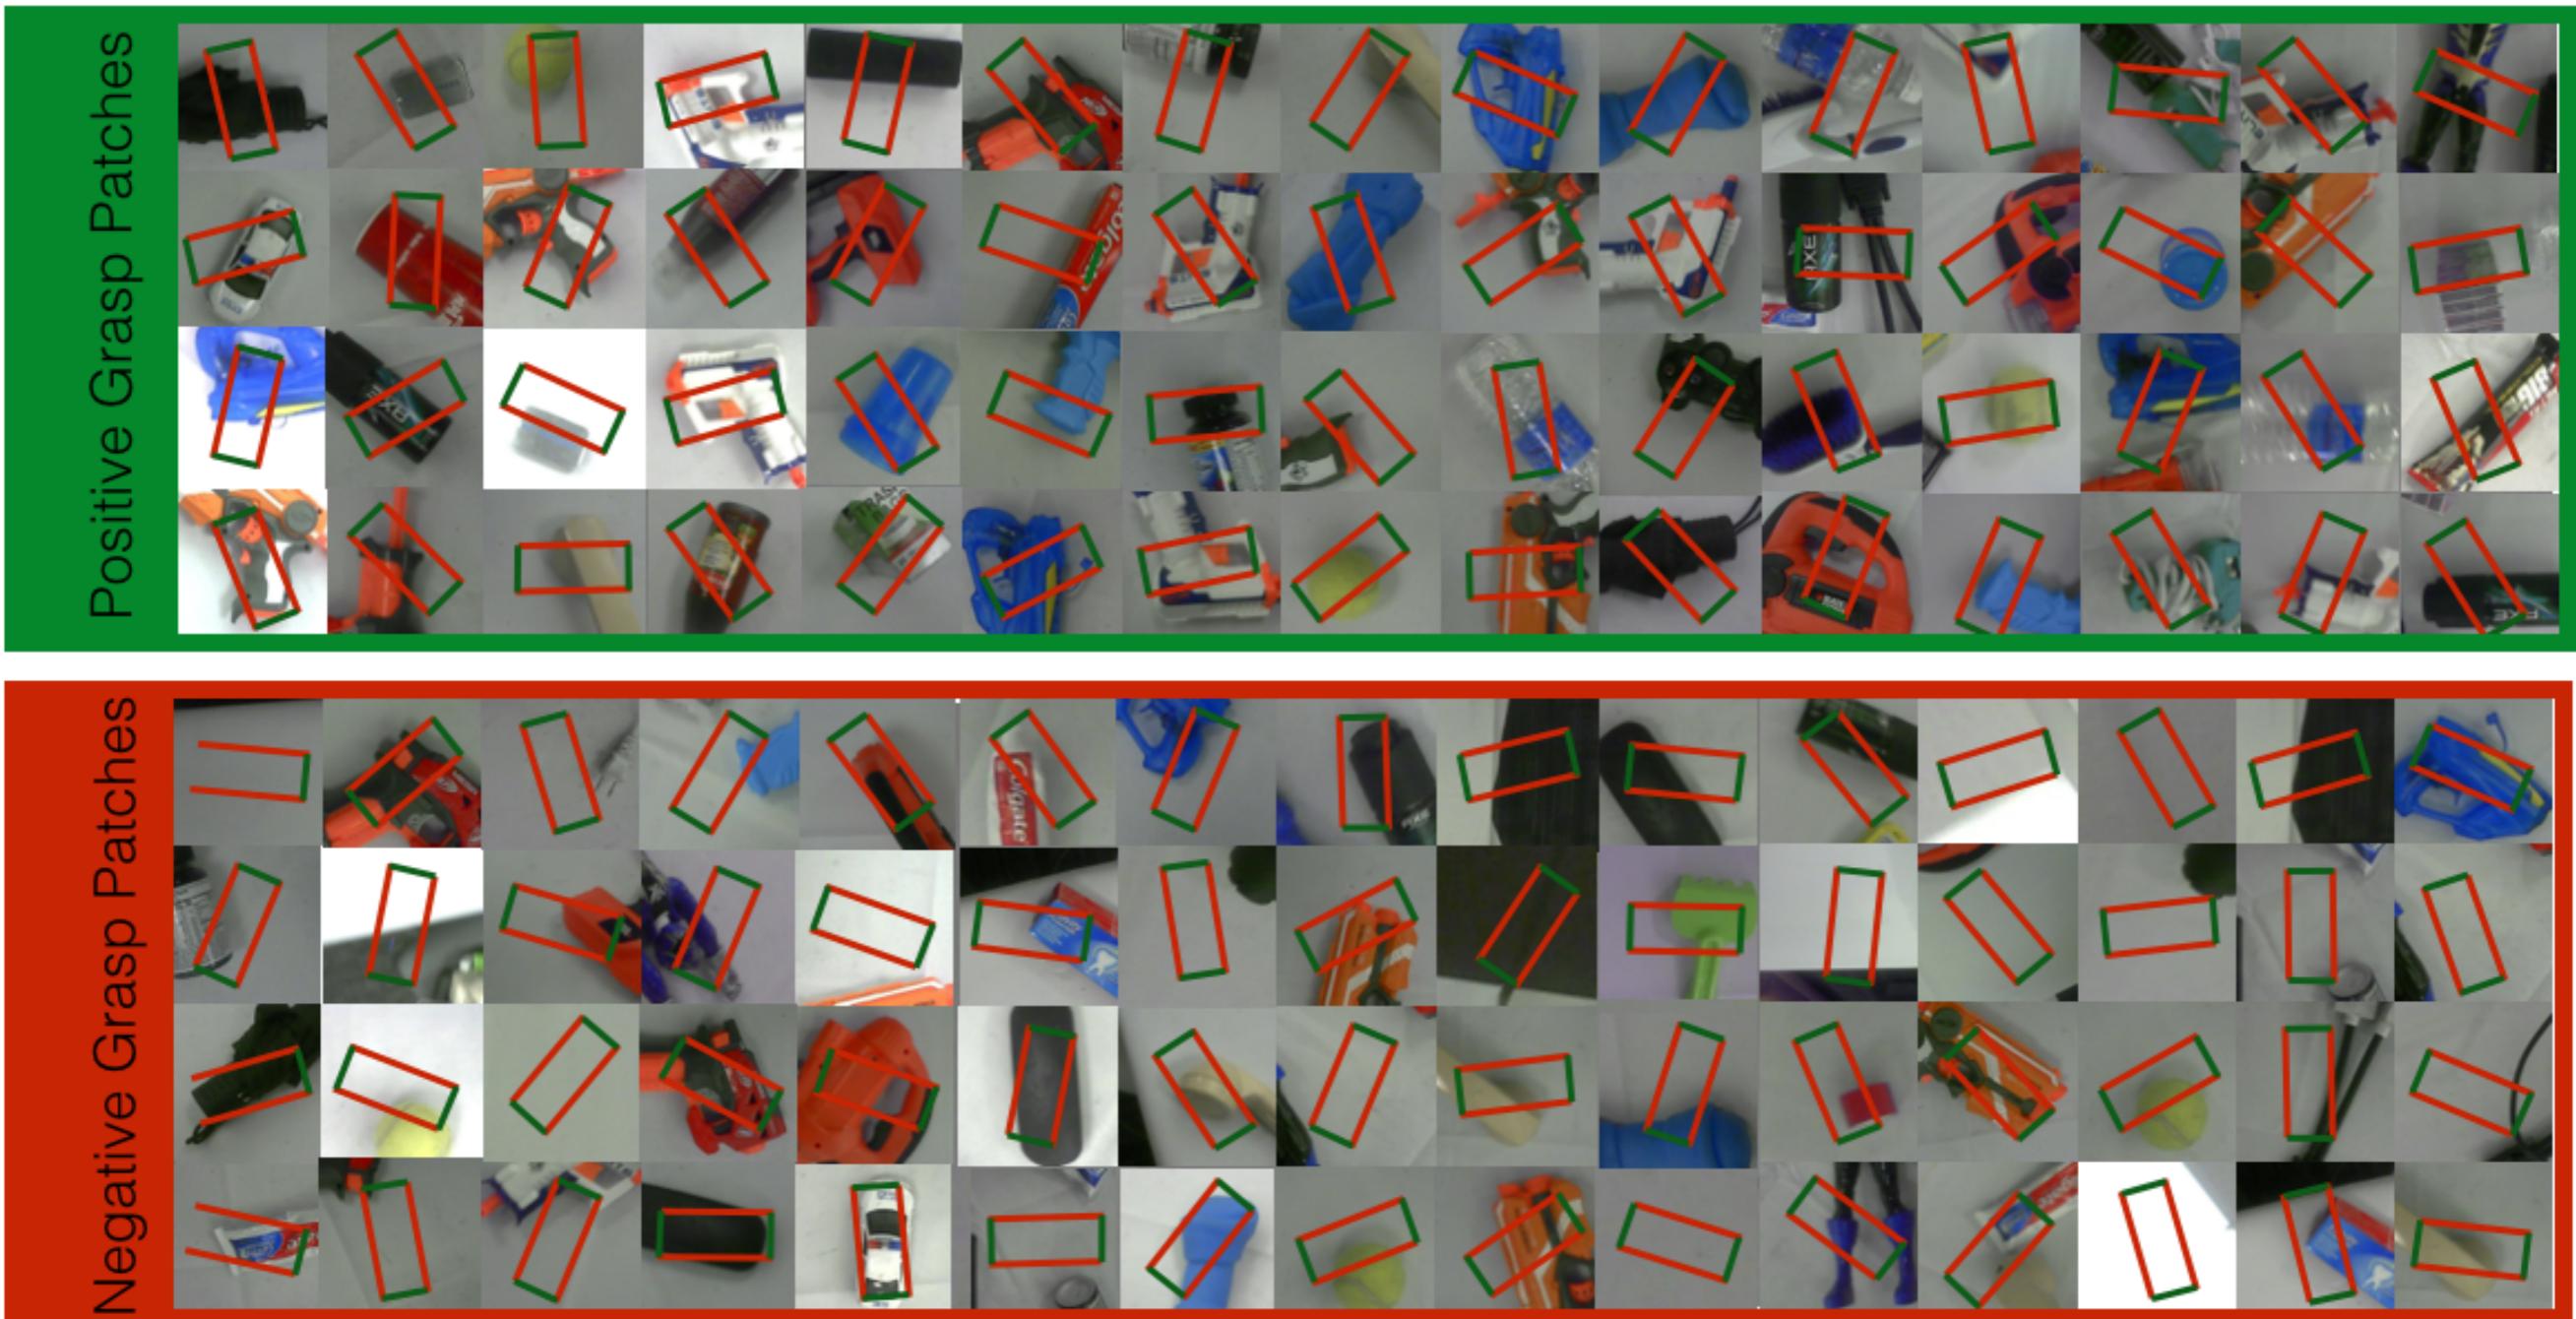
- Autonomous data collection:
 - force closure on gripper indicates grasp success / failure.
- Collect data over 700 hours.



Autonomous data collection and labeling

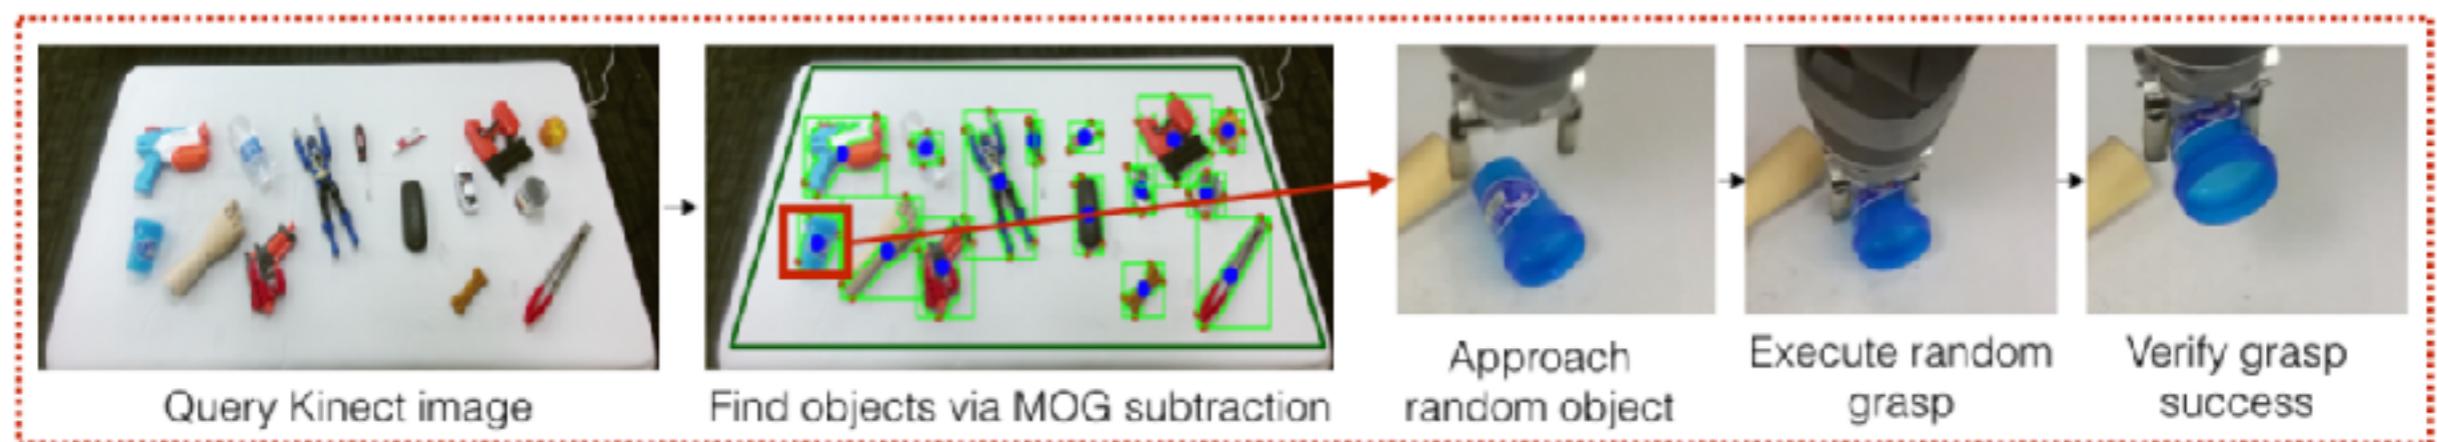


Collected Data



Data collection

- Autonomous data collection:
 - force closure on gripper indicates grasp success / failure.
 - Collect data over 700 hours.

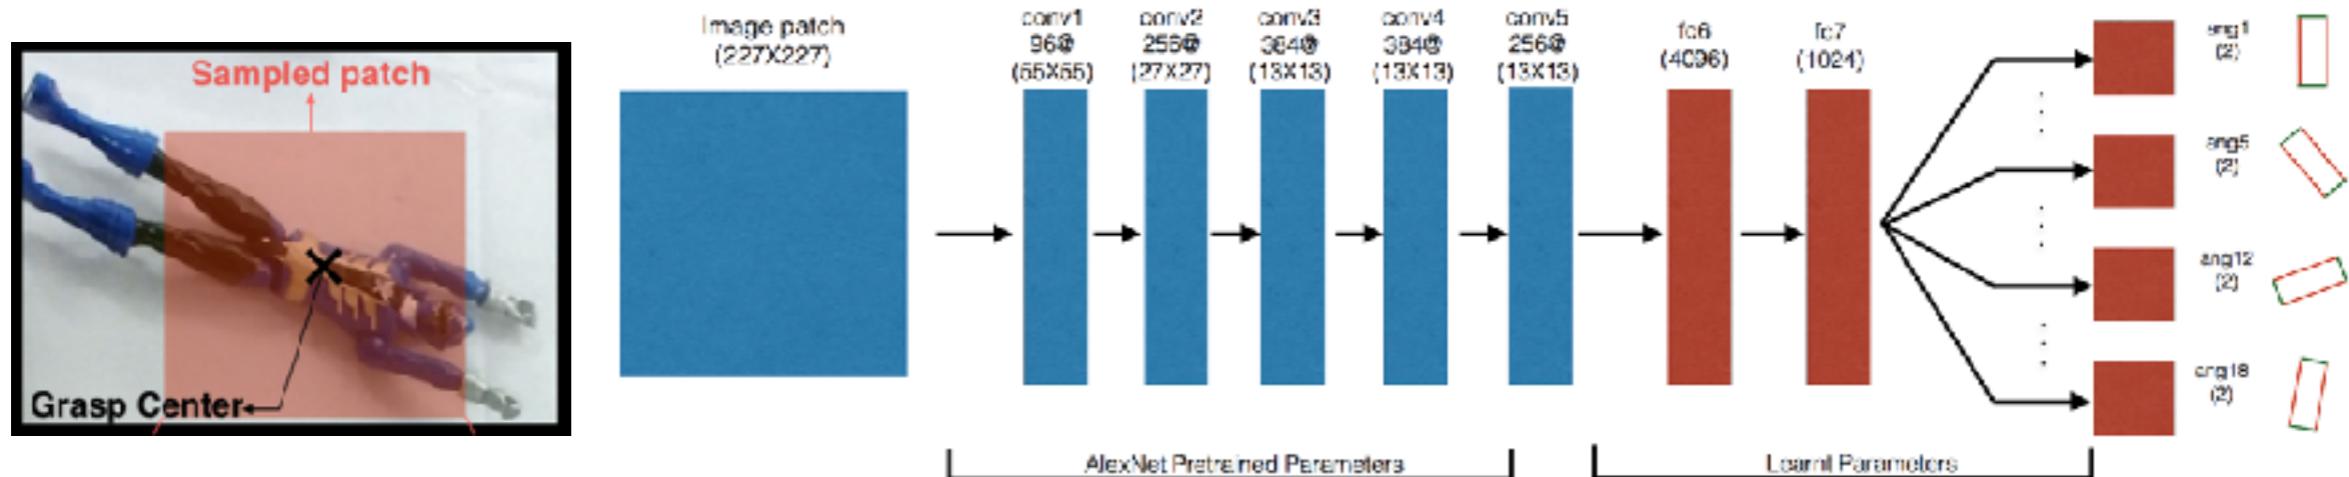


- Iterative data collection

Data Collection Type	Positive	Negative	Total	Grasp Rate
Random Trials	3,245	37,042	40,287	8.05%
Multi-Staged	2,807	4,500	7,307	38.41%
Test Set	214	2,759	2,973	7.19%
	6,266	44,301	50,567	

Training

- Train neural network for prediction grasp success rate



- $$L_B = \sum_{i=1}^B \sum_{j=1}^{N=18} \delta(j, \theta_i) \cdot \text{softmax}(A_{ji}, l_i)$$

Testing

- Off-line evaluation
 - measure graspability prediction on novel objects
- On-robot Grasping:
 - Grasp objects
 - Clearing clutter

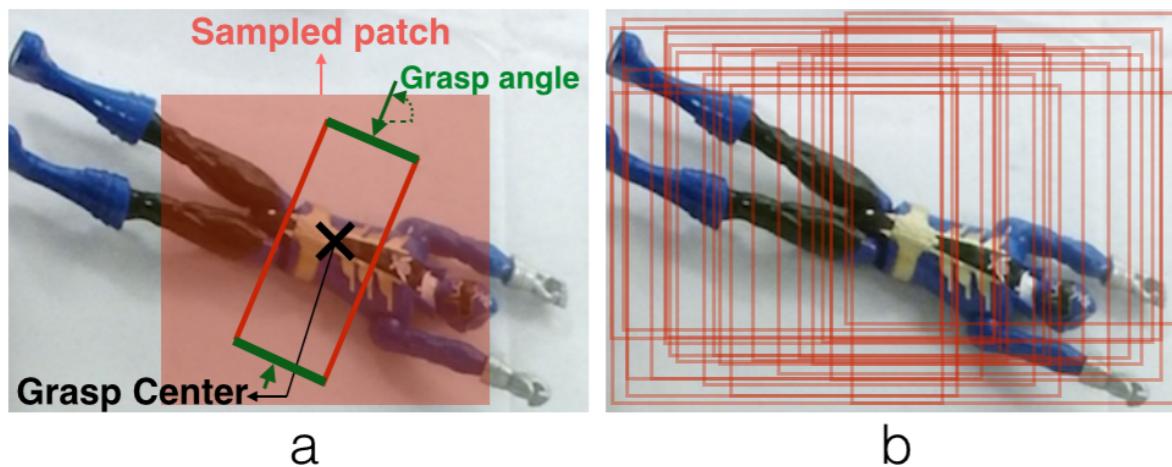


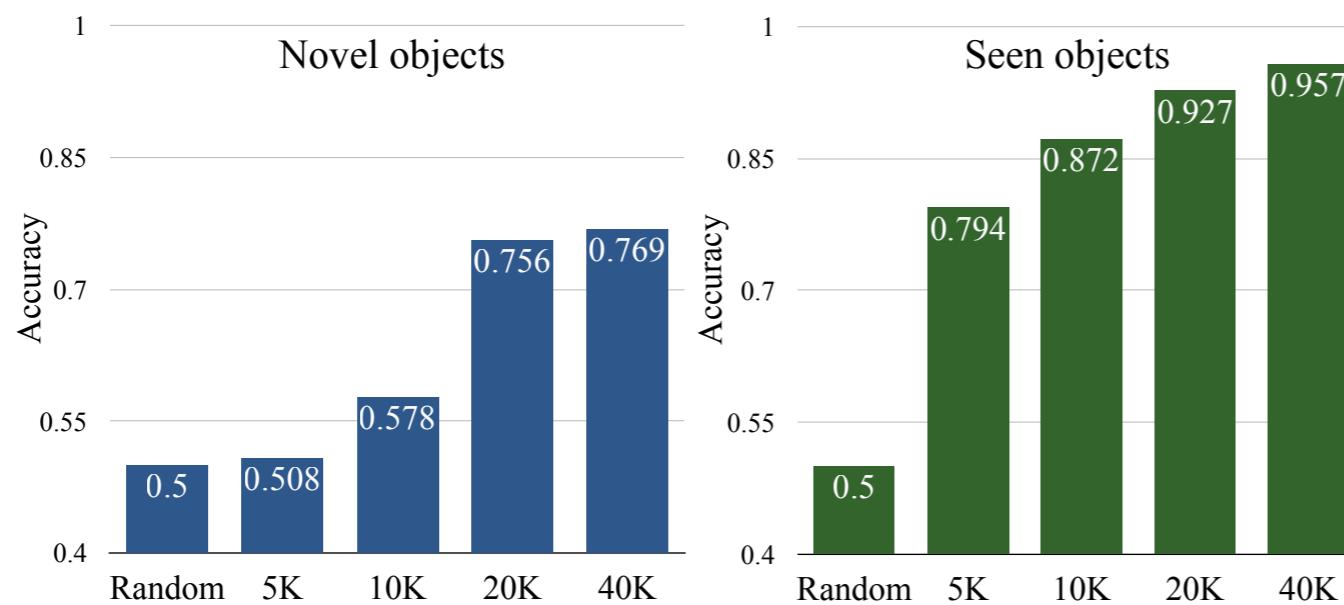
Fig. 3. (a) We use 1.5 times the gripper size image patch to predict the grasp-ability of a location and the angle at which it can be grasped. Visualization for showing the grasp location and the angle of gripper for grasping is derived from [8]. (b) At test time we sample patches at different positions and choose the top graspable location and corresponding gripper angle.

Off-line Testing

- Comparison to heuristic-methods

	Heuristic			Learning based			
	Min eigenvalue	Eigenvalue limit	Optimistic param. select	kNN	SVM	Deep Net (ours)	Deep Net + Multi-stage (ours)
Accuracy	0.534	0.599	0.621	0.694	0.733	0.769	0.795

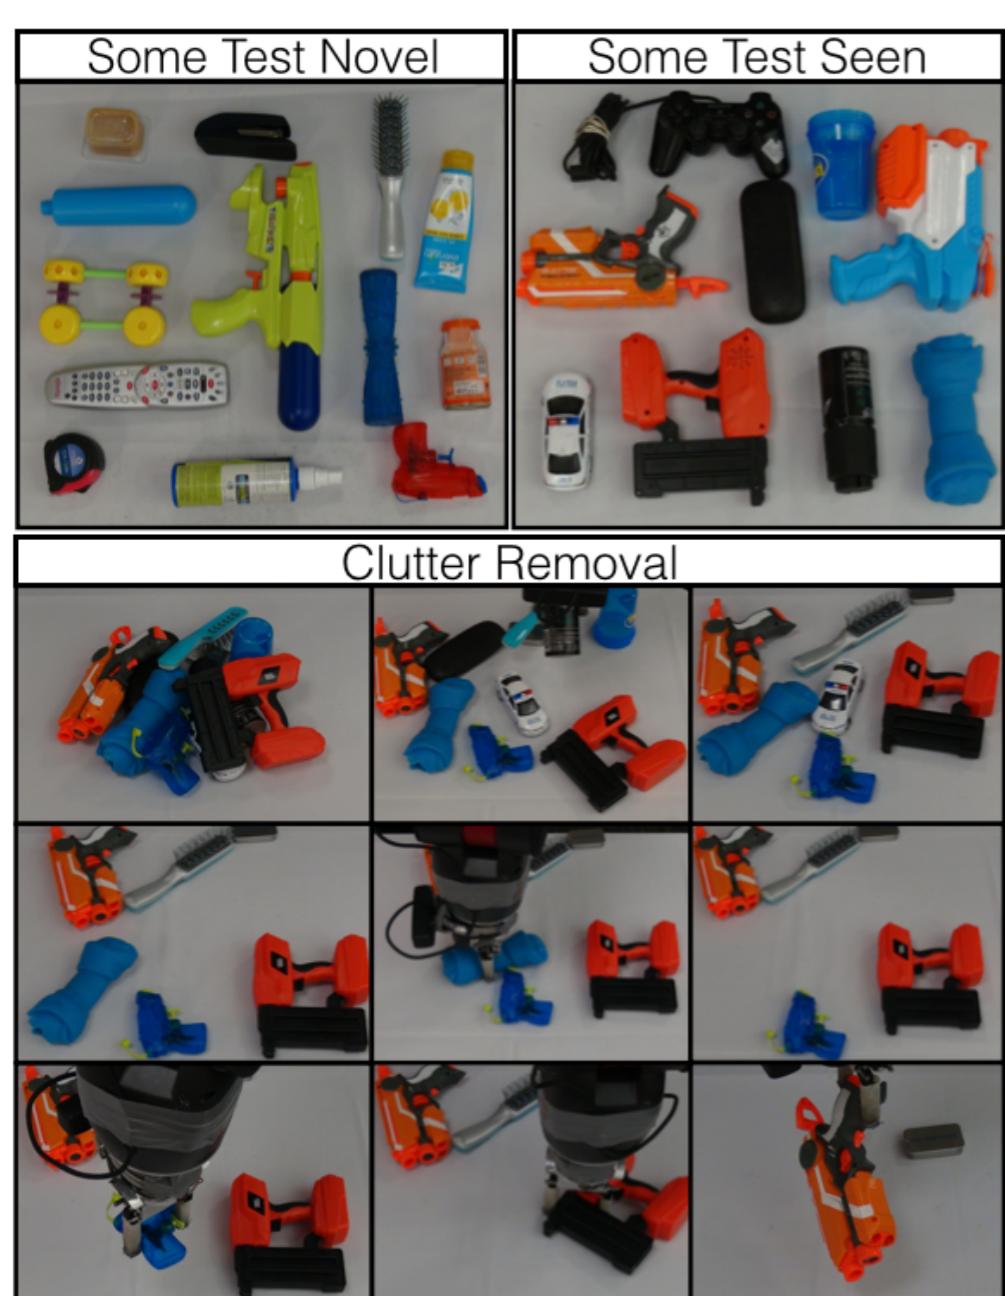
- Effect of amount of training data:

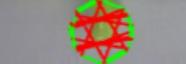
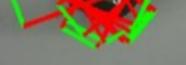
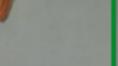
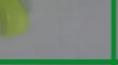
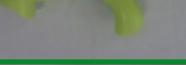
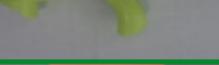
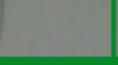
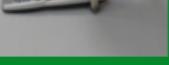
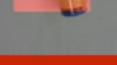


- Multi-stage data-gathering helps: accuracy goes up from 76.9 to 79.5.

On-line Testing

- Grasping rate for novel objects: 66%
 - Grasping rate for seen objects: 73%
 - Clutter removal: 10 object clutter, 26 interactions to clear clutter



Objects	Top Predictions	Executed Prediction	Robot Grasp	Objects	Failed Grasps
					
					
					
					
					
					
					

Learning Synergies between Pushing and Grasping with Self-supervised Deep Reinforcement Learning

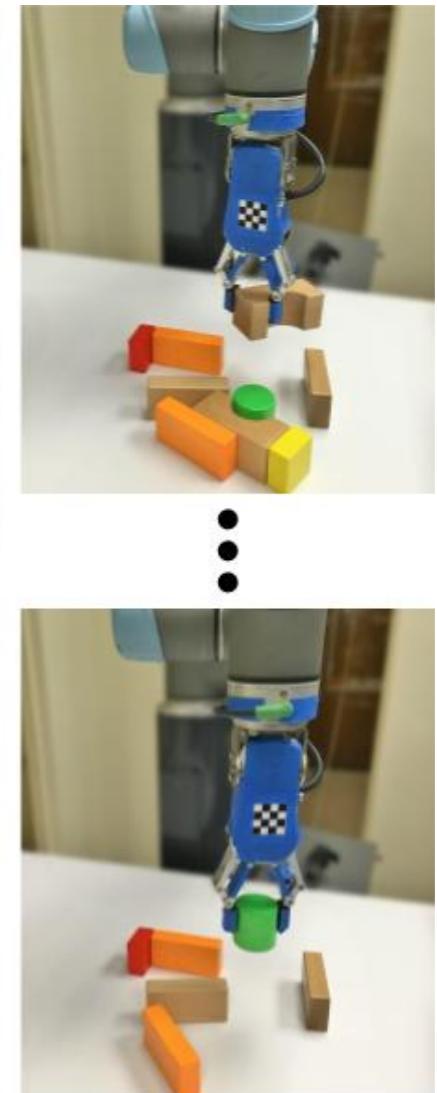
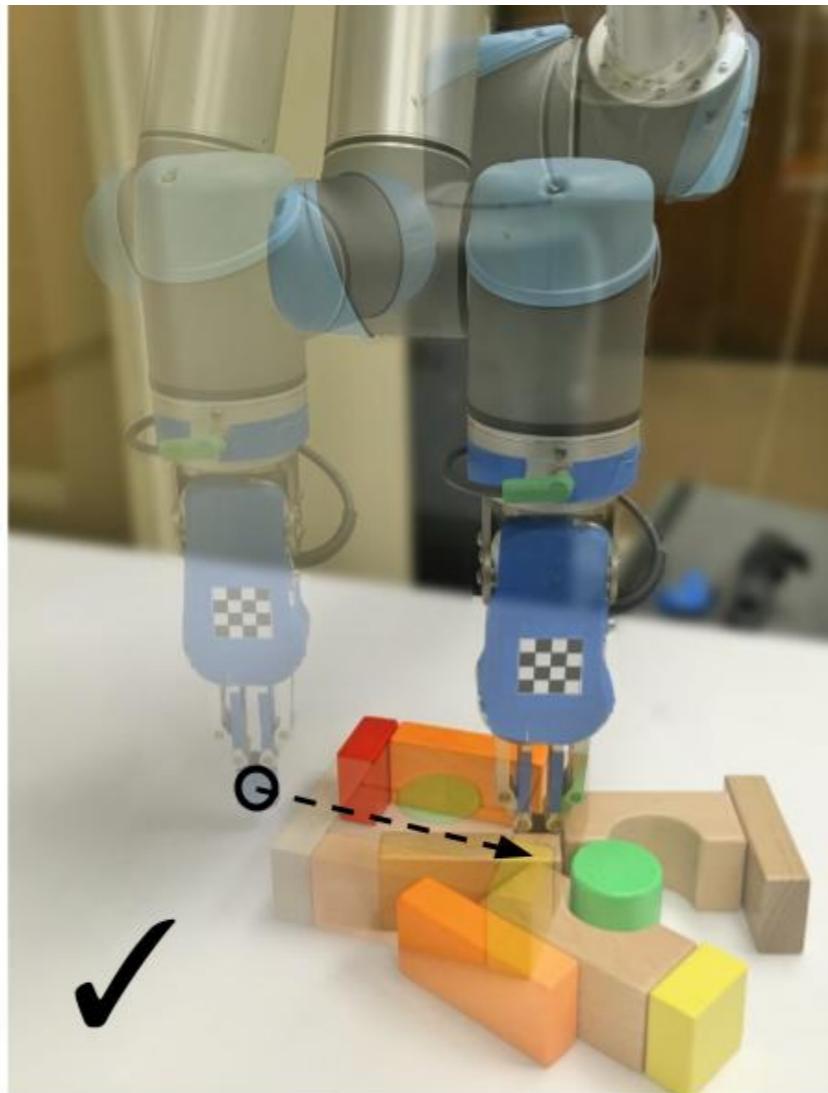
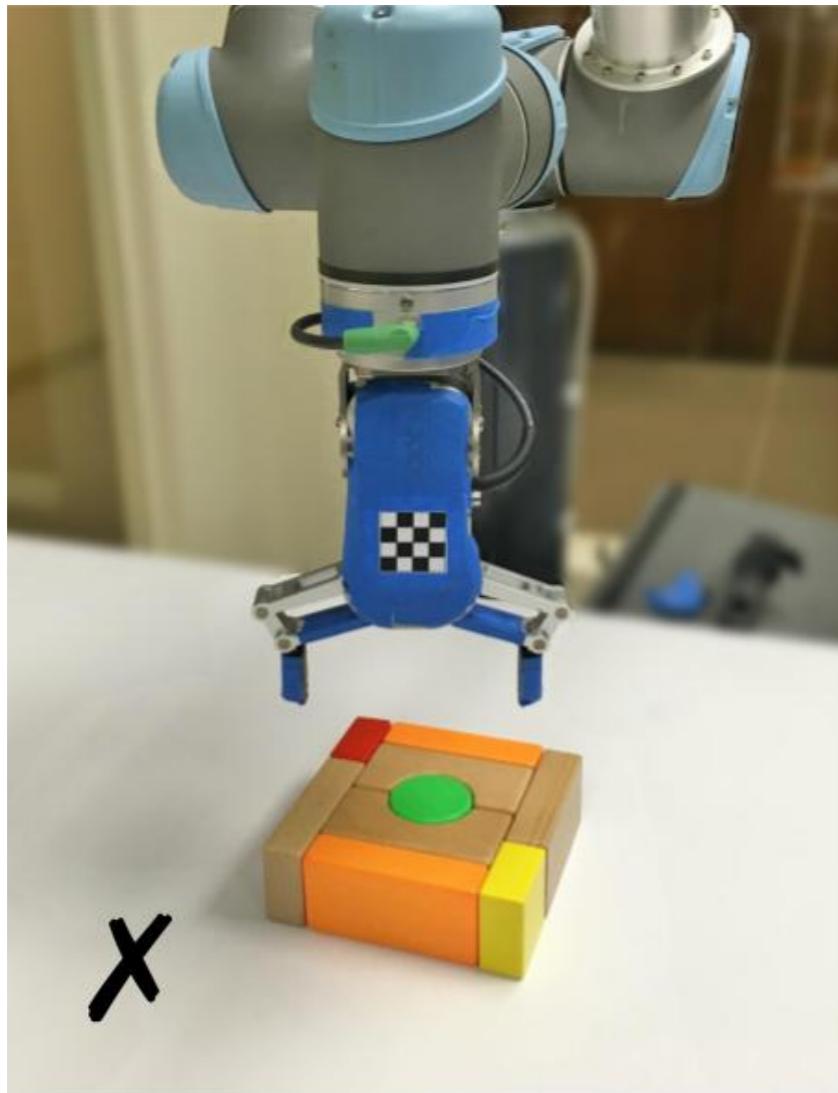
Andy Zeng^{1,2}, Shuran Song^{1,2}, Stefan Welker², Johnny Lee², Alberto Rodriguez³, Thomas Funkhouser^{1,2}

¹Princeton University

²Google

³Massachusetts Institute of Technology

<http://vpg.cs.princeton.edu>

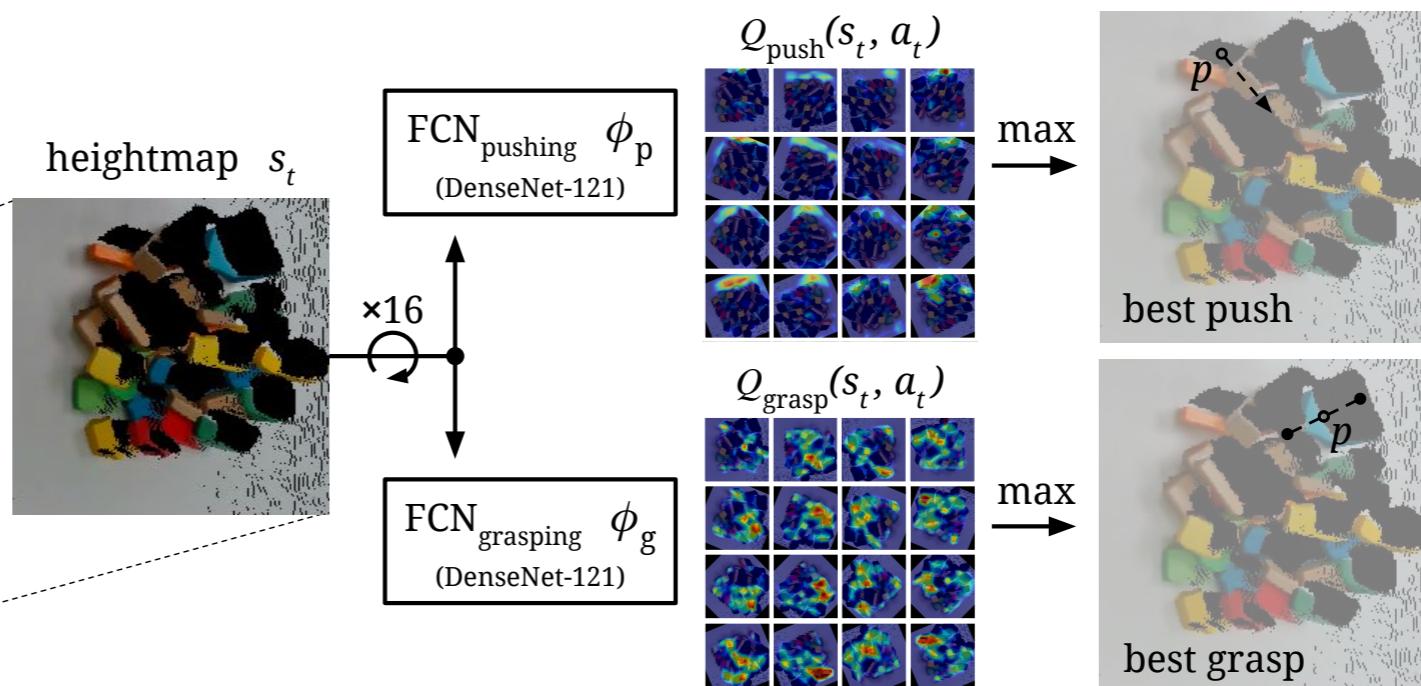
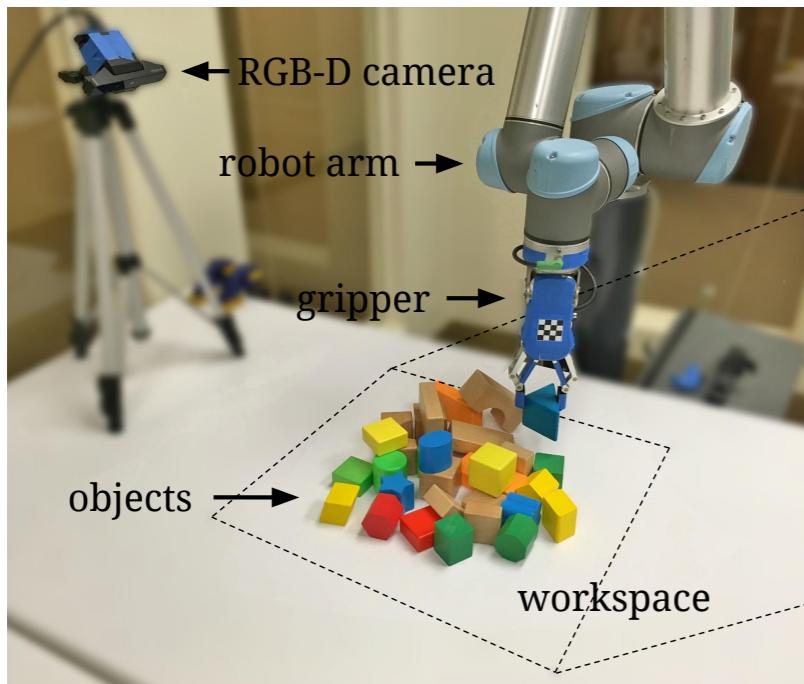


Summary

- Learning of pushing motion that is useful for grasping
- Clever parameterization of Q-functions
- Thorough evaluation of design choices in simulation

Q-functions

- Action space?
- Observation space?



- Rewards:
 - $R_g + 1$ if successfully grasped
 - $R_p + 0.5$ if push caused change in scene (SSD)
 - gamma = 0.5

Simulation Results

- Grasping random arrangements

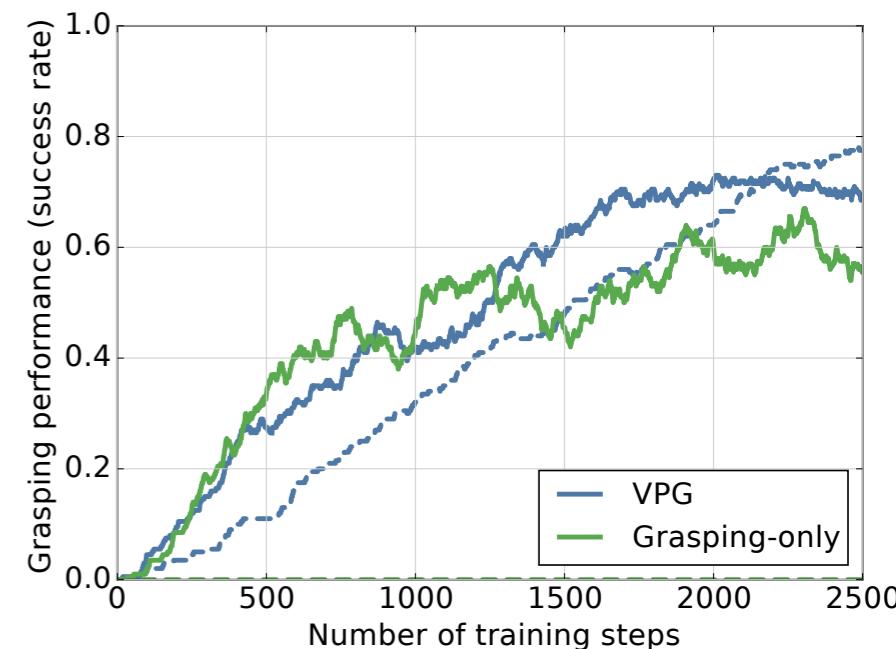
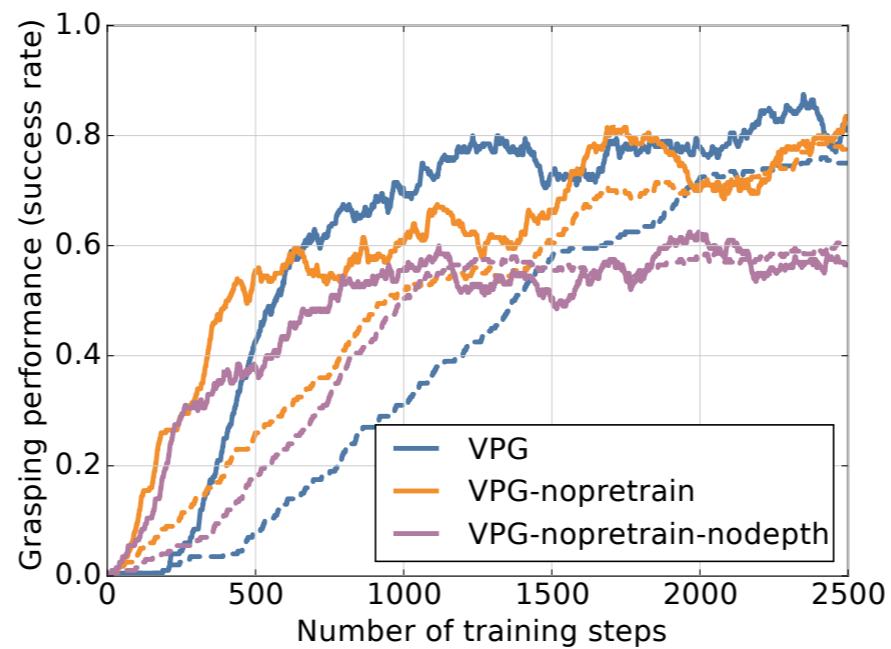
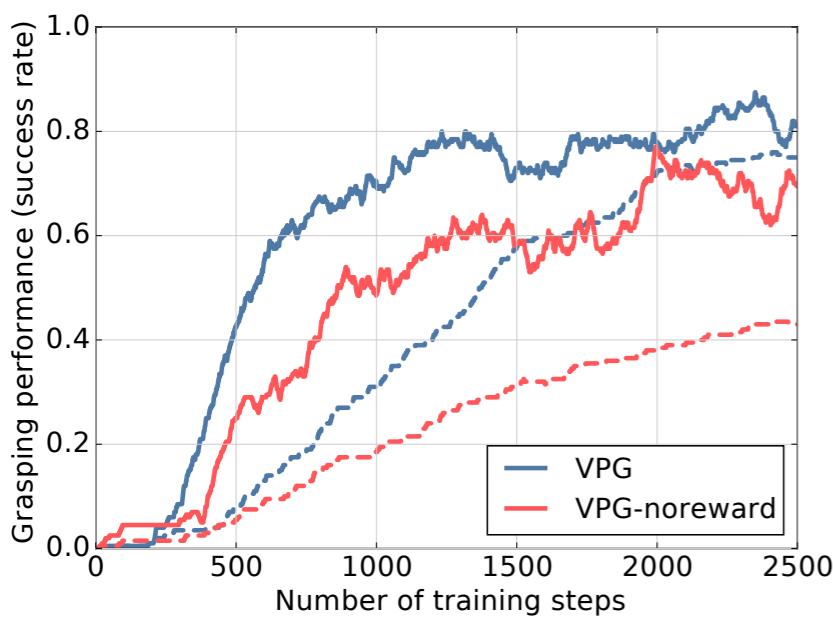
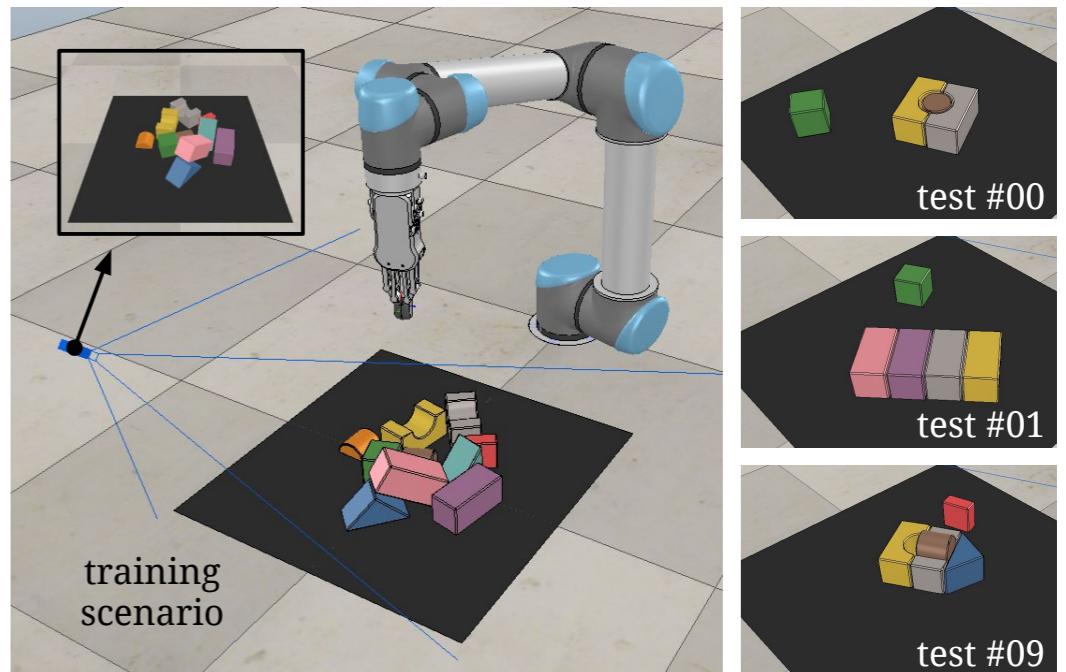
Method	Completion	Grasp Success	Action Efficiency
Grasping-only [8]	90.9	55.8	55.8
P+G Reactive	54.5	59.4	47.7
VPG	100.0	67.7	60.9

-

- Grasping adversarial arrangements

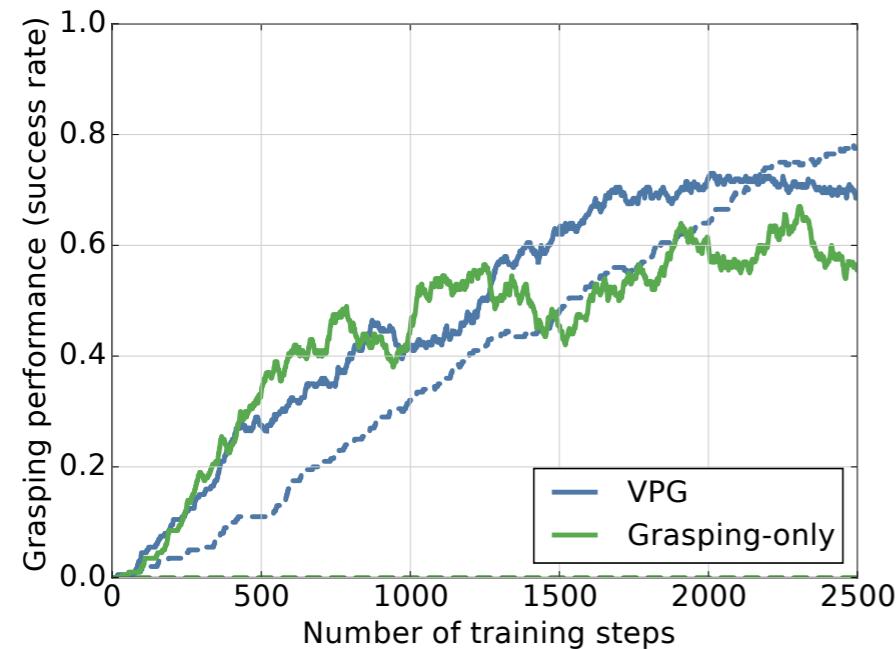
Method	Completion	Grasp Success	Action Efficiency
Grasping-only [8]	40.6	51.7	51.7
P+G Reactive	48.2	59.0	46.4
VPG	82.7	77.2	60.1

-



Method	Completion	Grasp Success	Action Efficiency
VPG-myopic	79.1	74.3	53.7
VPG	82.7	77.2	60.1

Real World Results



Method	Completion	Grasp Success	Action Efficiency
Grasping-only [8]	42.9	43.5	43.5
VPG	71.4	83.3	69.0

Solve a Related but Supervision-rich Problem



S. Levine et al. Learning Hand-Eye Coordination for Robotic Grasping with Deep Learning and Large-Scale Data Collection. ICLR 2017.

Solve a Related but Supervision-rich Problem

One-shot grasping often
leads to failed grasp attempts

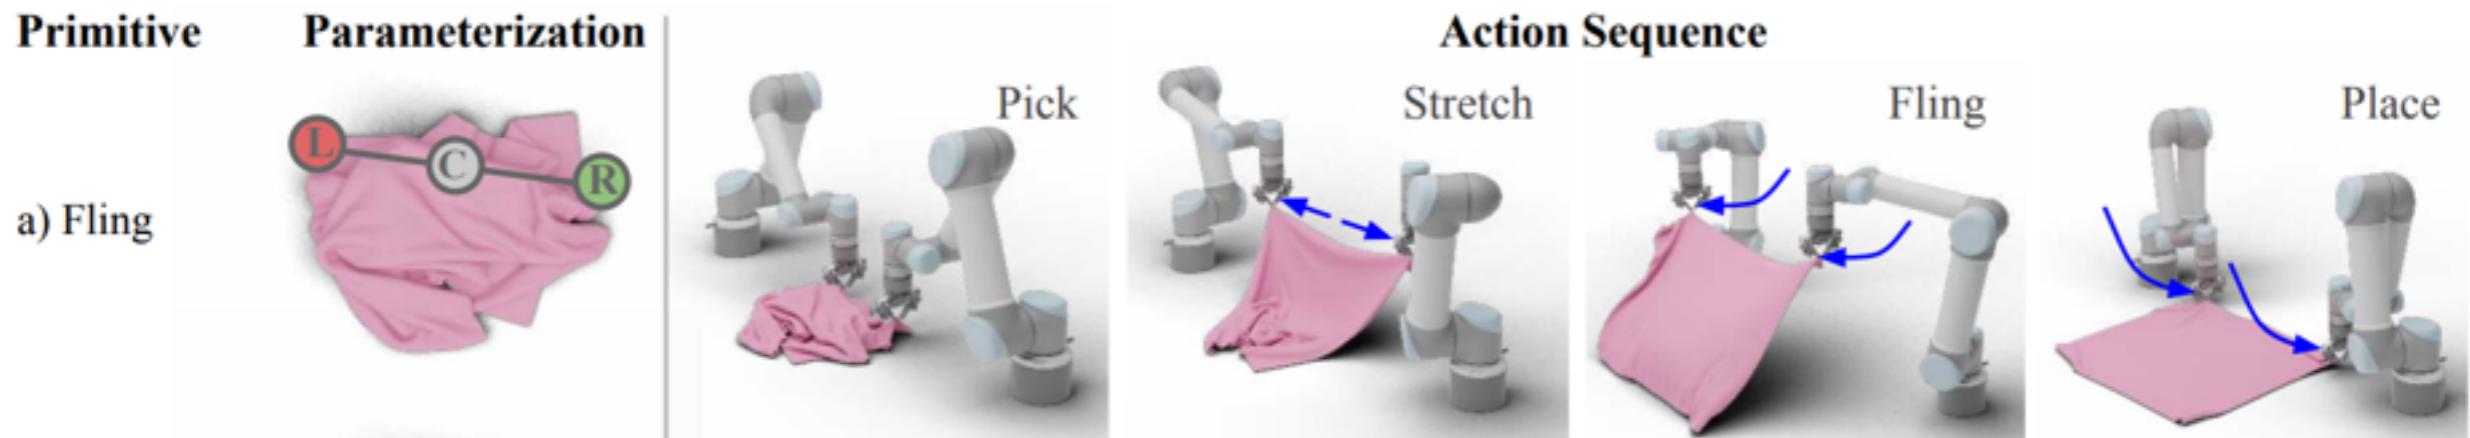
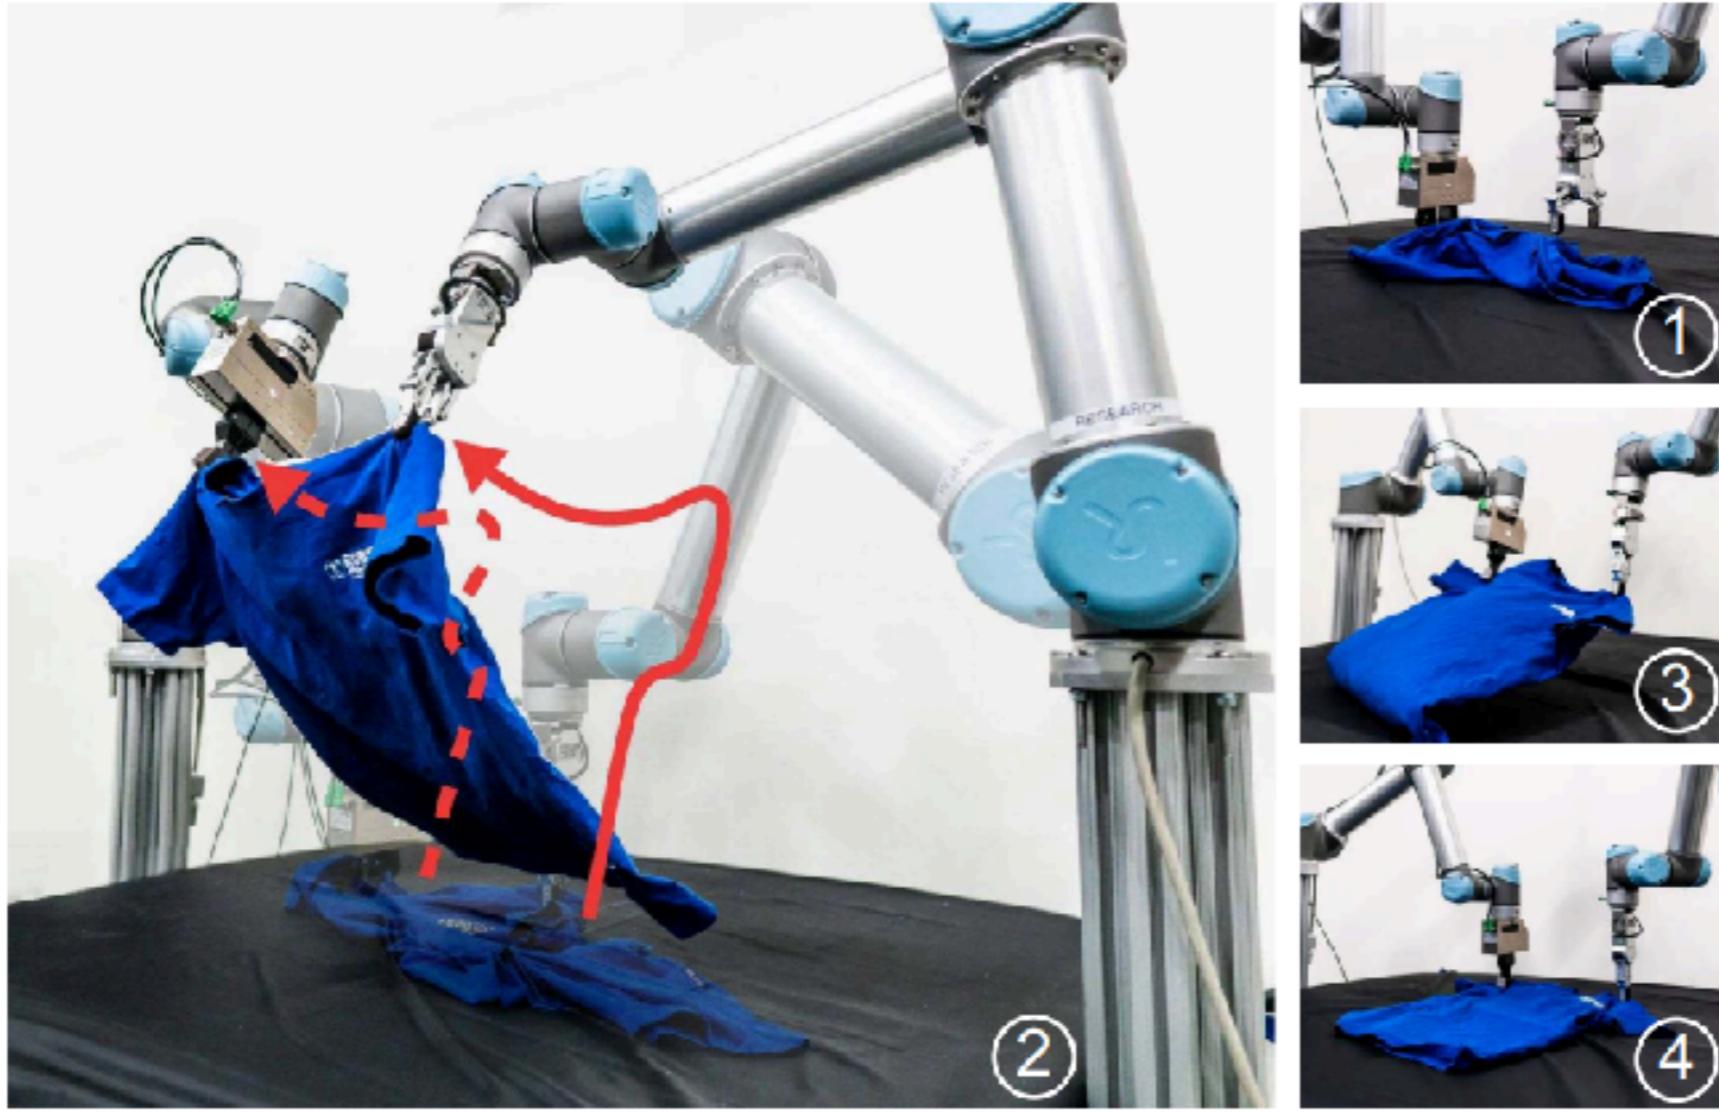


Solve a Related but Supervision-rich Problem

One-shot grasping often
leads to failed grasp attempts

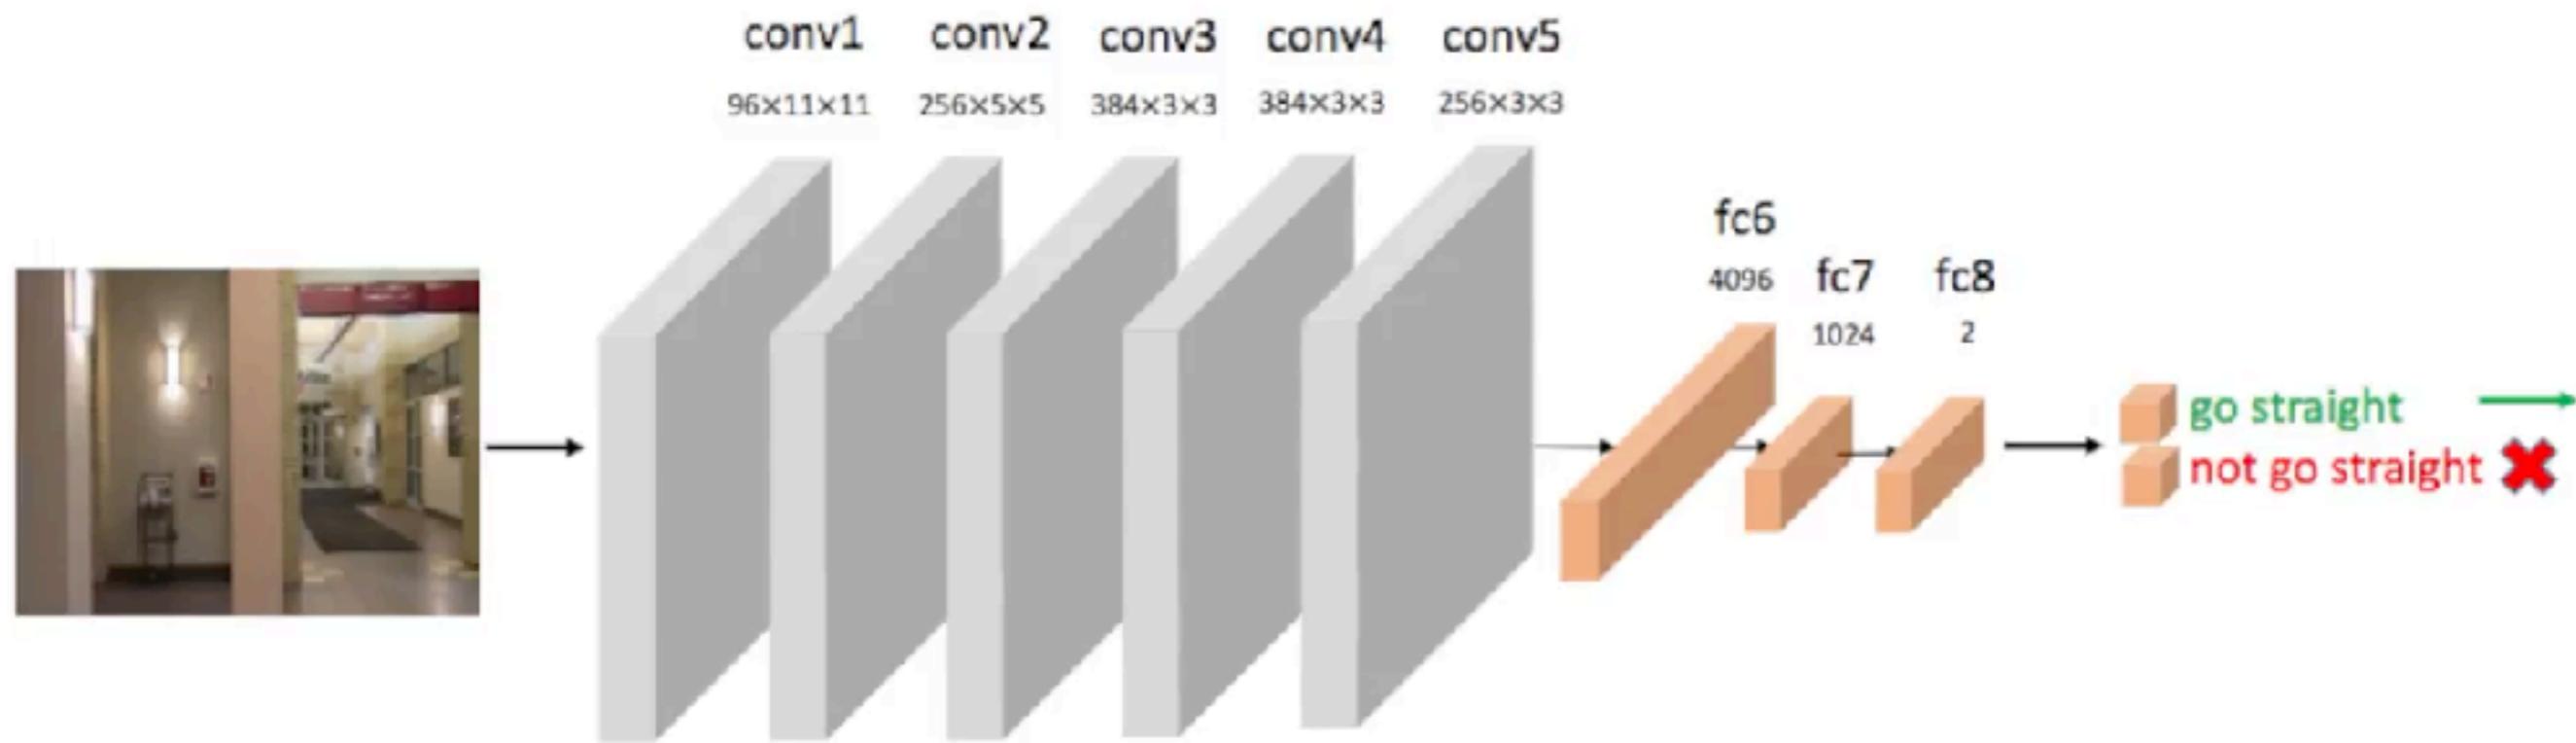


FlingBot



Solve a Related but Supervision-rich Problem

Training Network



Given the image trying to learn whether to **move forward** or **not**

Solve a Related but Supervision-rich Problem

Hallway With Chairs



Solve a Related but Supervision-rich Problem

Design setups that can allow rapid data gathering

Discussion

- A) Why may we want to use learning for grasping?
- B) What are the limitations of the techniques that we discussed today? In what scenarios can / can't we use them?

Thank you