

Sim2Real

Saurabh Gupta

Solving a RL Problem

Better Reward Signals

Sim2Real

Better Optimization

Convert into a
Supervised Training
Problem

Solve a Related but
Supervision-rich Problem

Build Models and Plan
with Them

Model-free RL
with sparse
rewards

Known reward,
known model.
Model-based RL



Sim2Real (Discussion)

Train policy in simulation, and then deploy in the real world

Let's assume we want to learn policies,

- A: what are some of the advantages of sim2real?
- B: what are some of the short-comings of sim2real?

Getting Sim2Real to Work

Minimizing the domain gap

- sensing
 - data augmentation (aka domain randomization)
 - building rich simulators
 - use modalities that are easier to simulate
 - domain adaptation
- actuation
 - operating at the appropriate level of abstraction

Building Rich Simulators

Simulator based on scans of
Real World Environments

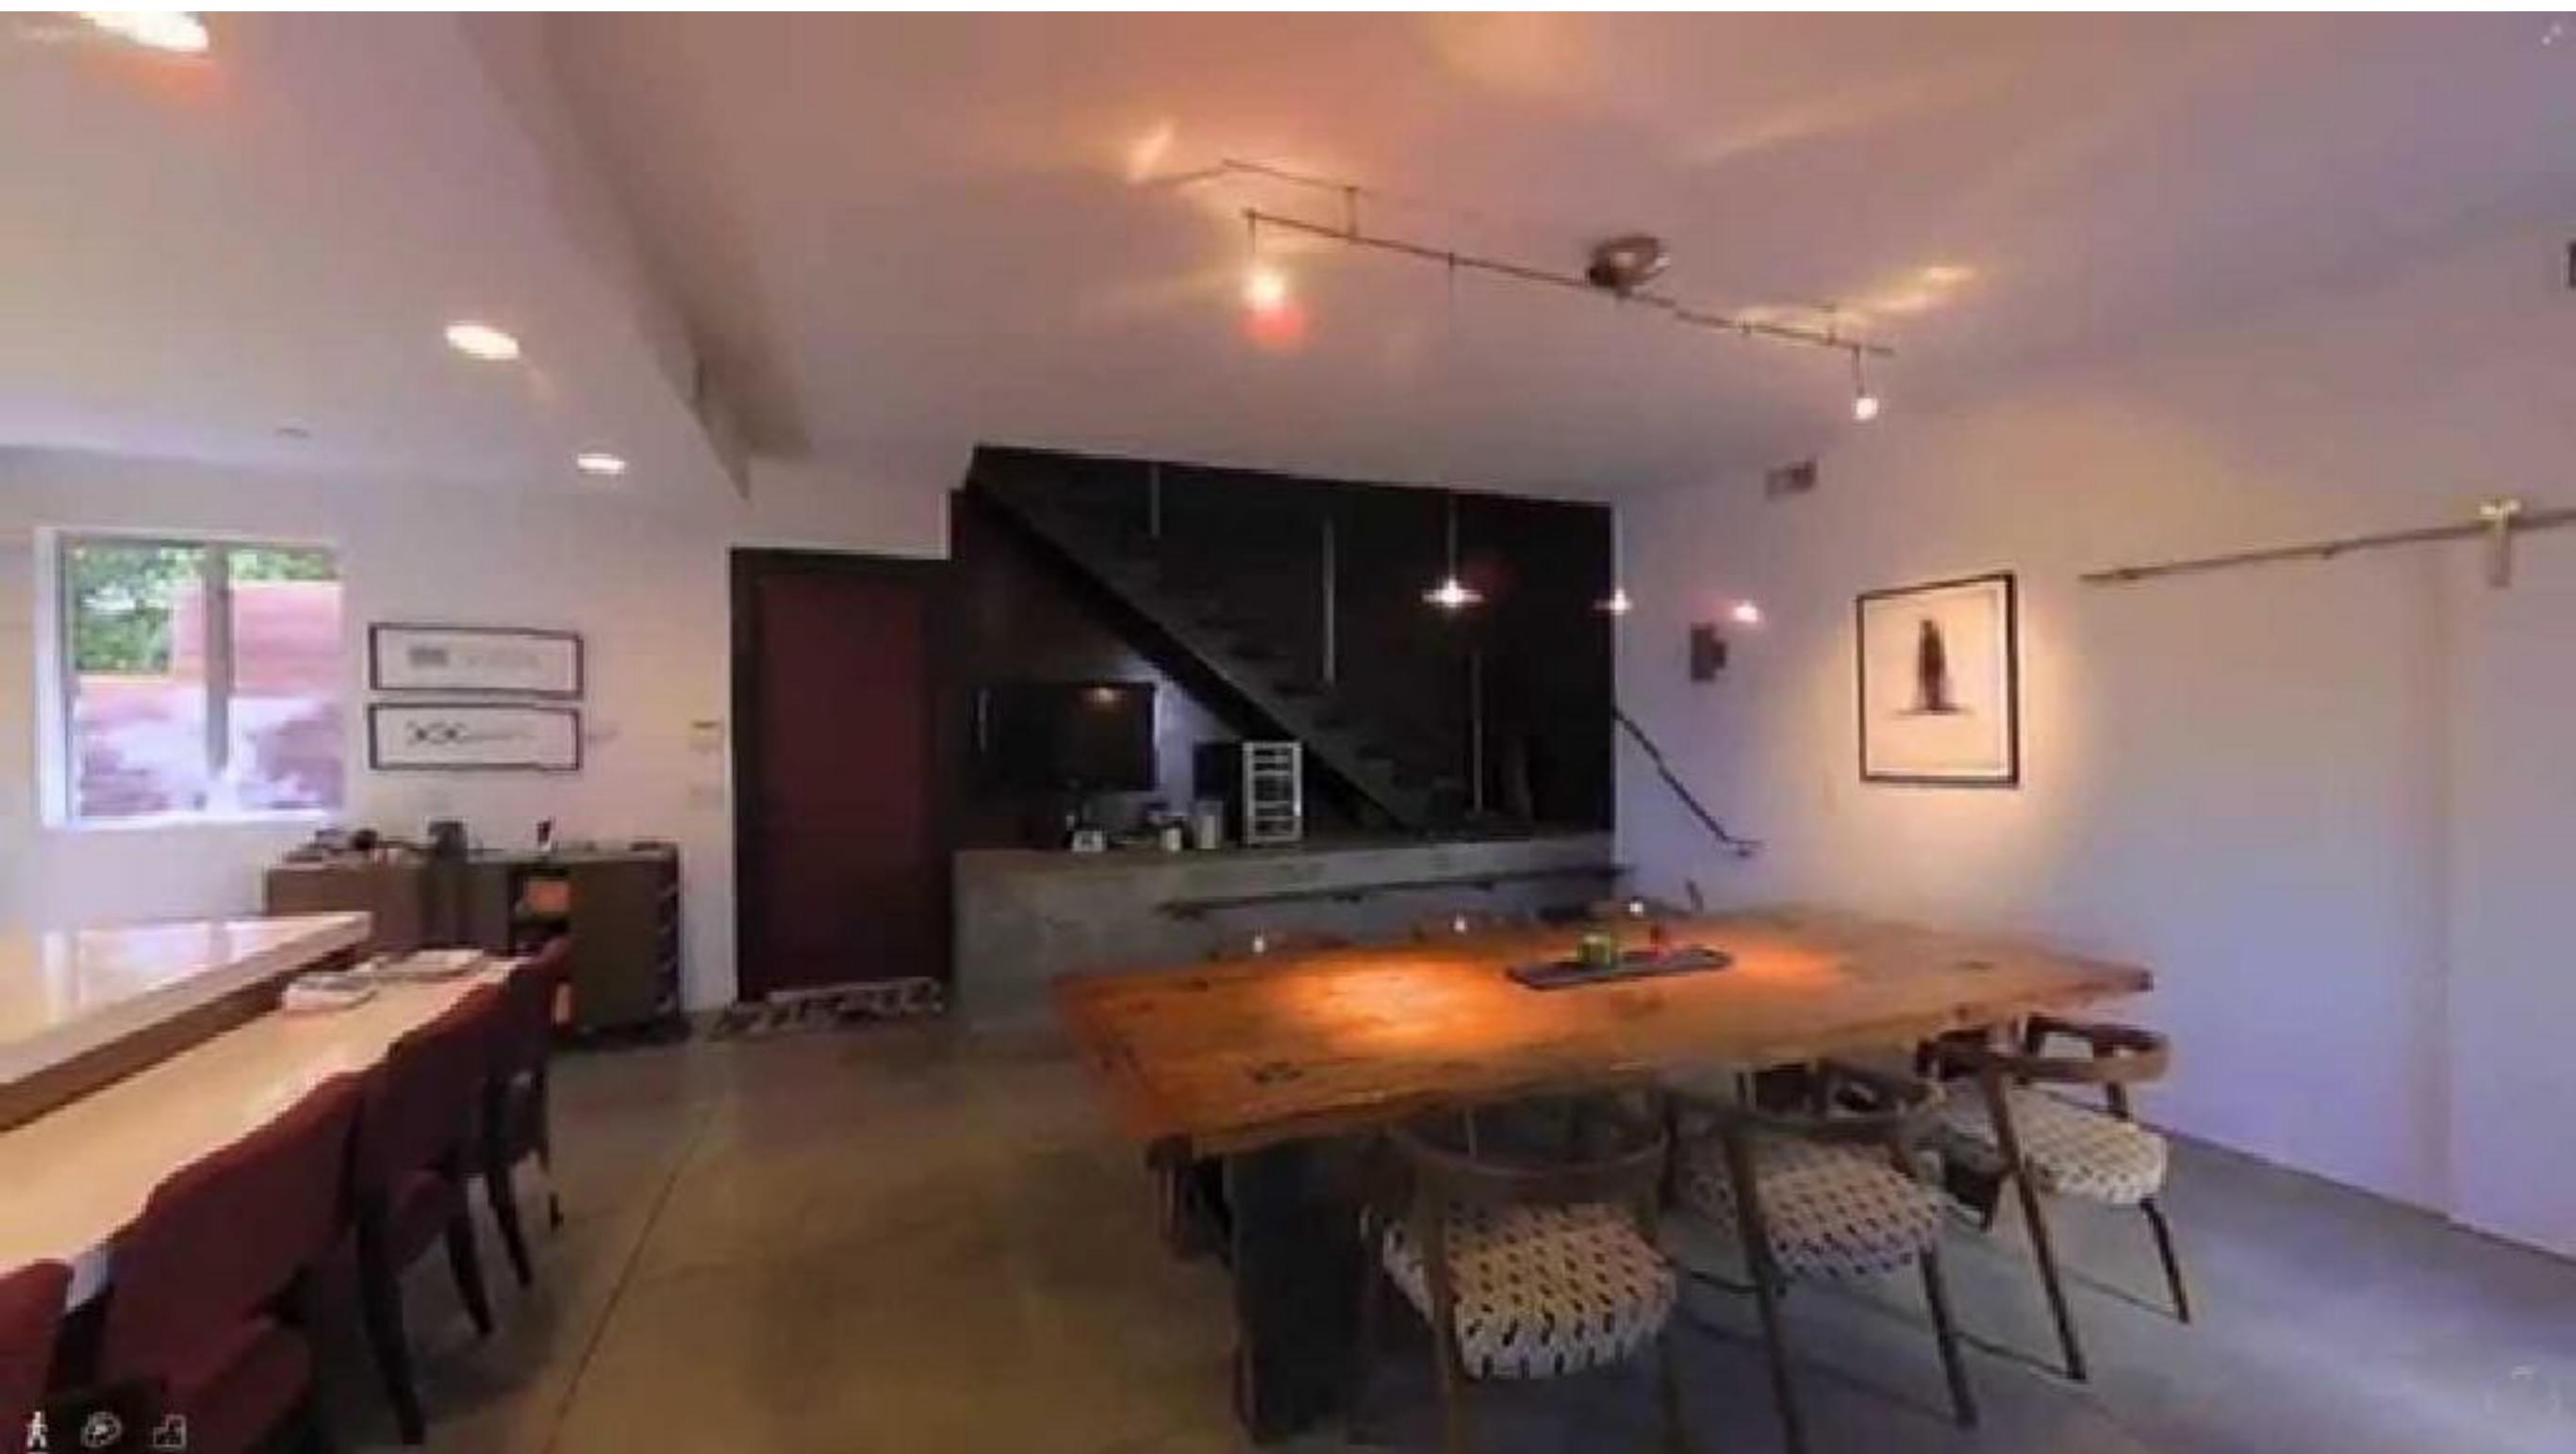
Simulate robot views
and motion



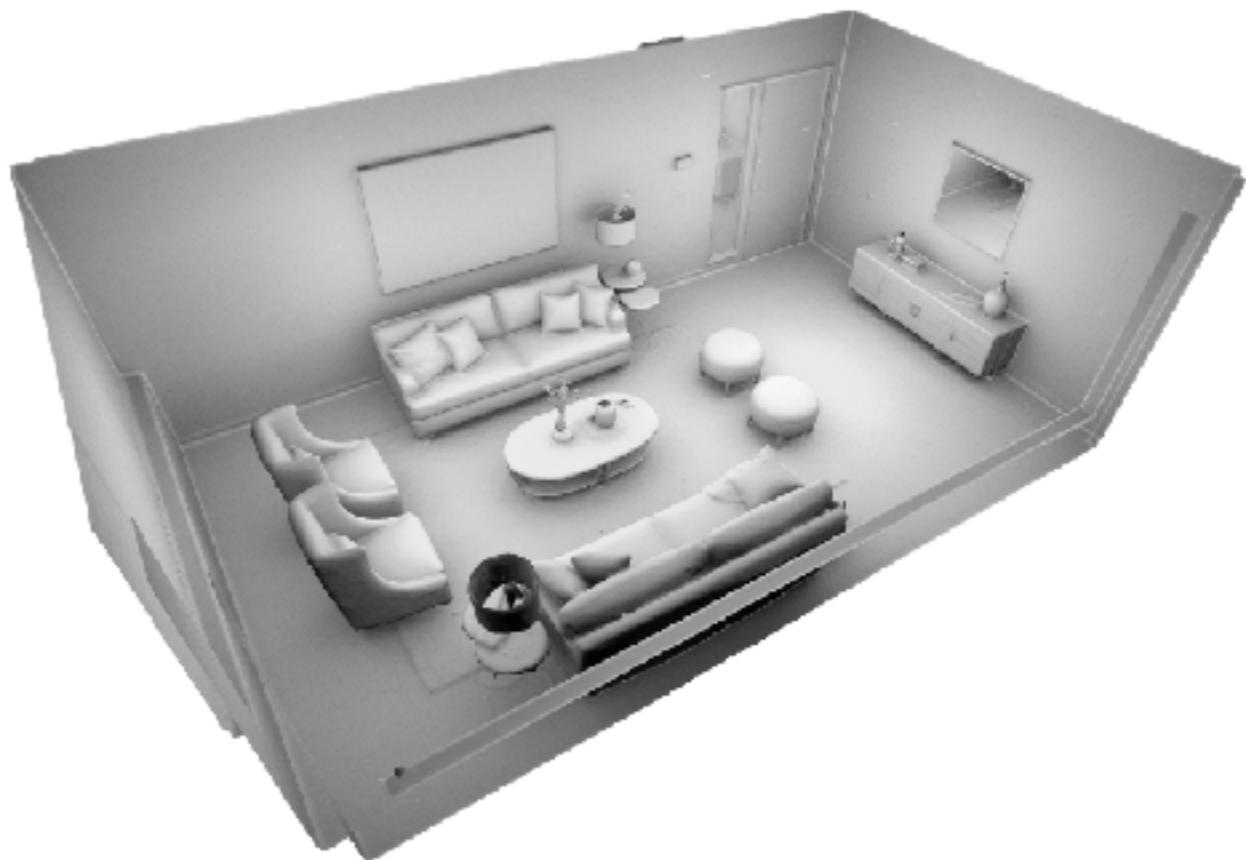
Compute ground
truth traversability



Building Rich Simulators



Building Rich Simulators



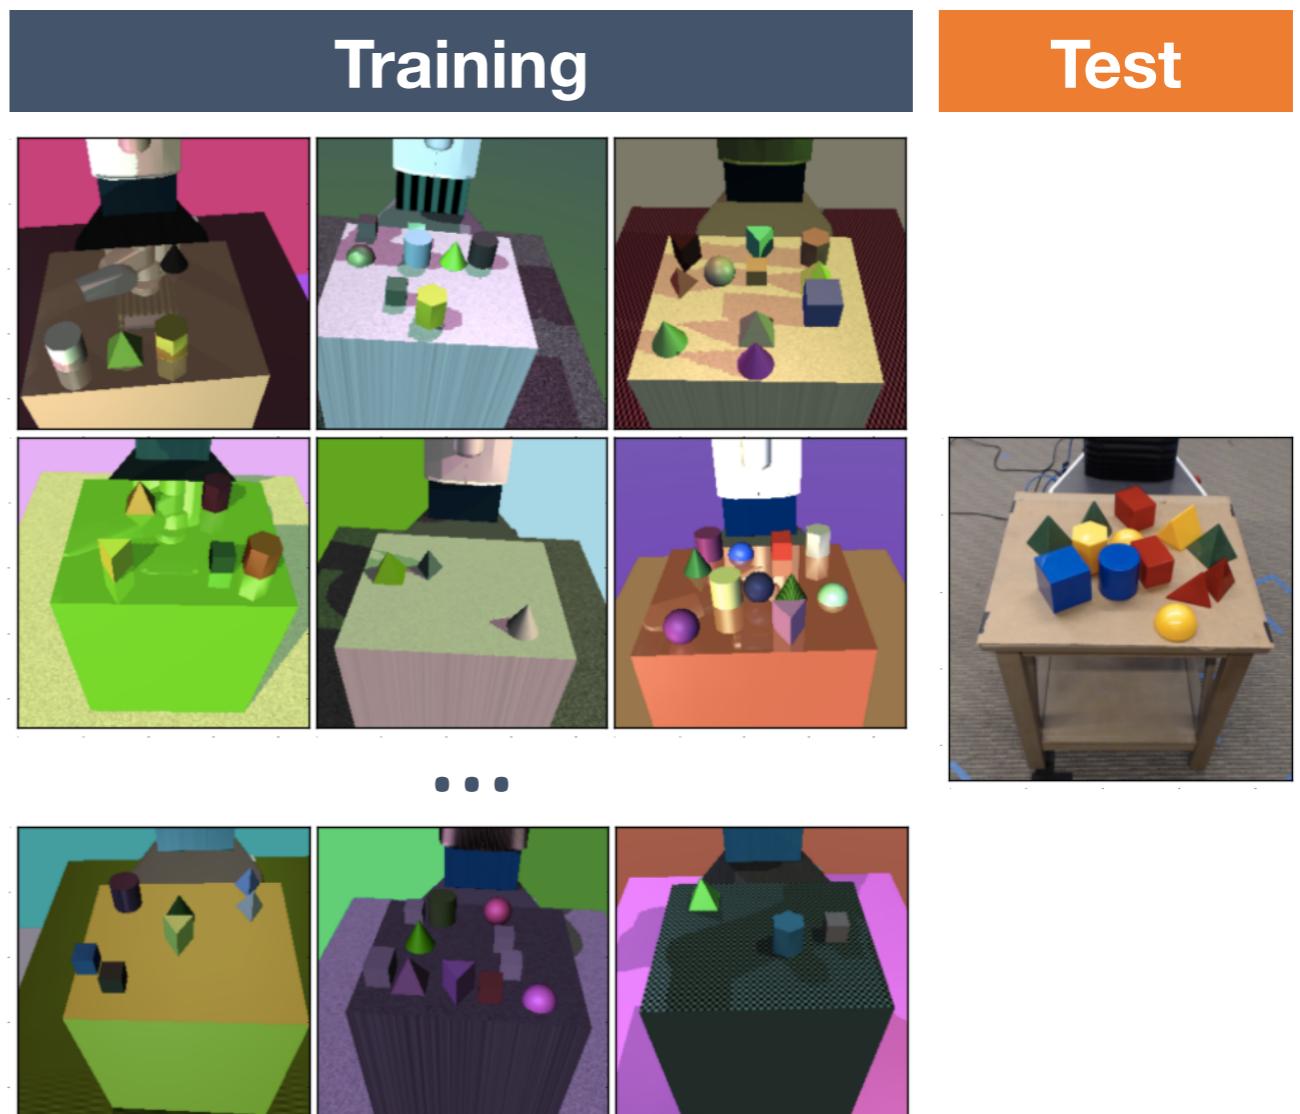
The Replica Dataset: A Digital Replica of Indoor Spaces

Data Augmentation

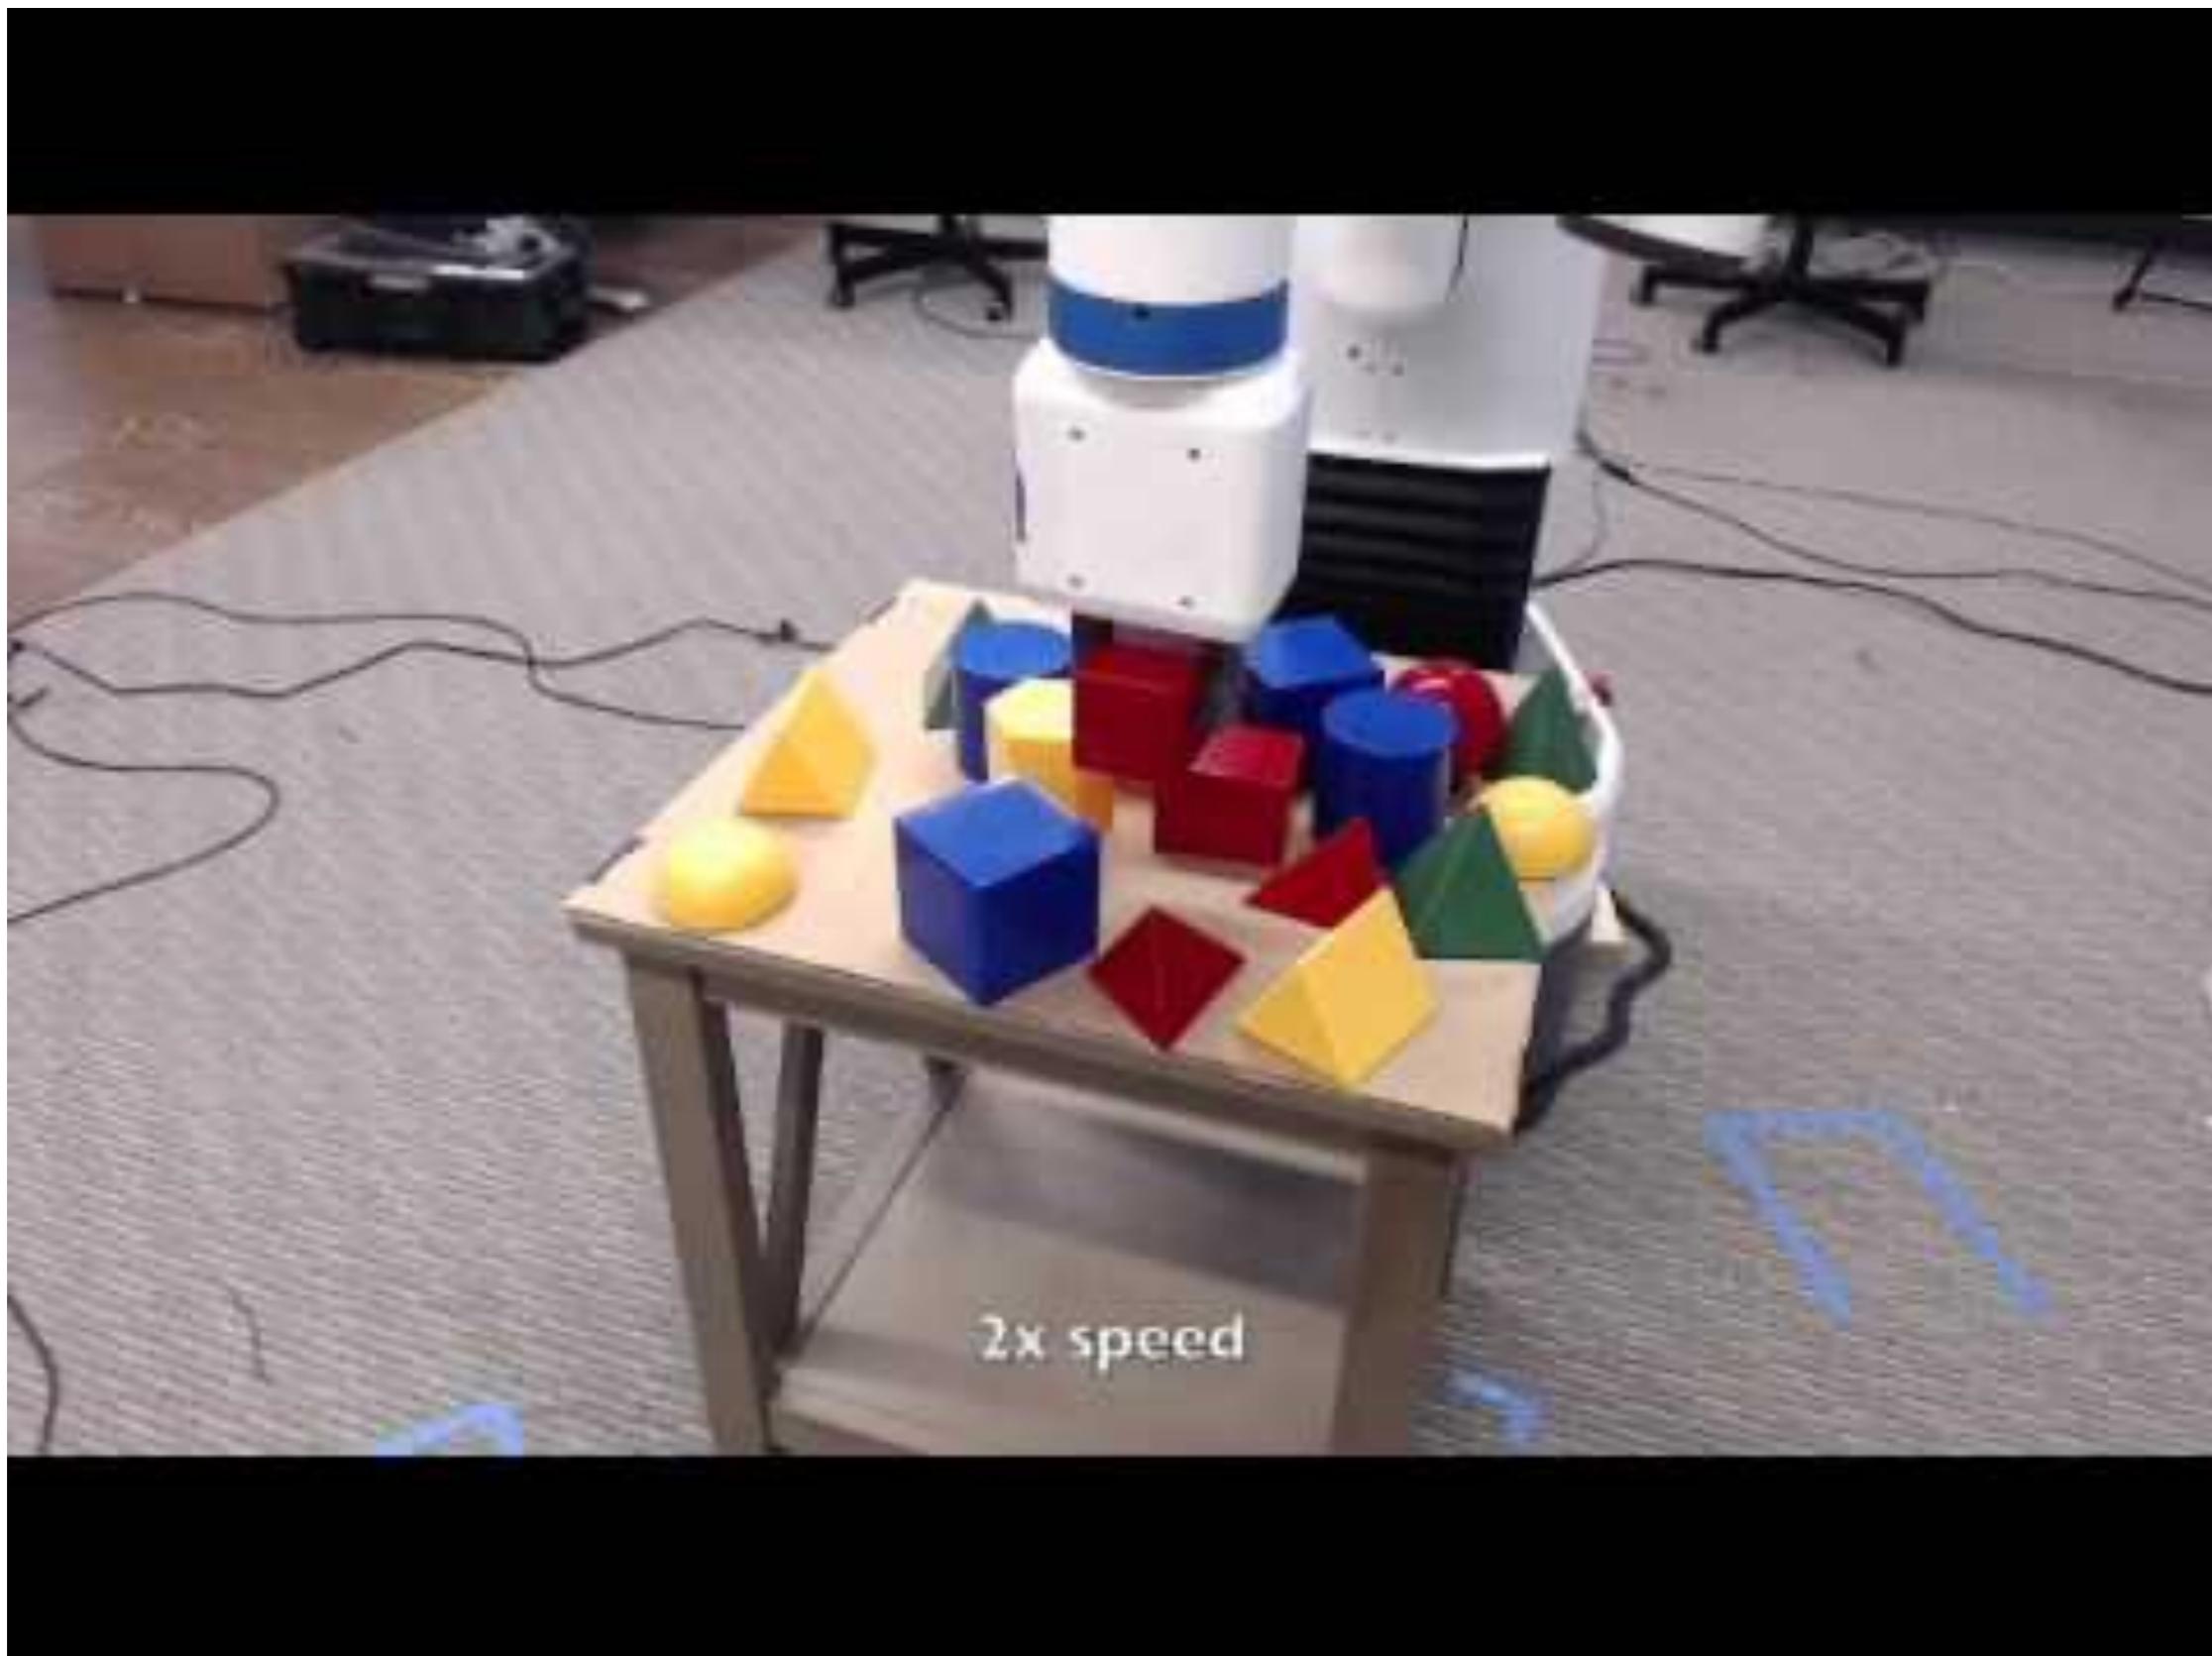
Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World

Josh Tobin¹, Rachel Fong², Alex Ray², Jonas Schneider², Wojciech Zaremba², Pieter Abbeel³

Demonstrate that with enough variations during training, models trained in simulation can transfer to real world for estimating object locations.



Data Augmentation



Data Augmentation

CAD²RL: Real Single-Image Flight Without a Single Real Image

Fereshteh Sadeghi

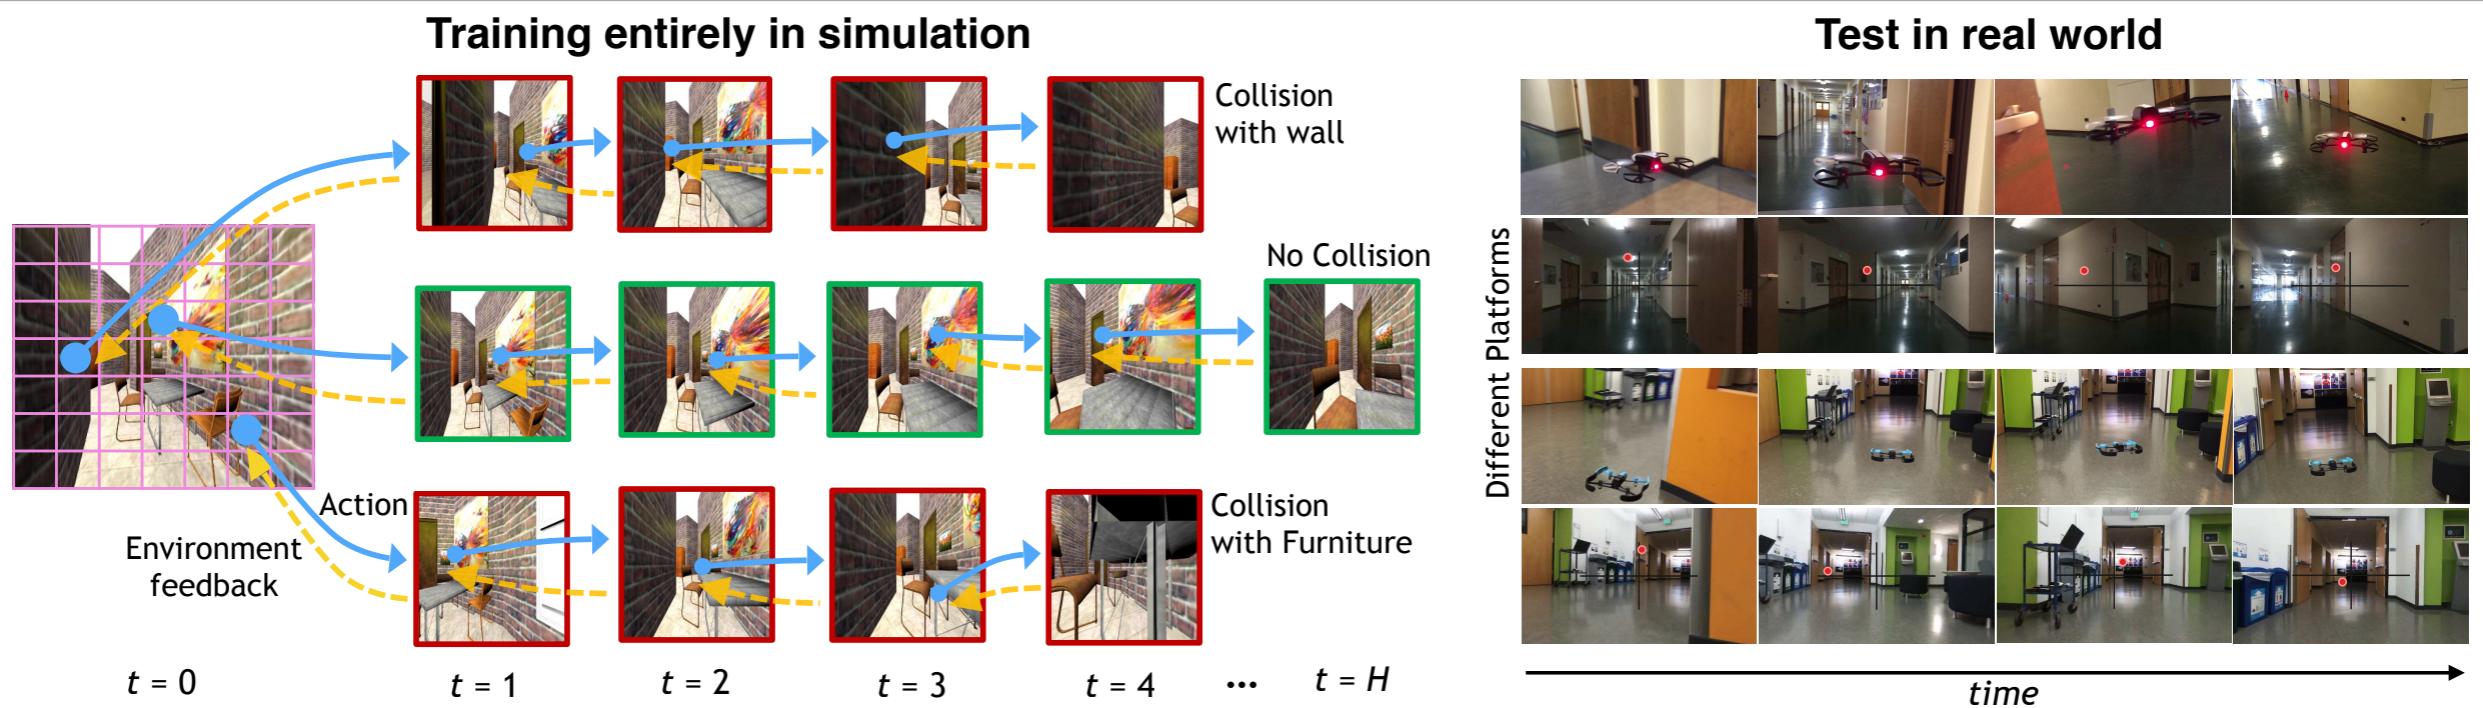
University of Washington

fsadeghi@cs.washington.edu

Sergey Levine

University of California, Berkeley

svlevine@eecs.berkeley.edu

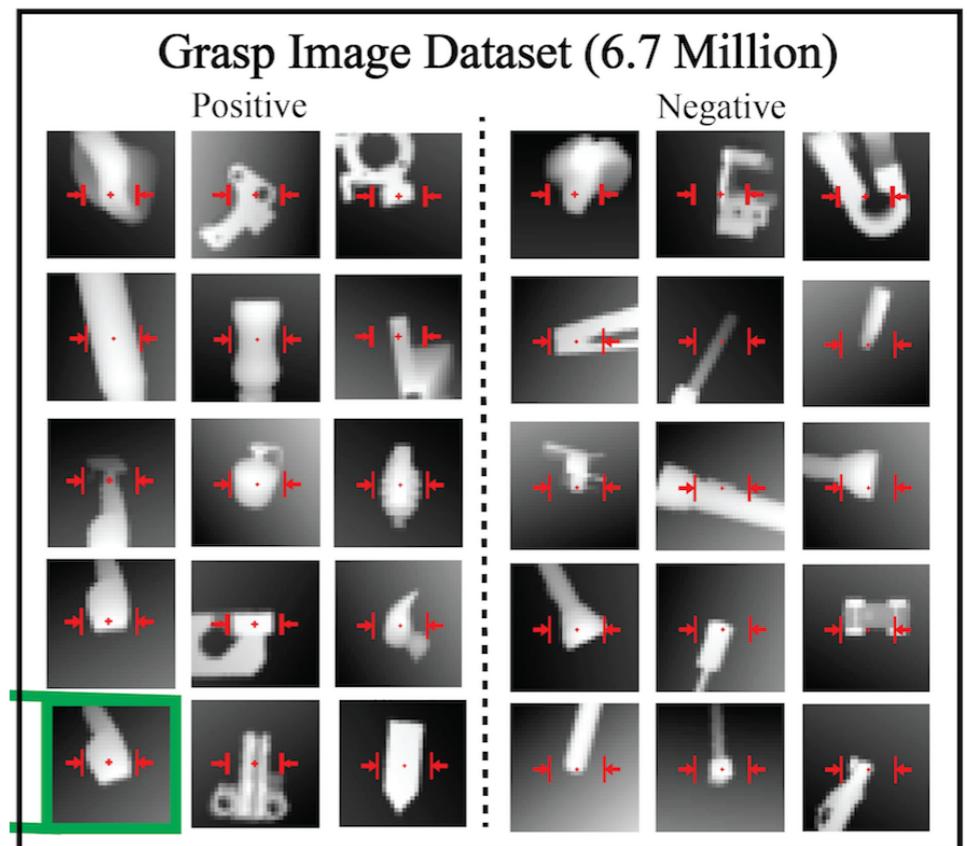
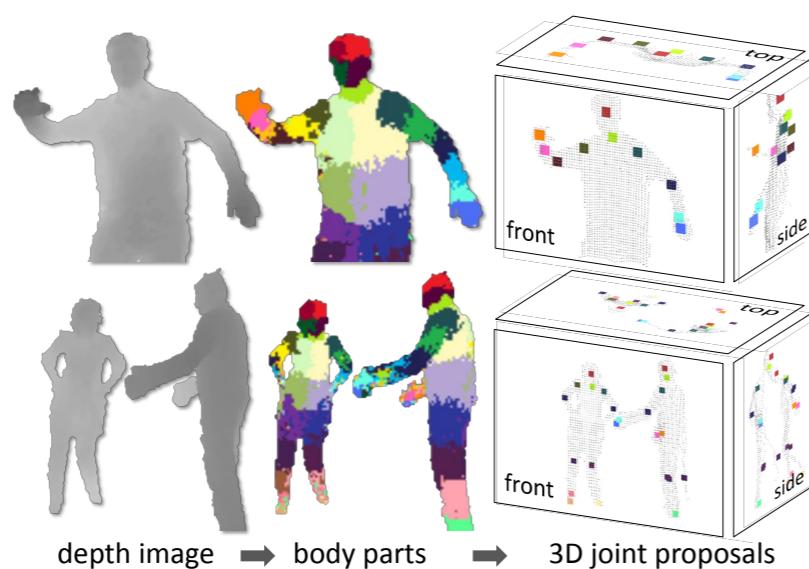


Data Augmentation

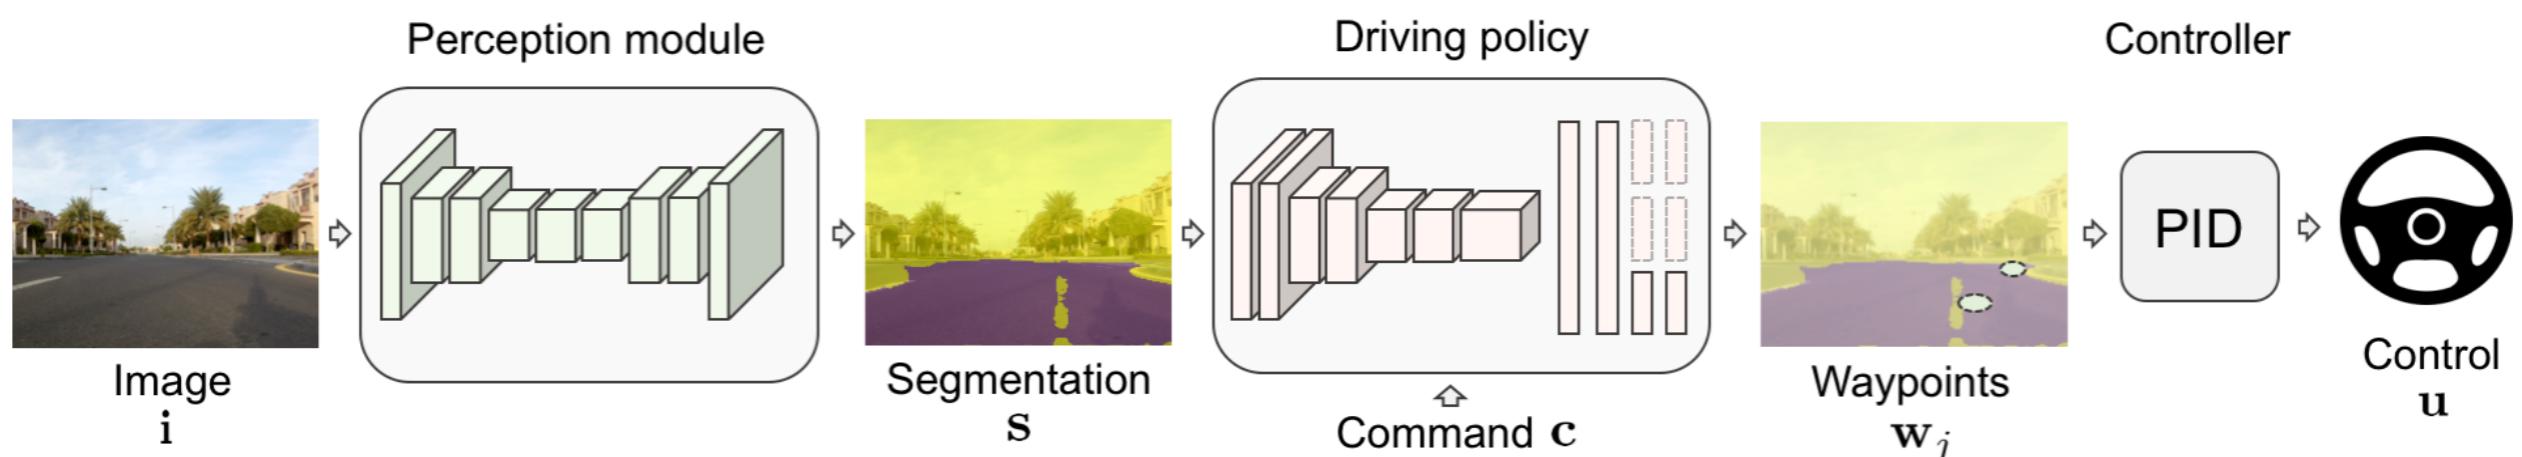


Modalities that have smaller domain gap

eg. depth



eg. semantic segmentation masks



Real-Time Human Pose Recognition in Parts from Single Depth Images. Shotton et al. CVPR 2011

DexNet 2.0. Mahler et al. 2017

Driving Policy Transfer via Modularity and Abstraction. Müller et al. CoRL 2018

Domain Adaptation

ML techniques to learn representations that are invariant to domain gap



Source image (GTA5)



Target image (CityScapes)

Pixel accuracy on target
Source-only: 54.0%
Adapted (ours): 83.6%



Source images (SVHN)

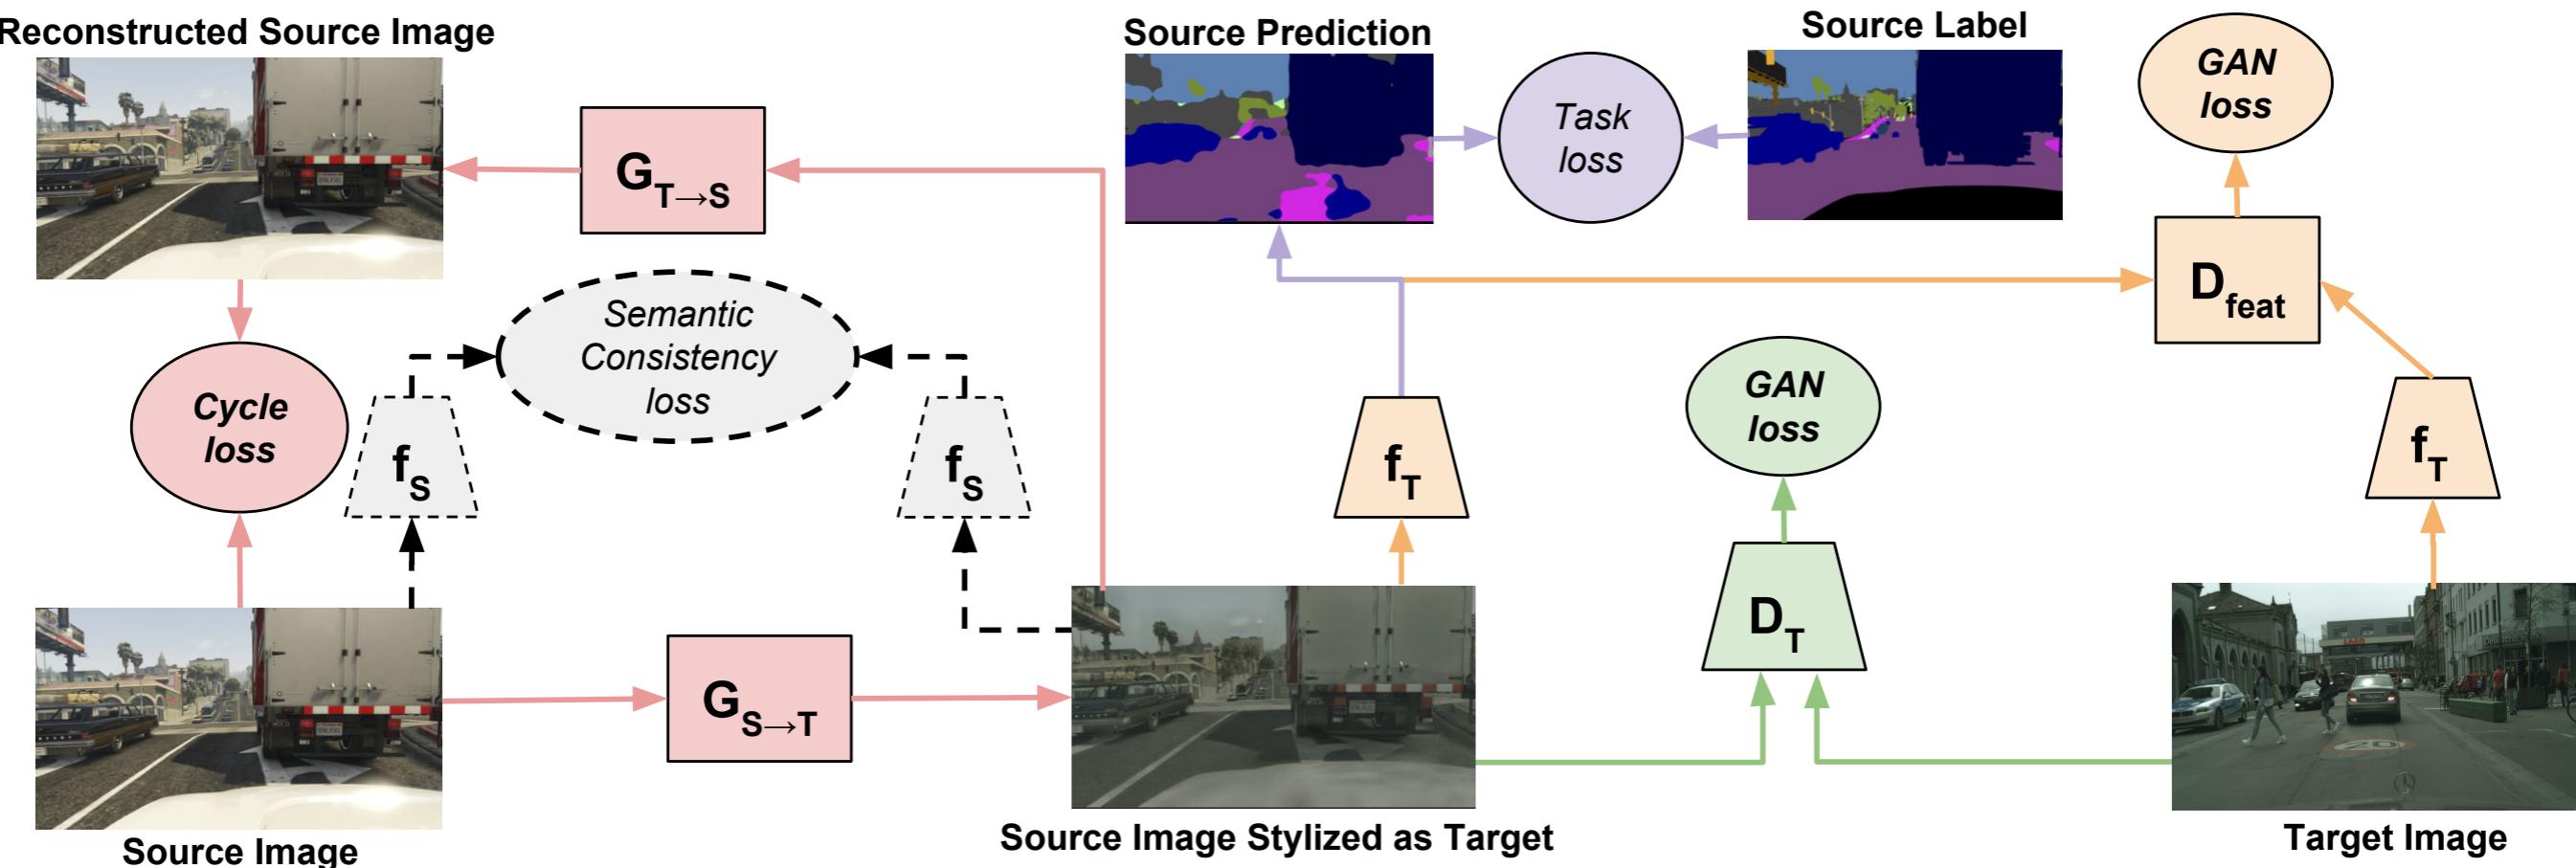


Target images (MNIST)

Accuracy on target
Source-only: 67.1%
Adapted (ours): 90.4%

Domain Adaptation

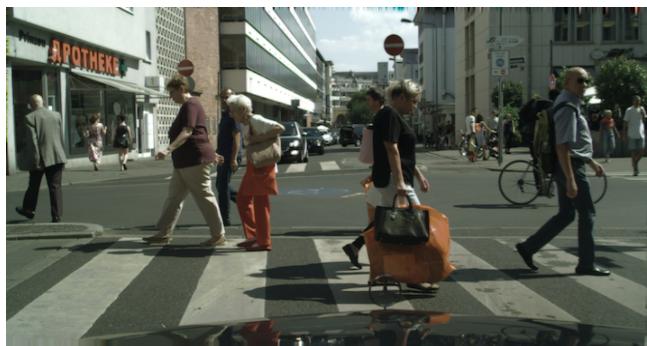
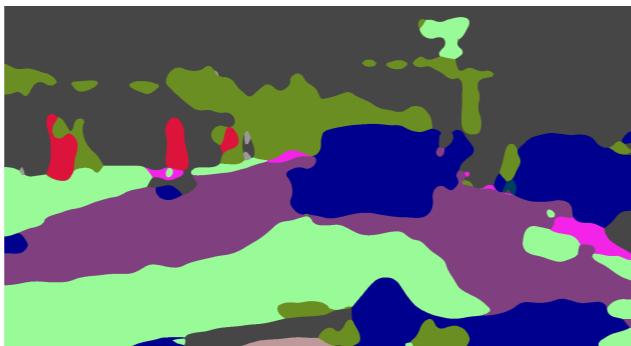
ML techniques to learn representations that are invariant to domain gap



Adversarial discriminative domain adaptation ETzeng, J Hoffman, K Saenko, T Darrell. ICCV 2017.
Cycada: Cycle-consistent adversarial domain adaptation. J Hoffman, et al. ICML 2018.

Domain Adaptation

ML techniques to learn representations that are invariant to domain gap



(a) Test Image

(b) Source Prediction

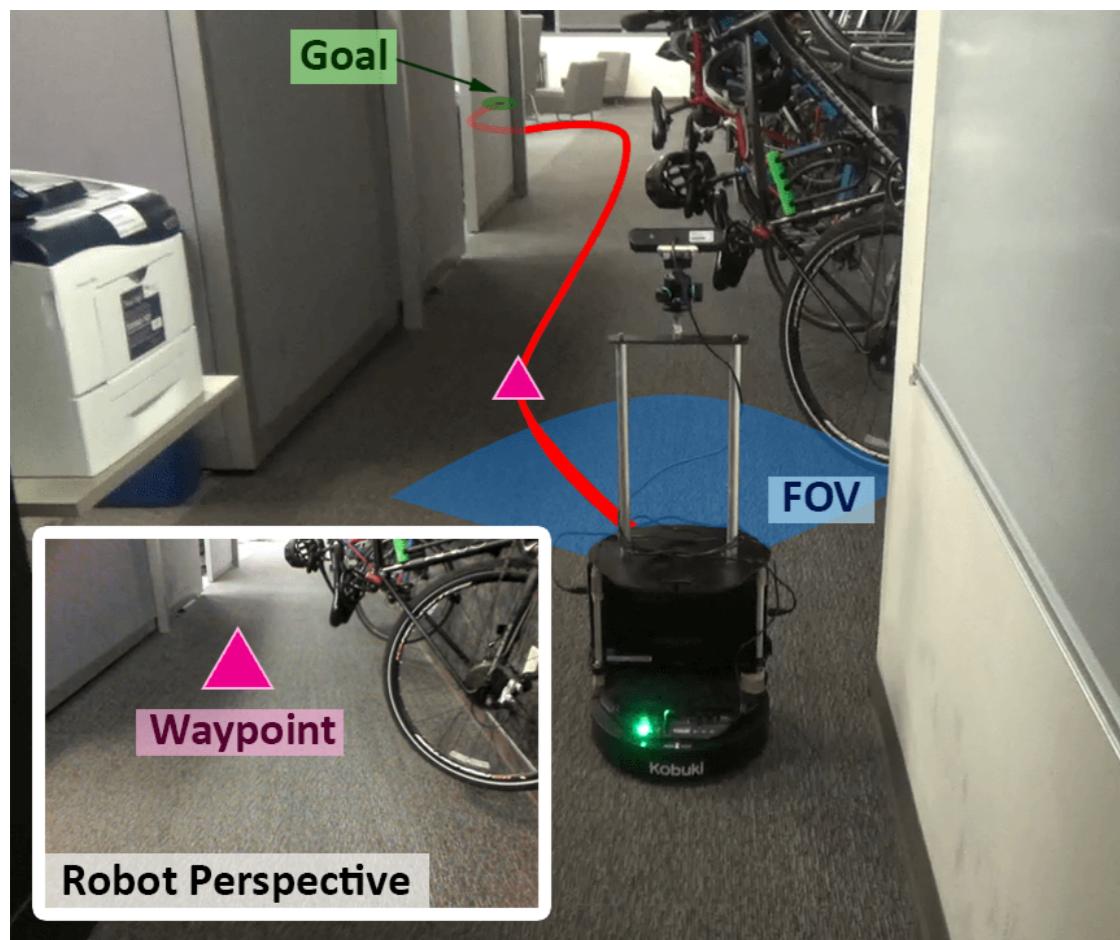
(c) CyCADA Prediction

(d) Ground Truth

Operating at the right abstraction level

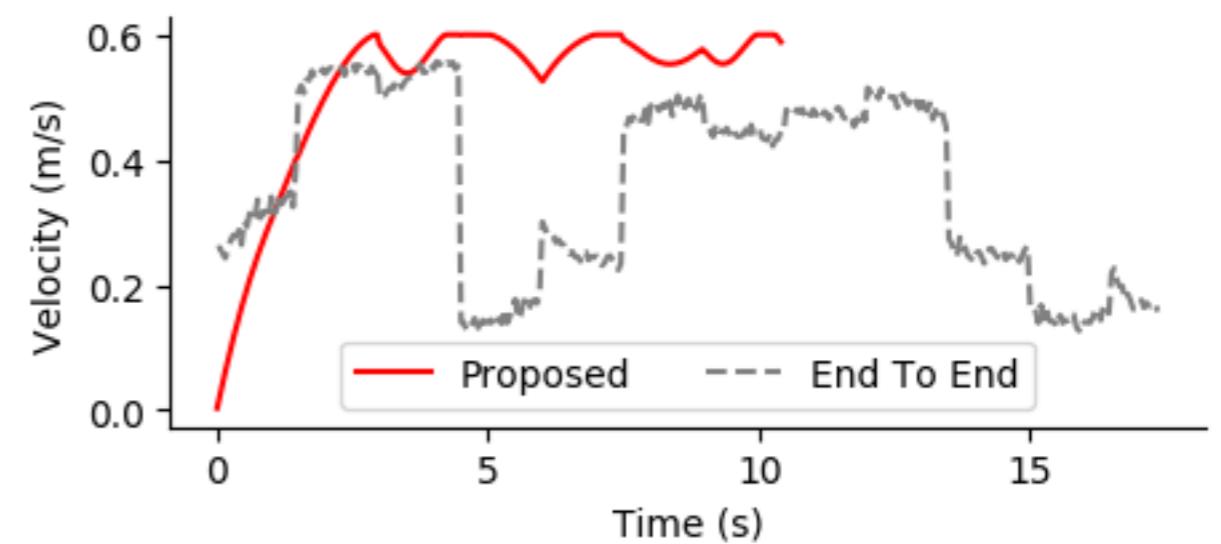
Combining Optimal Control and Learning for Visual Navigation in Novel Environments

Somil Bansal^{*1} Varun Tolani^{*1} Saurabh Gupta² Jitendra Malik^{1,2} Claire Tomlin¹



Output:

- a) way-points that are tracked using a LQR controller
- b) angular and linear velocities



Operating at the right abstraction level

LEARNING LOCOMOTION SKILLS USING DEEPRL: DOES THE CHOICE OF ACTION SPACE MATTER?

Xue Bin Peng & Michiel van de Panne

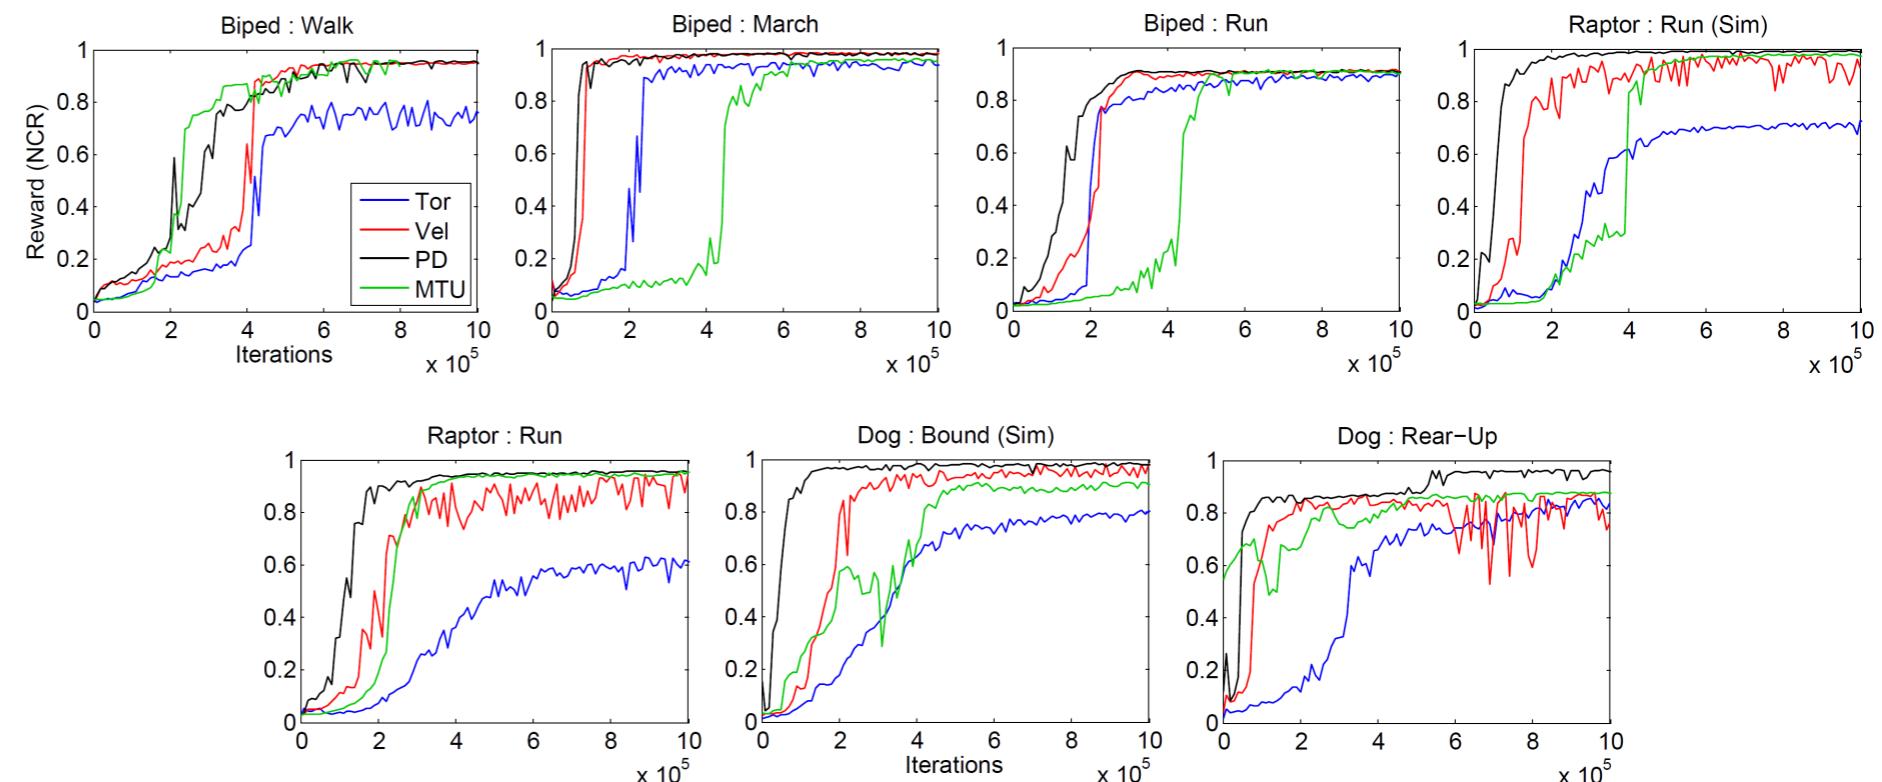
Department of Computer Science
University of the British Columbia
Vancouver, Canada
`{xbpeng, van}@cs.ubc.ca`

Policy action space:

- a) torques
- b) target joint velocities
- c) target joint angles
- d) muscle activations

$$\tau^i = k_d^i(\hat{q}^i - \dot{q}^i)$$

$$k_p^i(\hat{q}^i - q^i) + k_d^i(\hat{\dot{q}}^i - \dot{q}^i)$$



Getting Sim2Real to Work

Minimizing the domain gap

- sensing
 - data augmentation (aka domain randomization)
 - building rich simulators
 - use modalities that are easier to simulate
 - domain adaptation
- actuation
 - operating at the appropriate level of abstraction

Learning Agile and Dynamic Motor Skills for Legged Robots

JEMIN HWANGBO^{1*}, JOONHO LEE¹, ALEXEY DOSOVITSKIY², DARIO BELLICOSO¹, JOONHO LEE¹, VASSILIOS TSOUNIS¹, VLADLEN KOLTUN², AND MARCO HUTTER¹

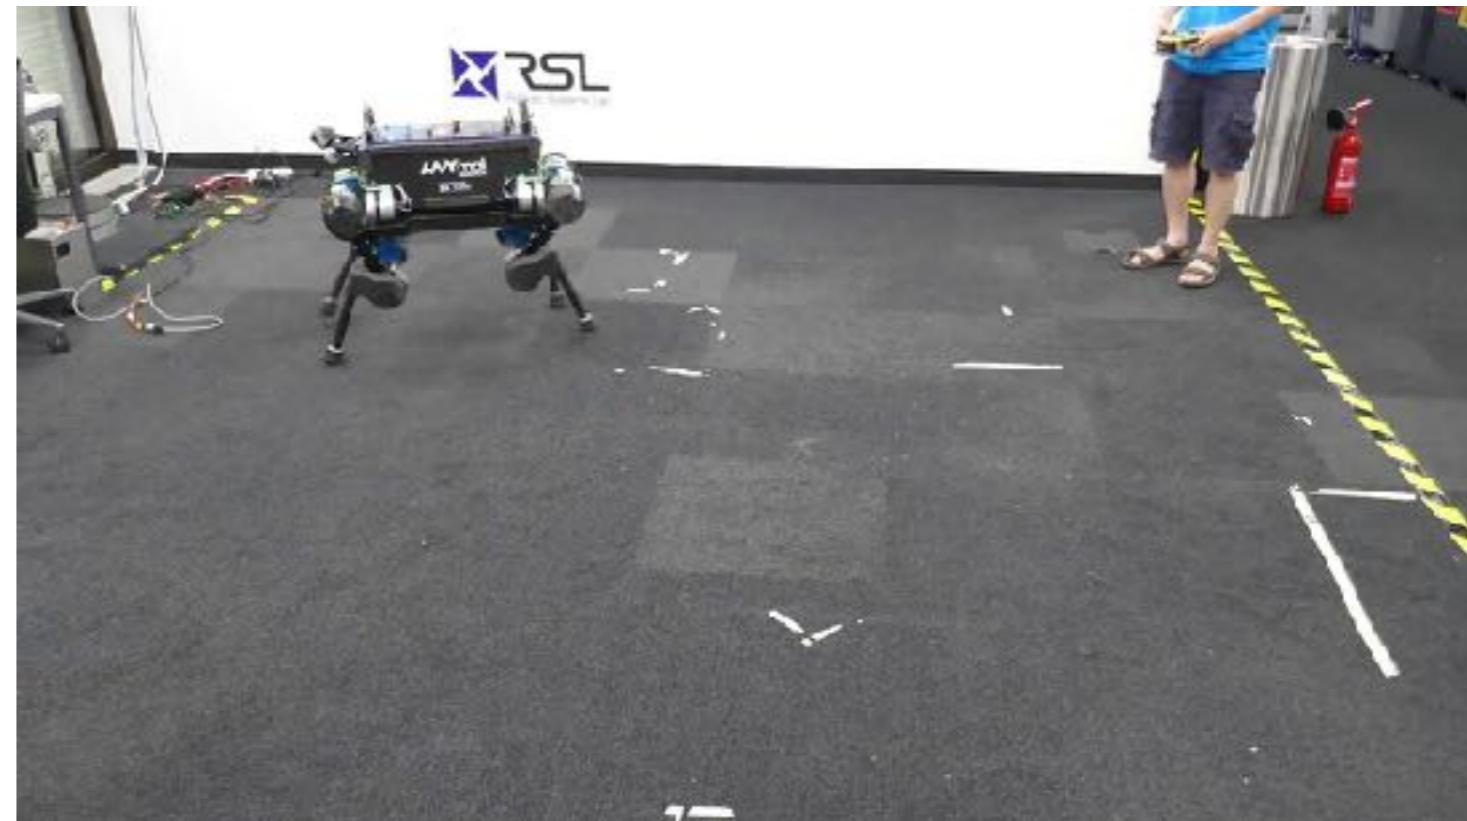
¹*Robotic Systems Lab, ETH Zurich*

²*Intelligent Systems Lab, Intel*

*Corresponding author: jemin.hwangbo@gmail.com

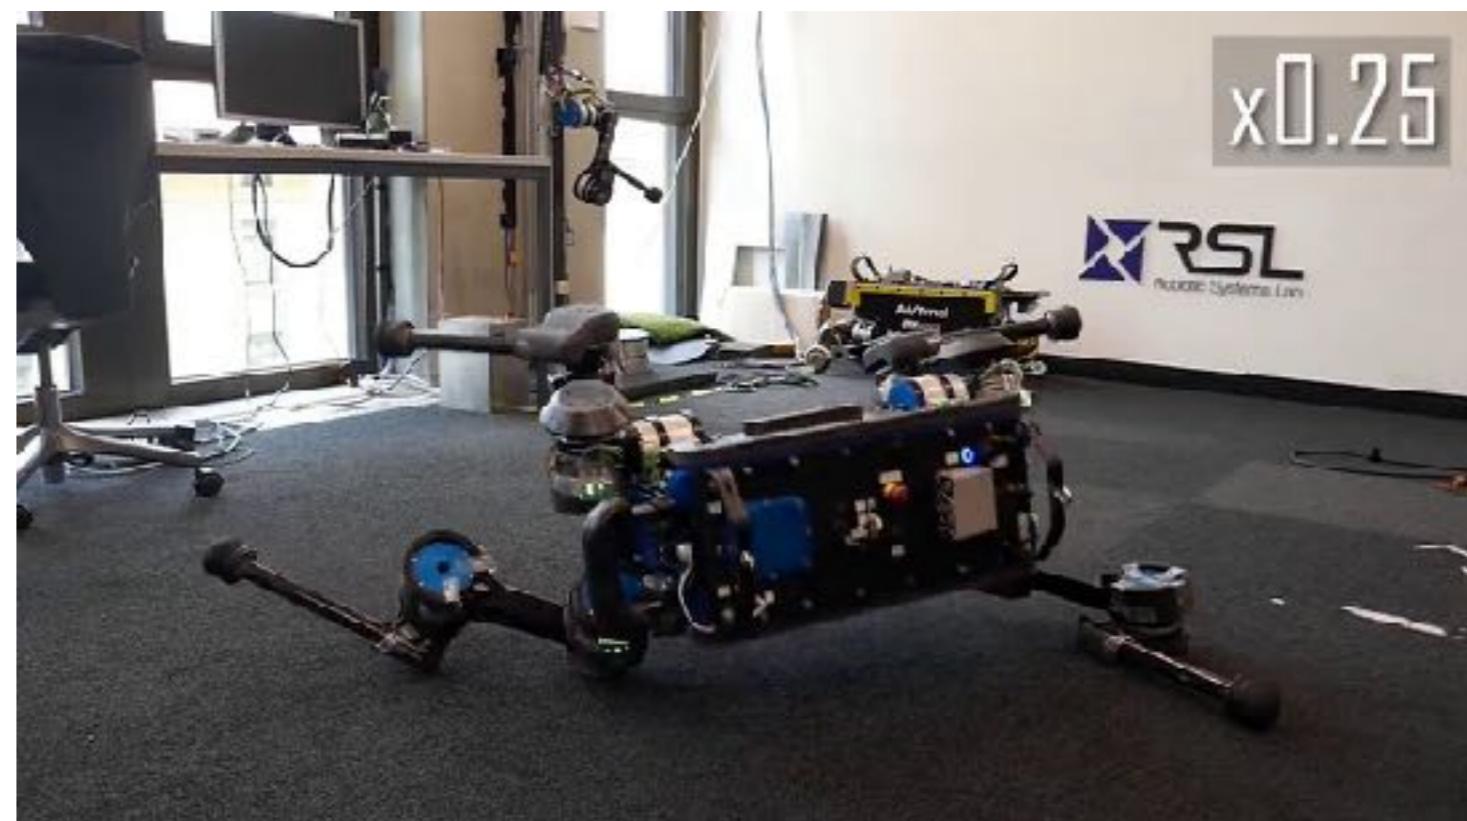
- Successful demonstration of sim2real for quadruped locomotion
- RL inside a learned simulation for successful transfer
- Case study of how to get sim2real to work in realistic scenarios
- Combining classical physics contact dynamics with learned models for actuation

ANYmal

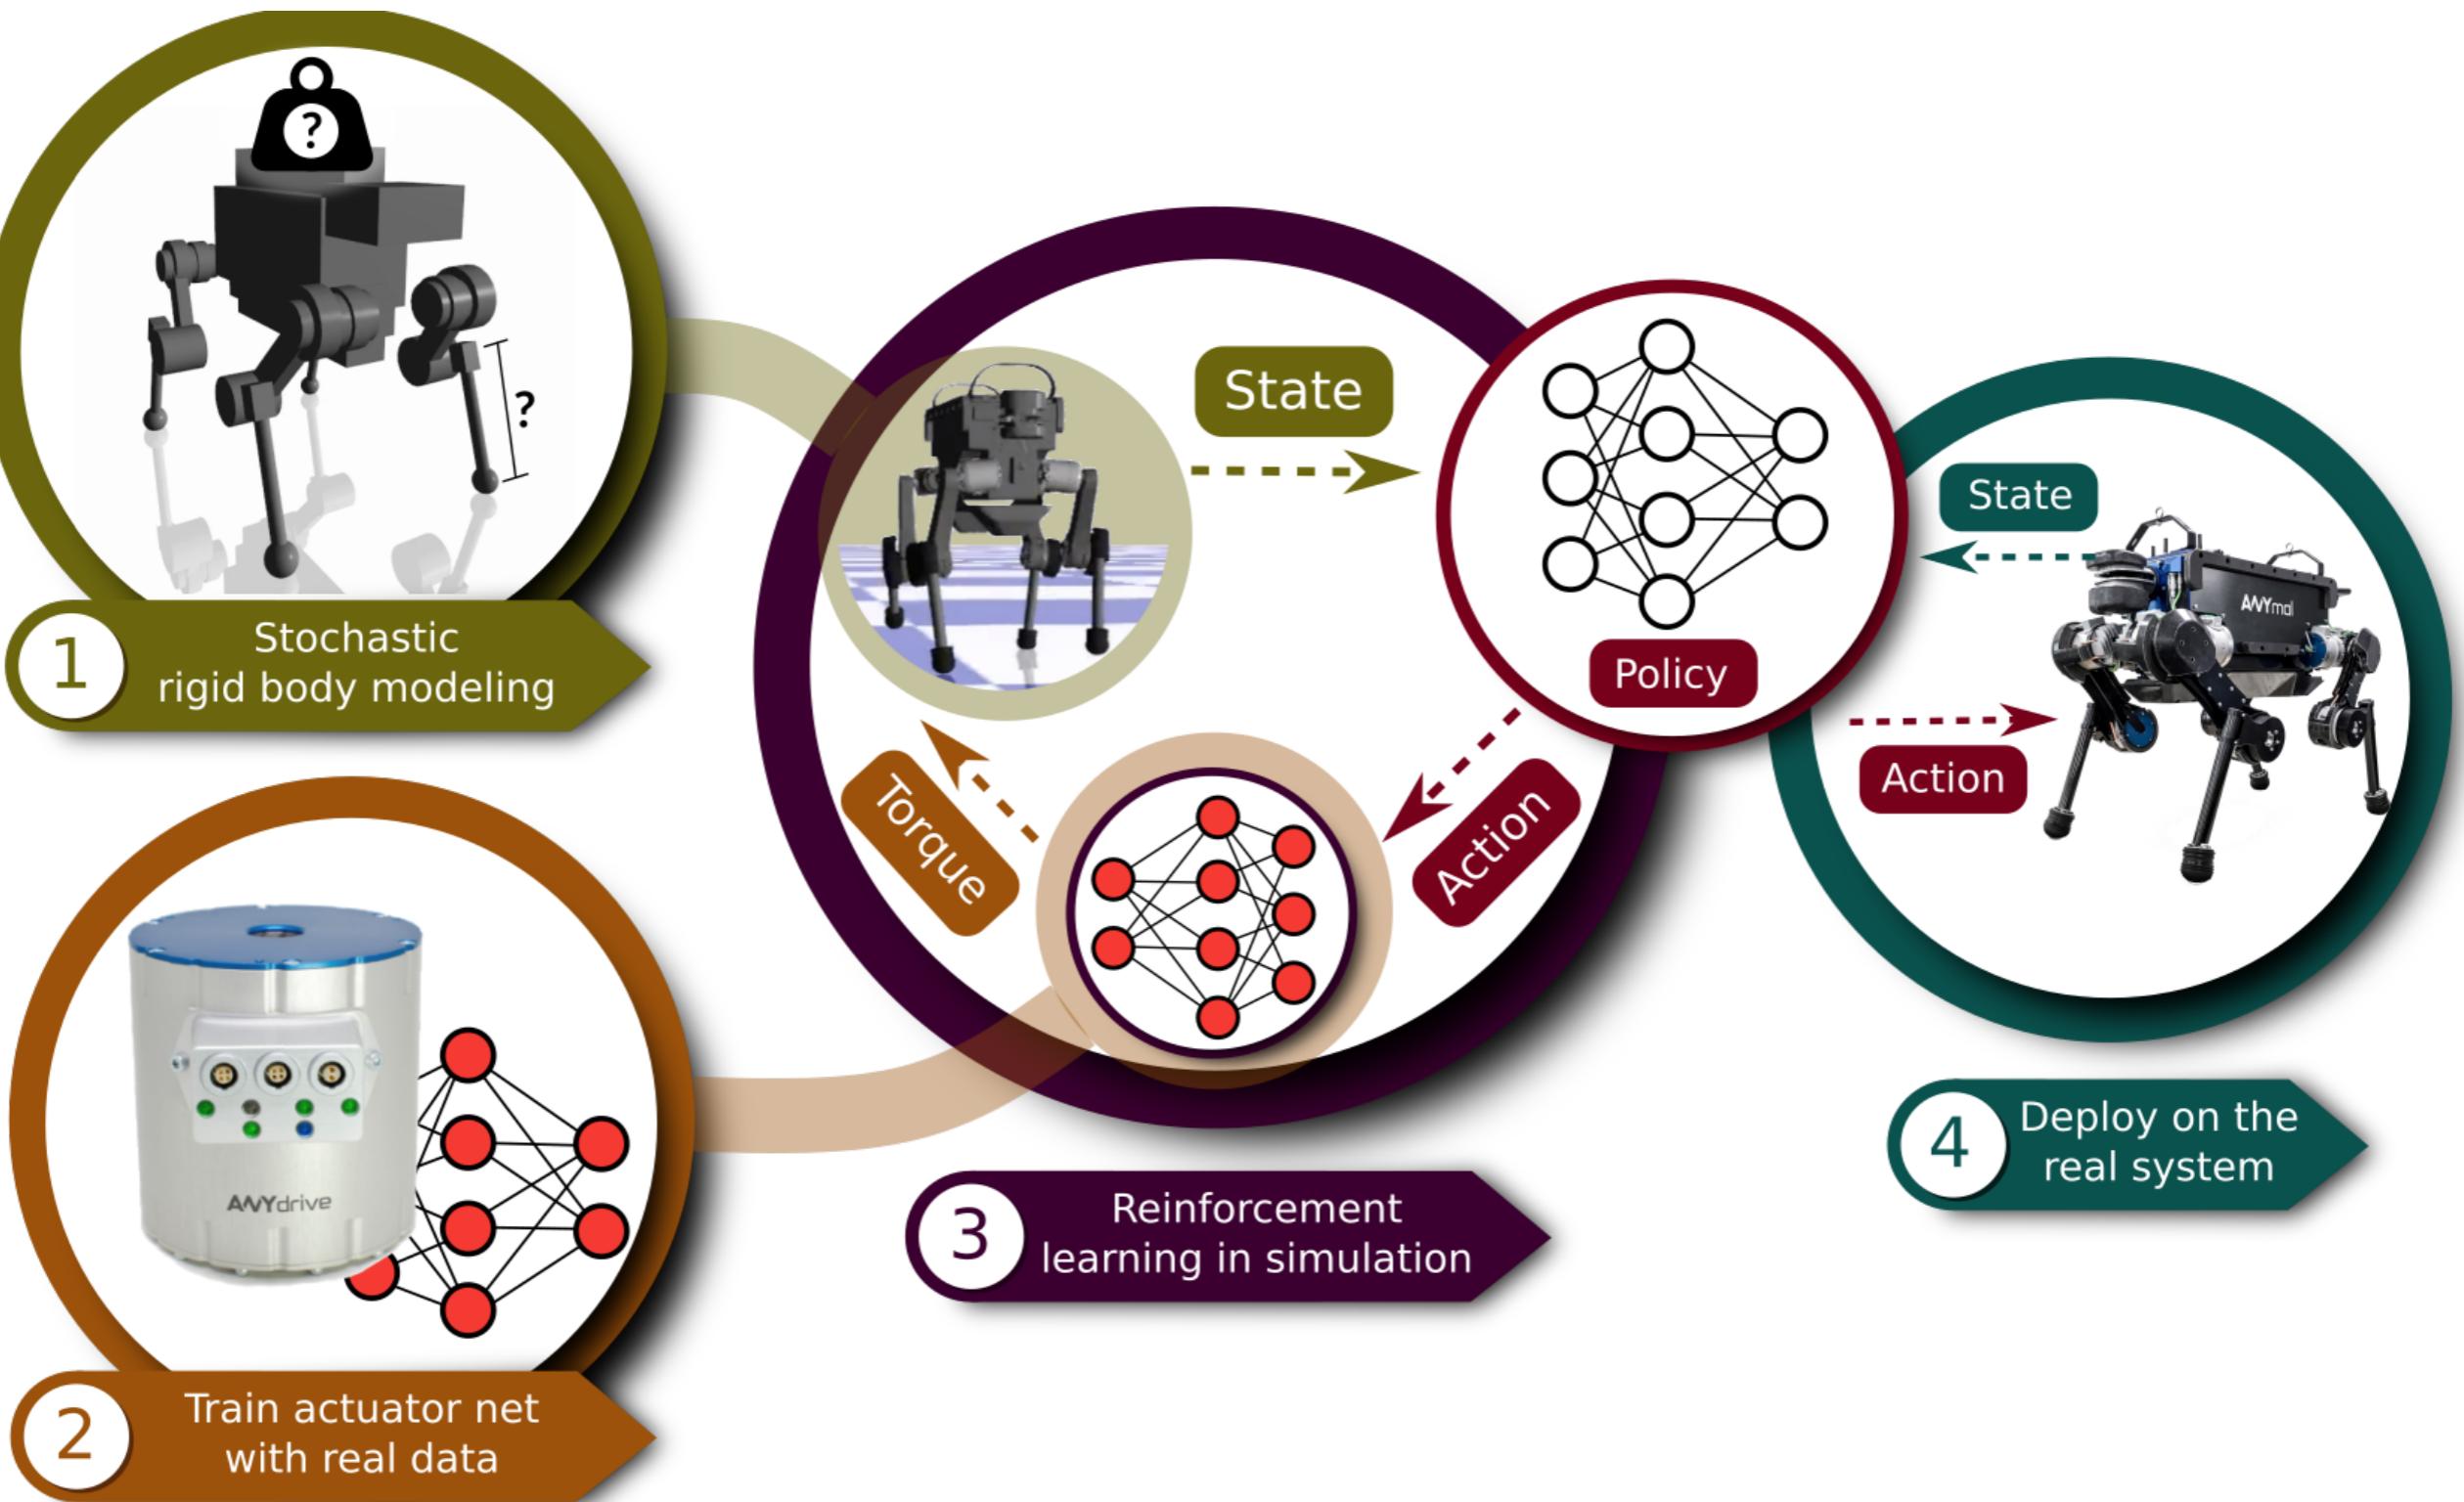


Controlled via SEA motors,
hard to analytically model

- Faster, more energy efficient motion than hand-crafted controllers
- Robust recovery motions

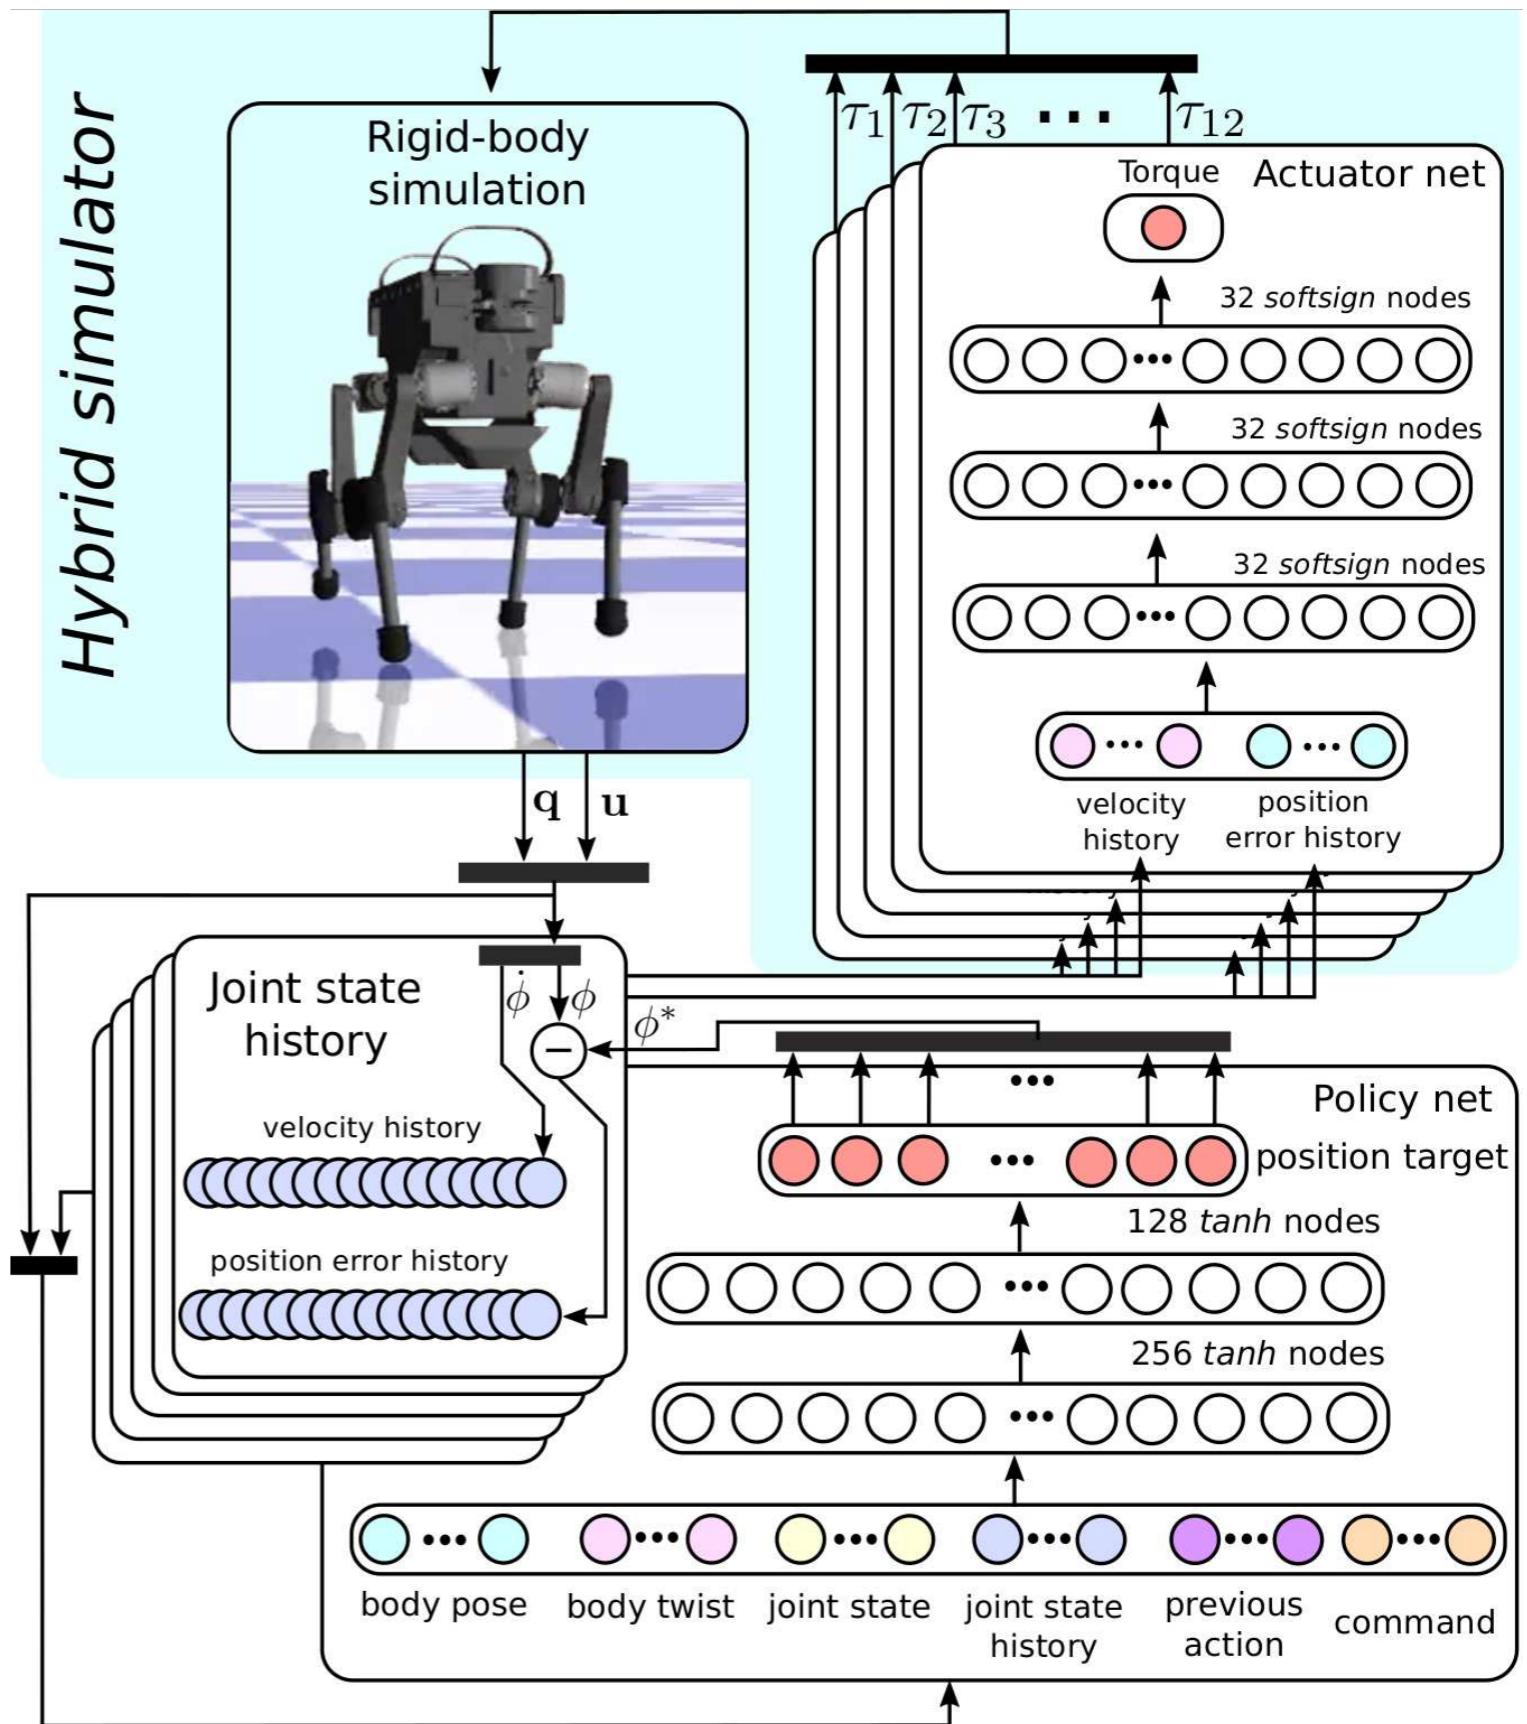


Overall Pipeline

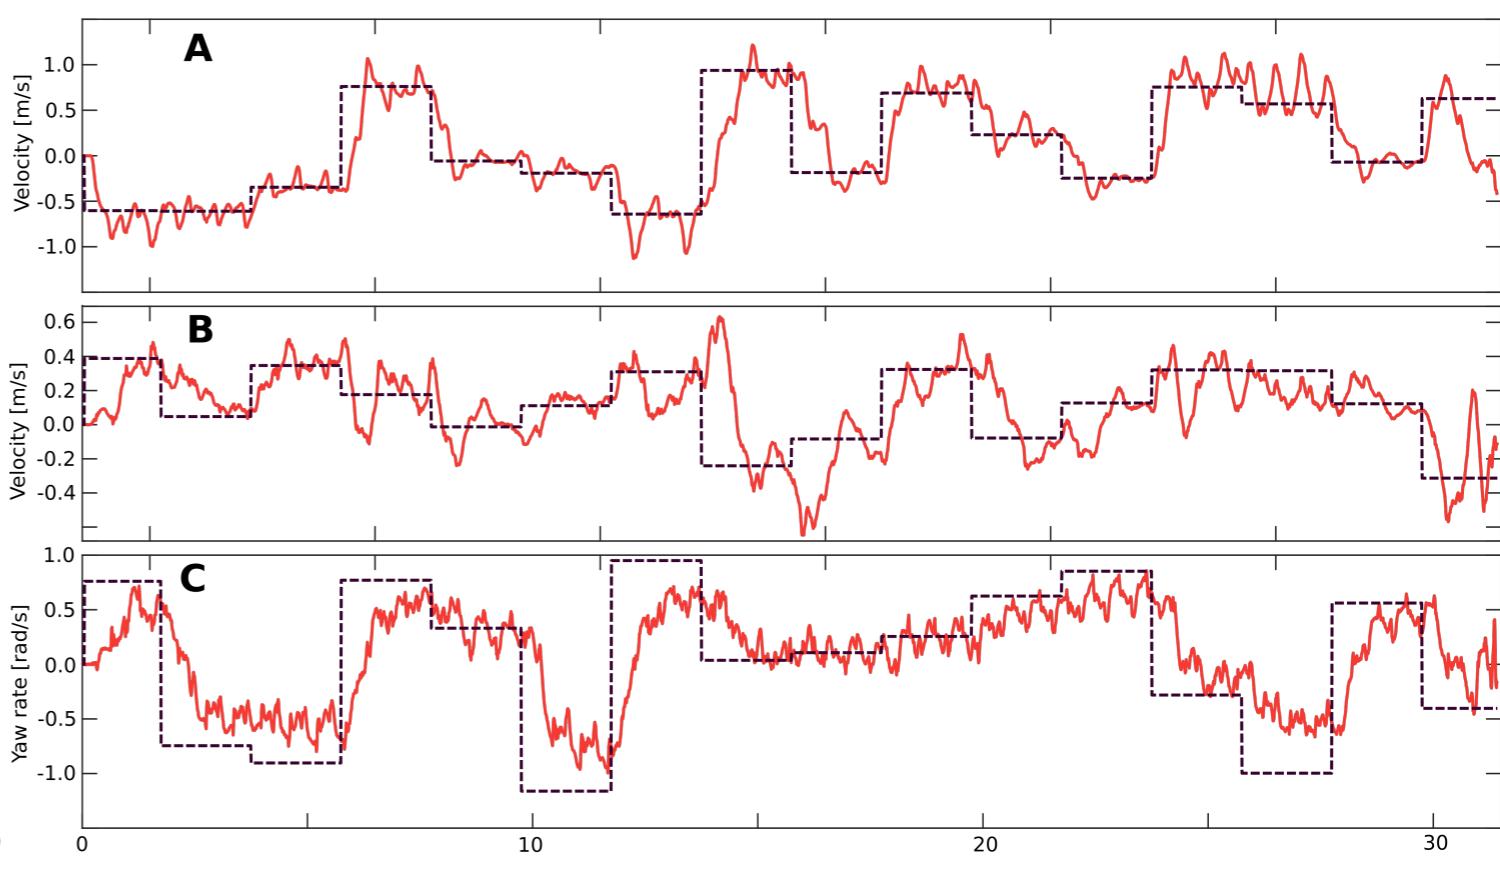
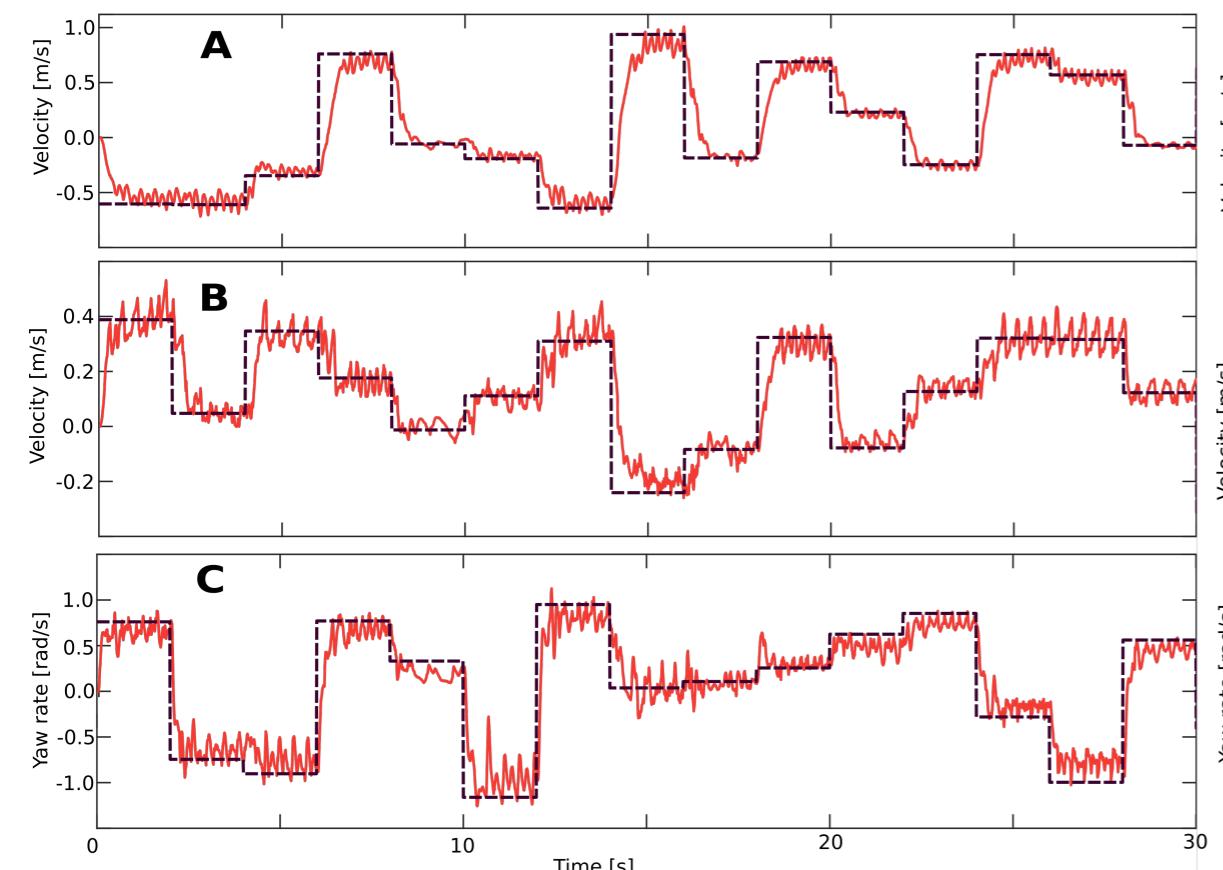


Actuator Net

- Training:
 - Collect data on physical platform by executing hand-crafted gaits
 - varying frequencies, speeds
 - Record commanded joint locations, and associated torques

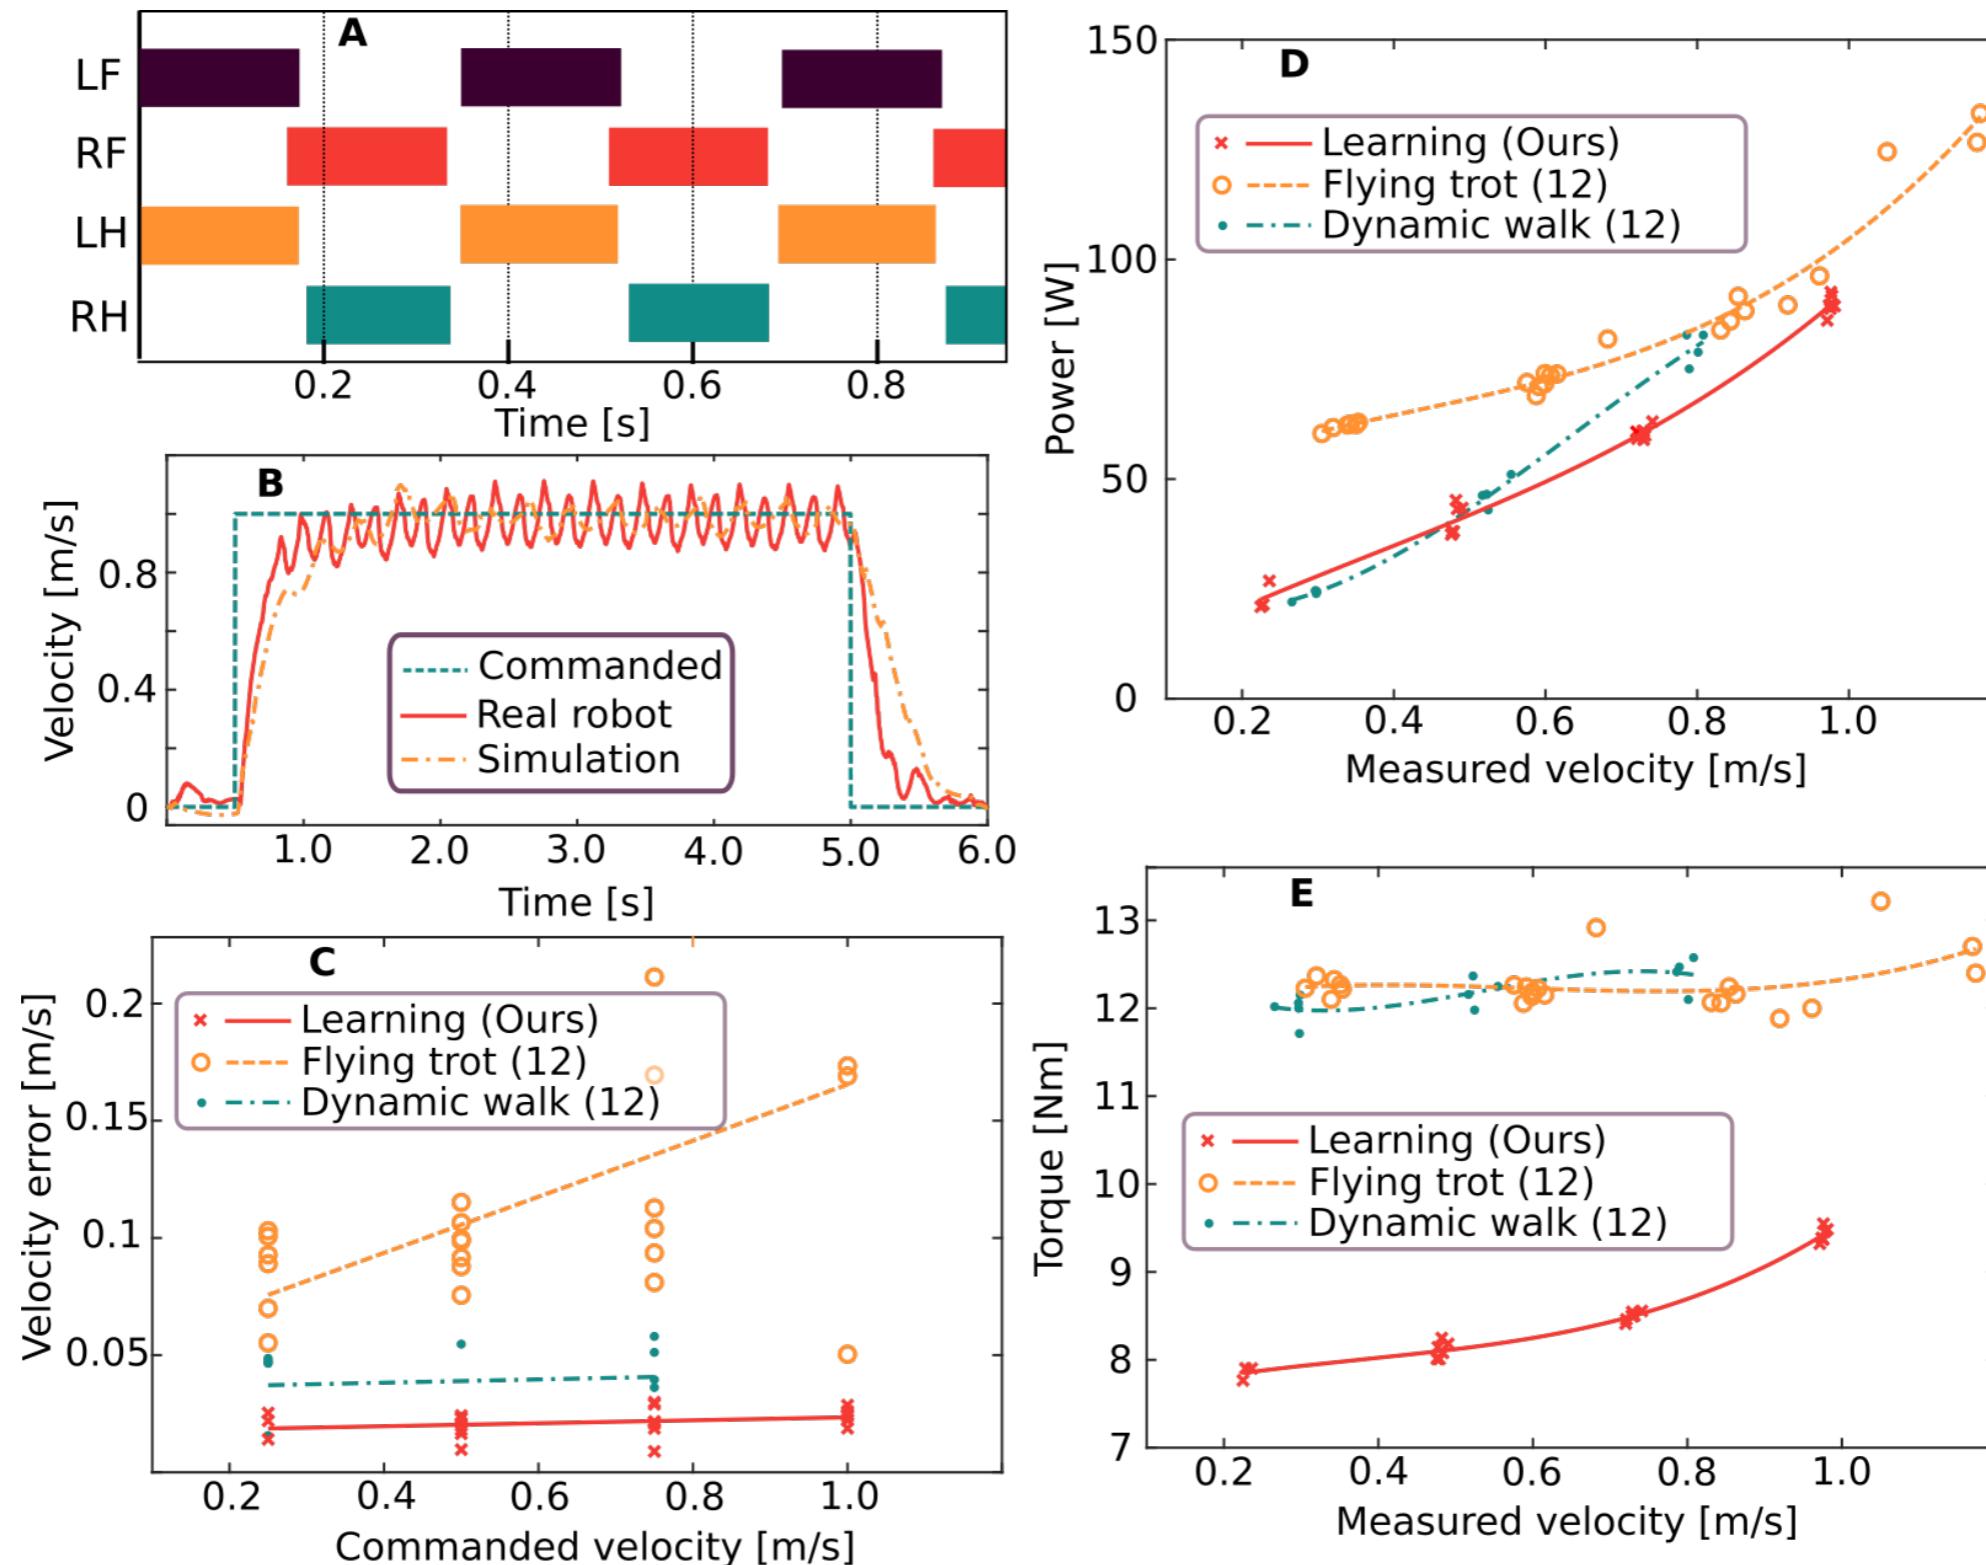


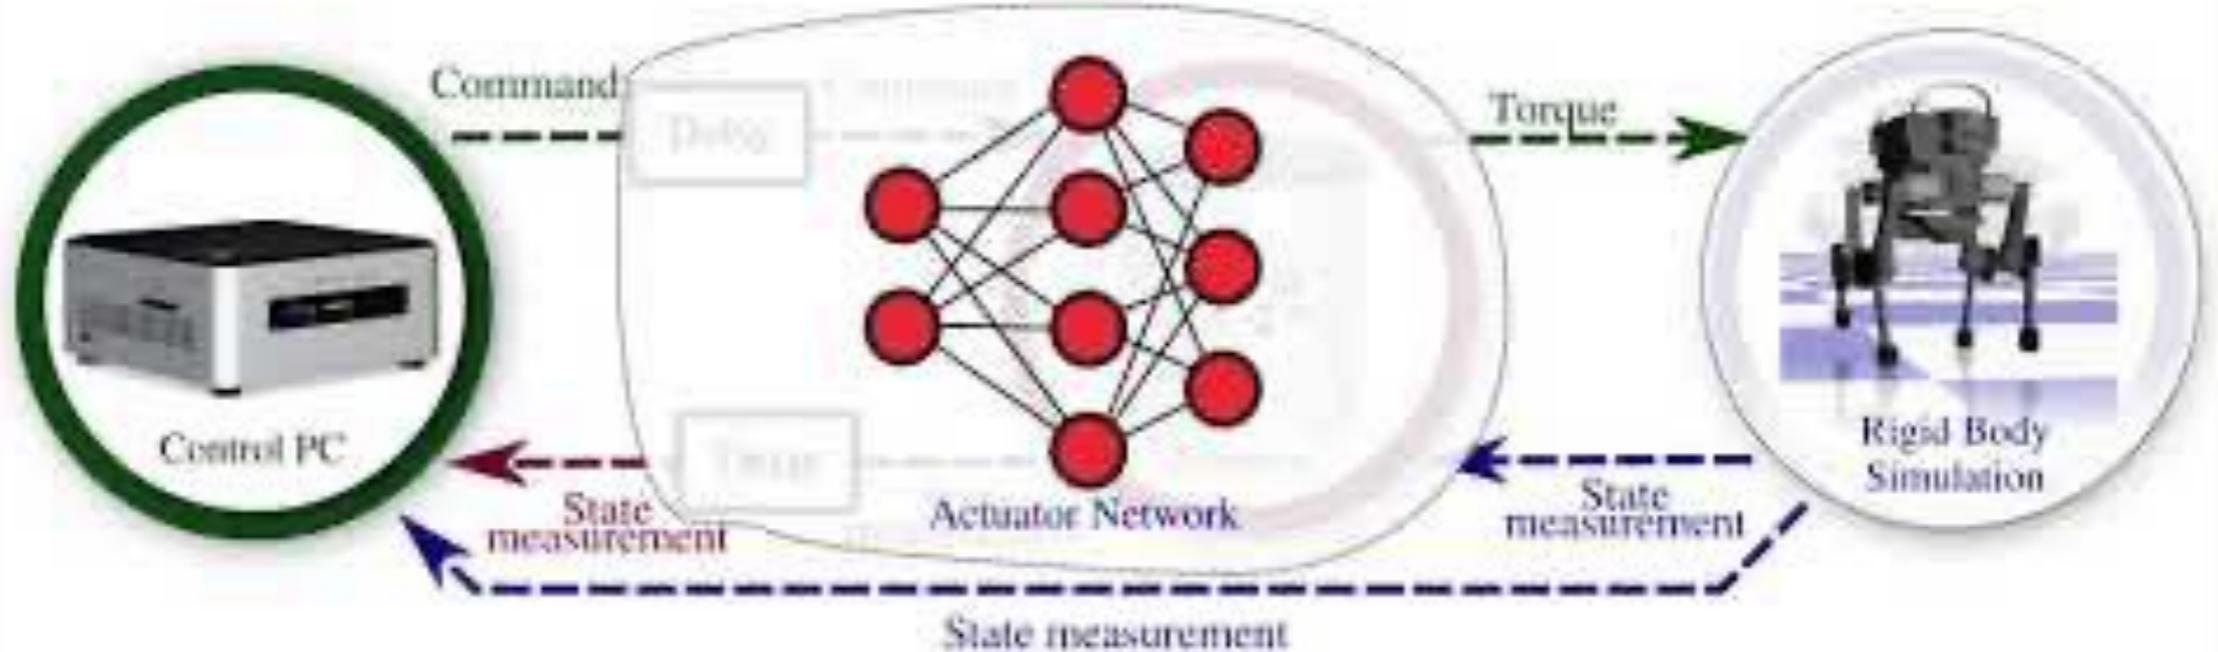
Control Performance at Test-time



Control Performance at Test-time

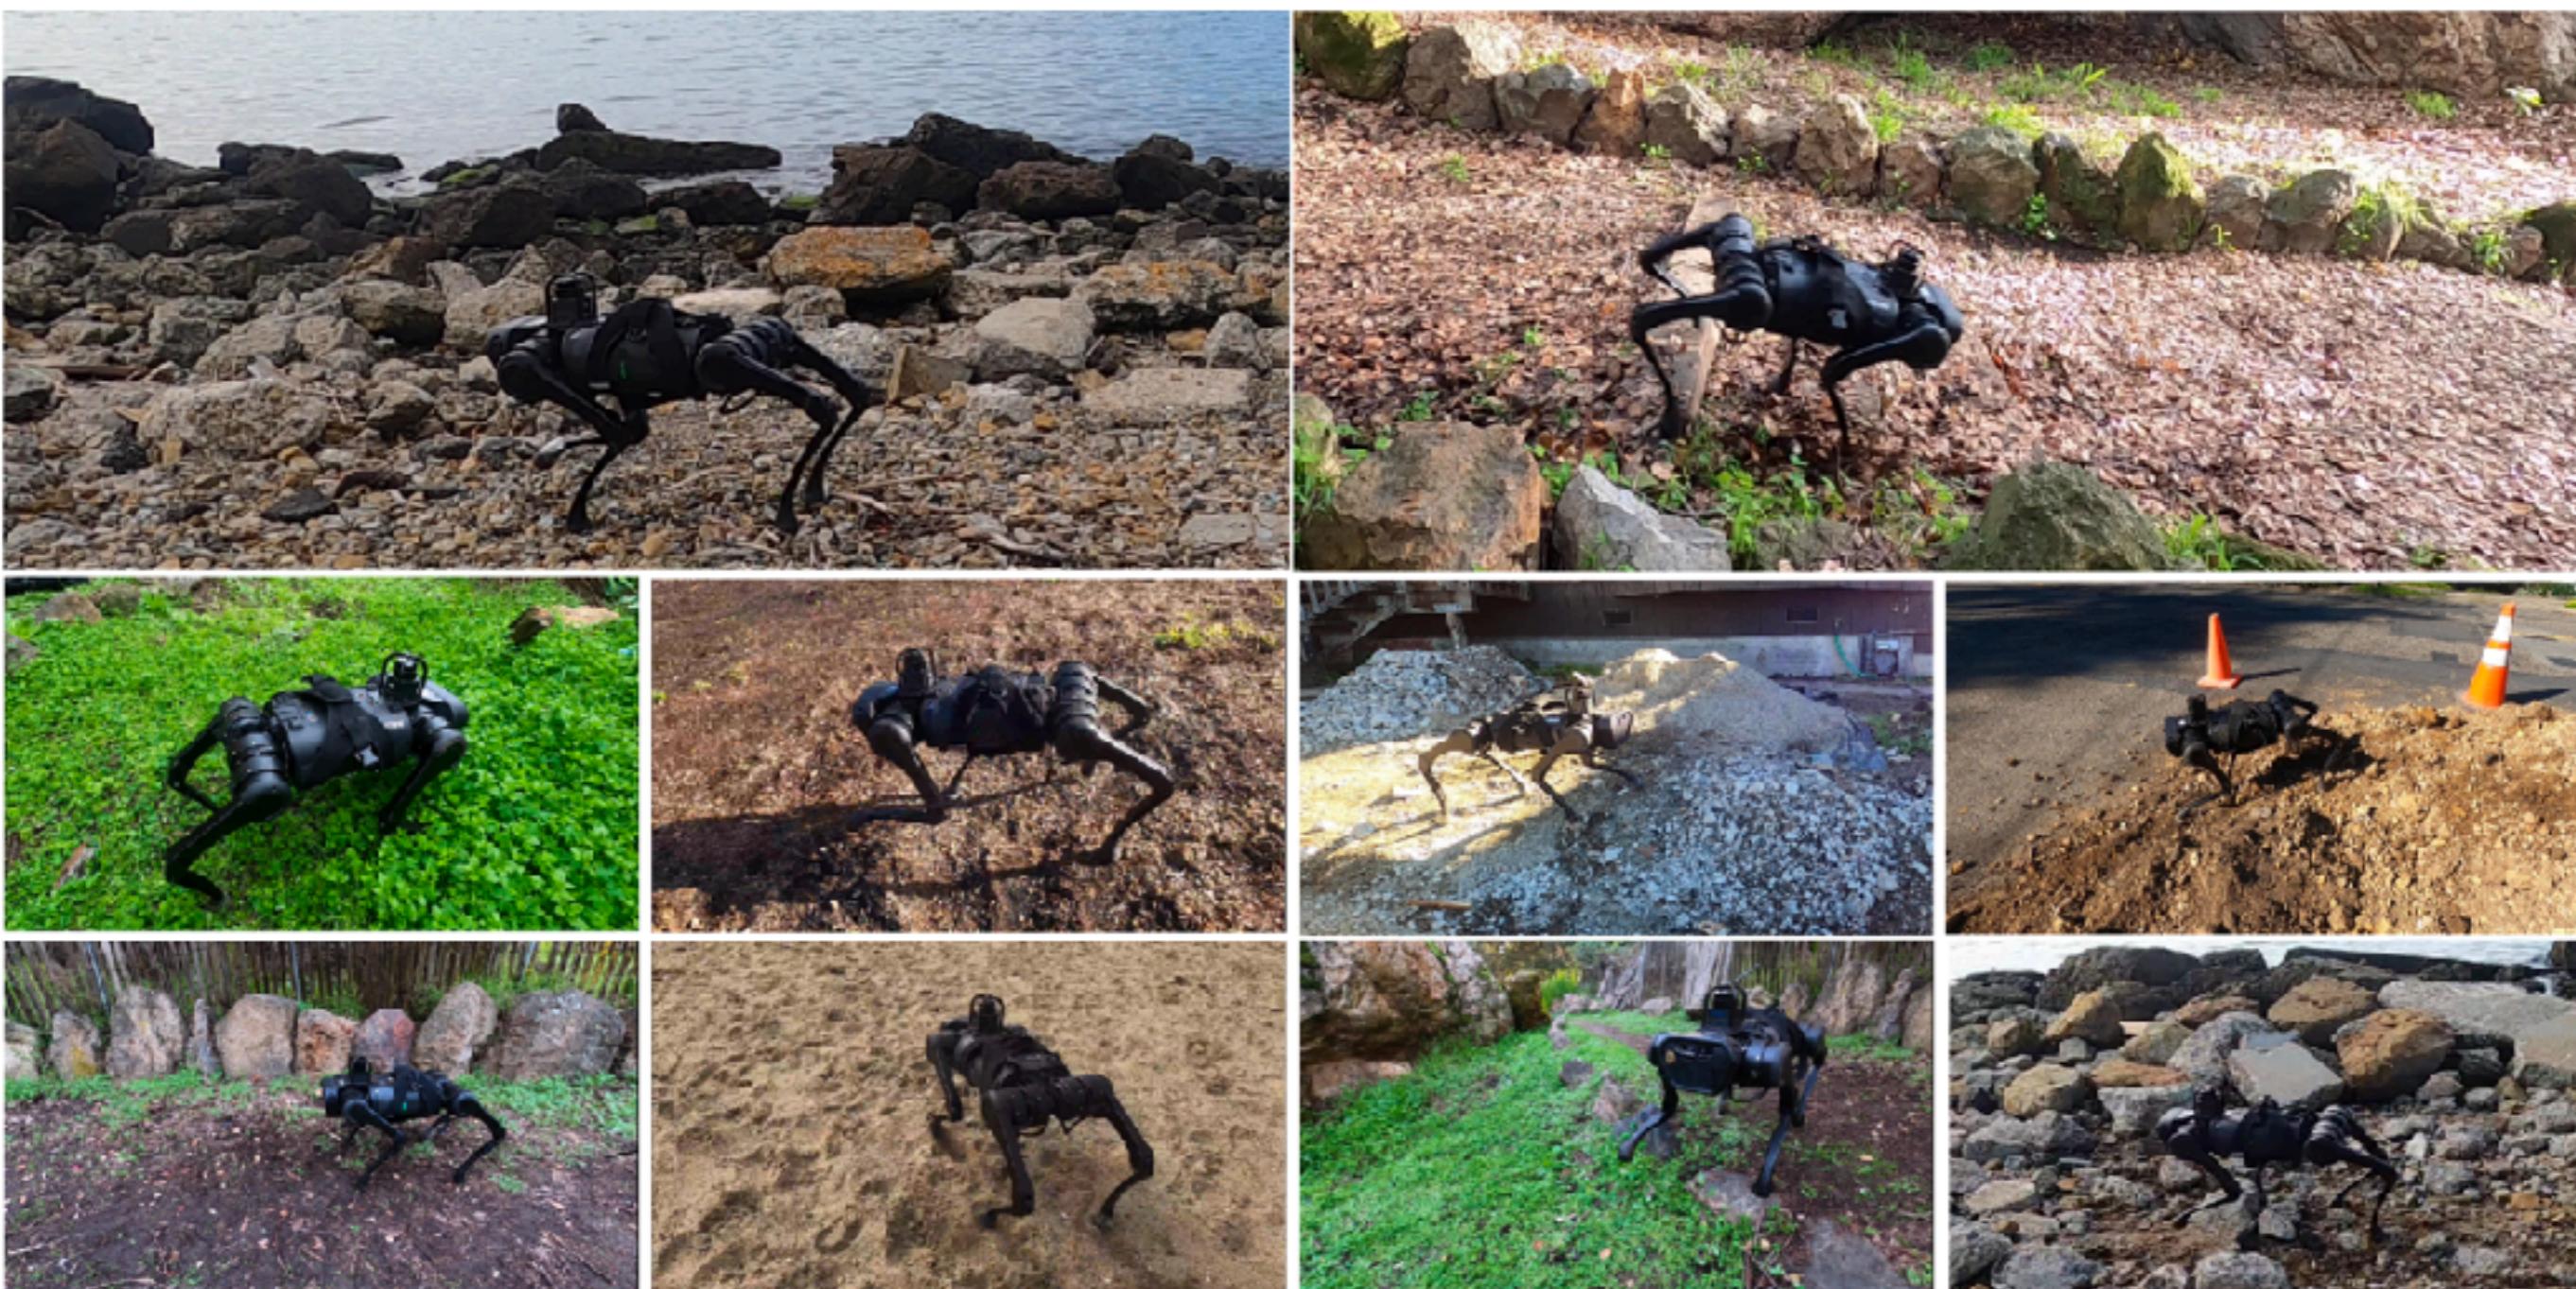
Efficient test-time motion





To this end, we train a neural network representing this complex dynamics with data from the real robot.

RMA (Rapid Motor Adaptation)



Motivation

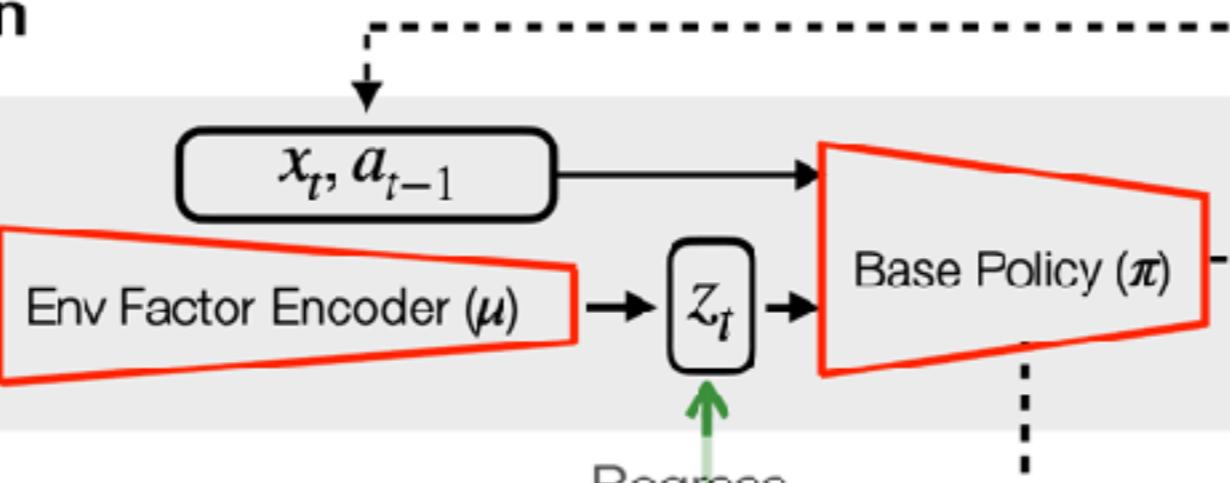
- Hand-designing controllers for complex systems is hard
 - Employ learning in simulation
- Learning with domain randomization causes overly-cautious behavior; walking on different terrains may need different behavior
 - Adapt controllers to the situation you are in
- System identification
 - Full system identification may not be necessary

RMA (Rapid Motor Adaptation)

A) Training in Simulation

Phase 1

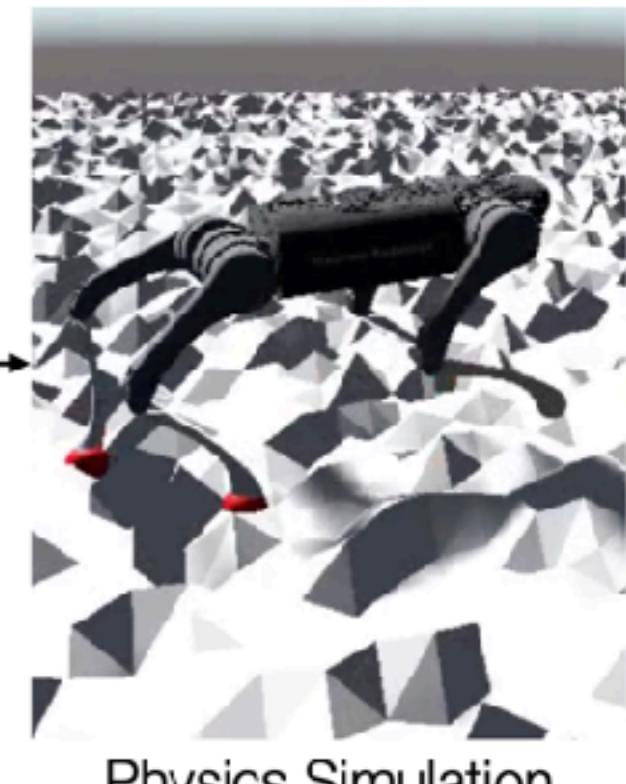
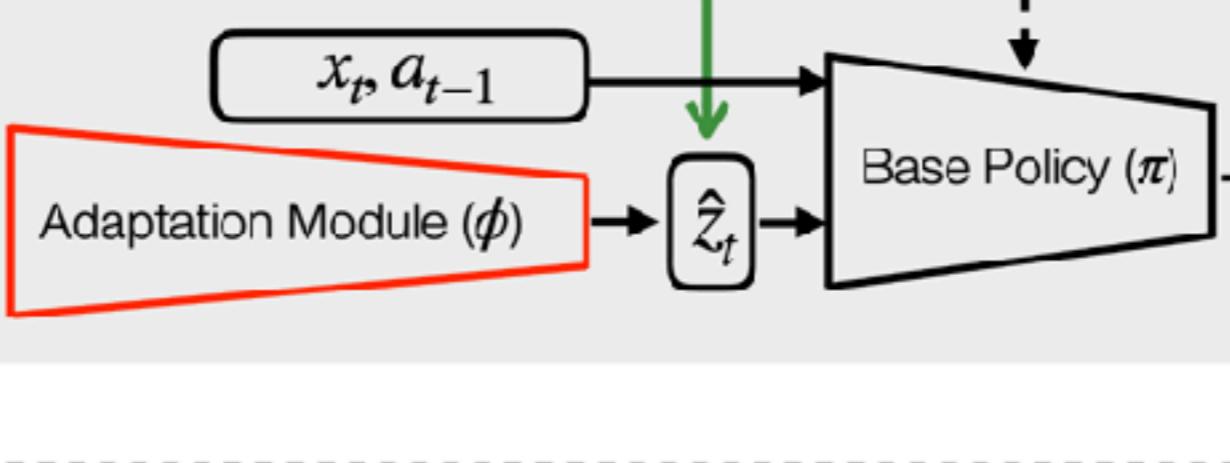
Mass, COM, Friction
Terrain Height
Motor Strength
 (e_t)



*Trainable Modules in Red

Phase 2

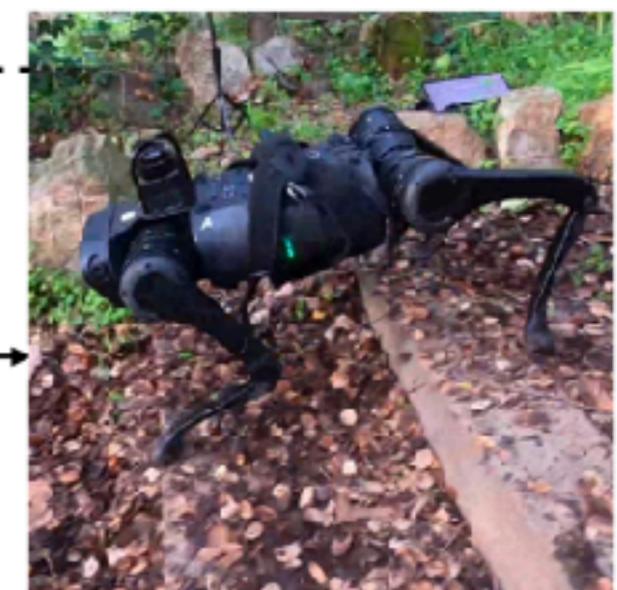
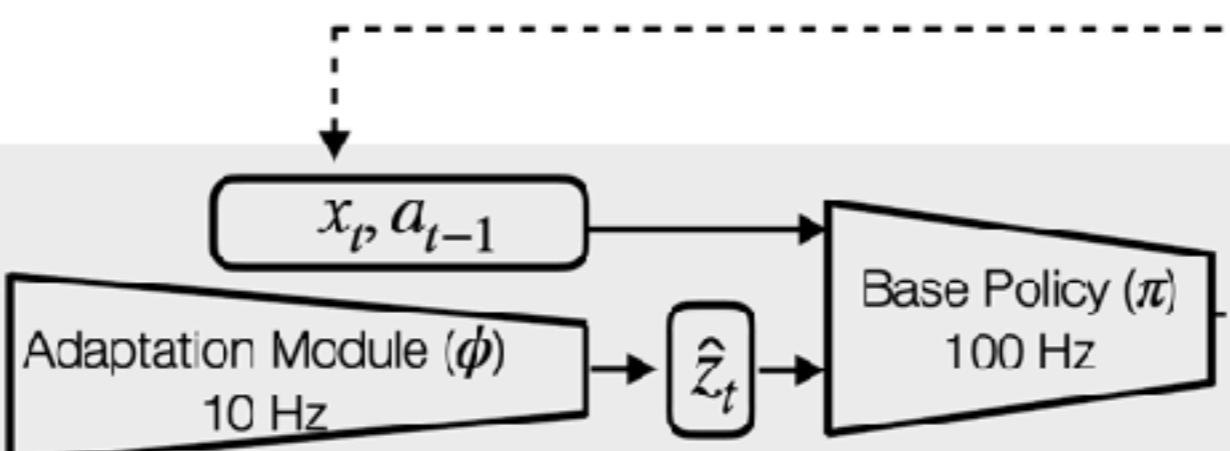
x_{t-51}, a_{t-51}
:
 x_{t-1}, a_{t-1}



Physics Simulation

B) Deployment

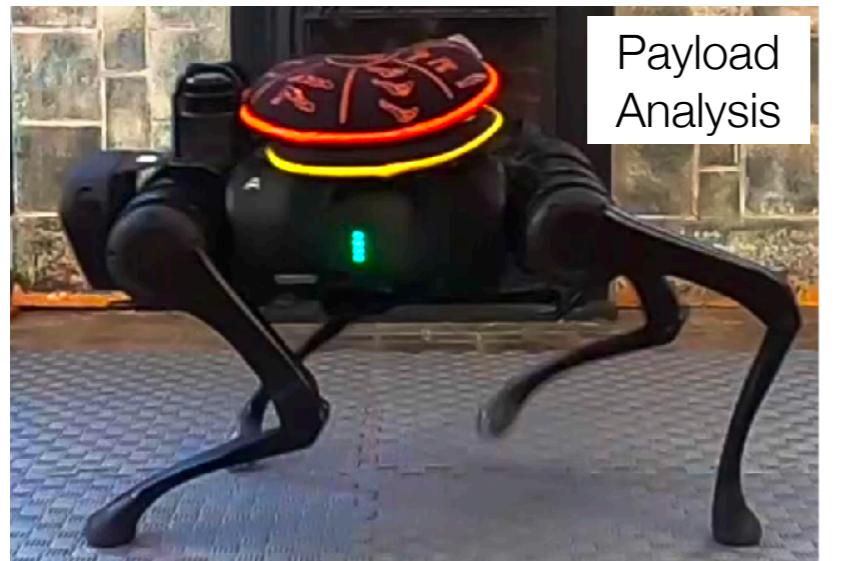
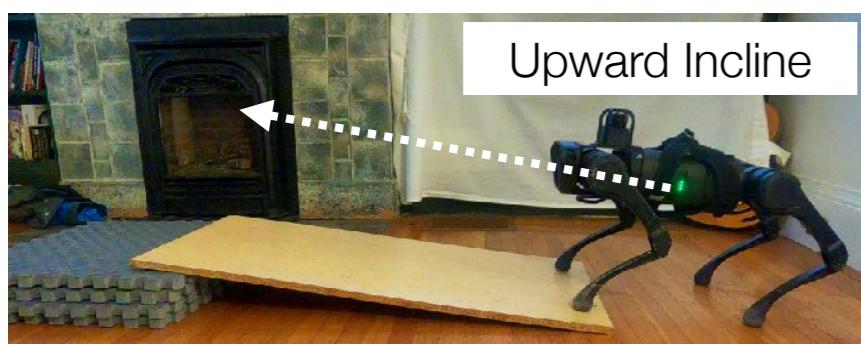
x_{t-50}, a_{t-51}
:
 x_t, a_{t-1}



PID Controllers

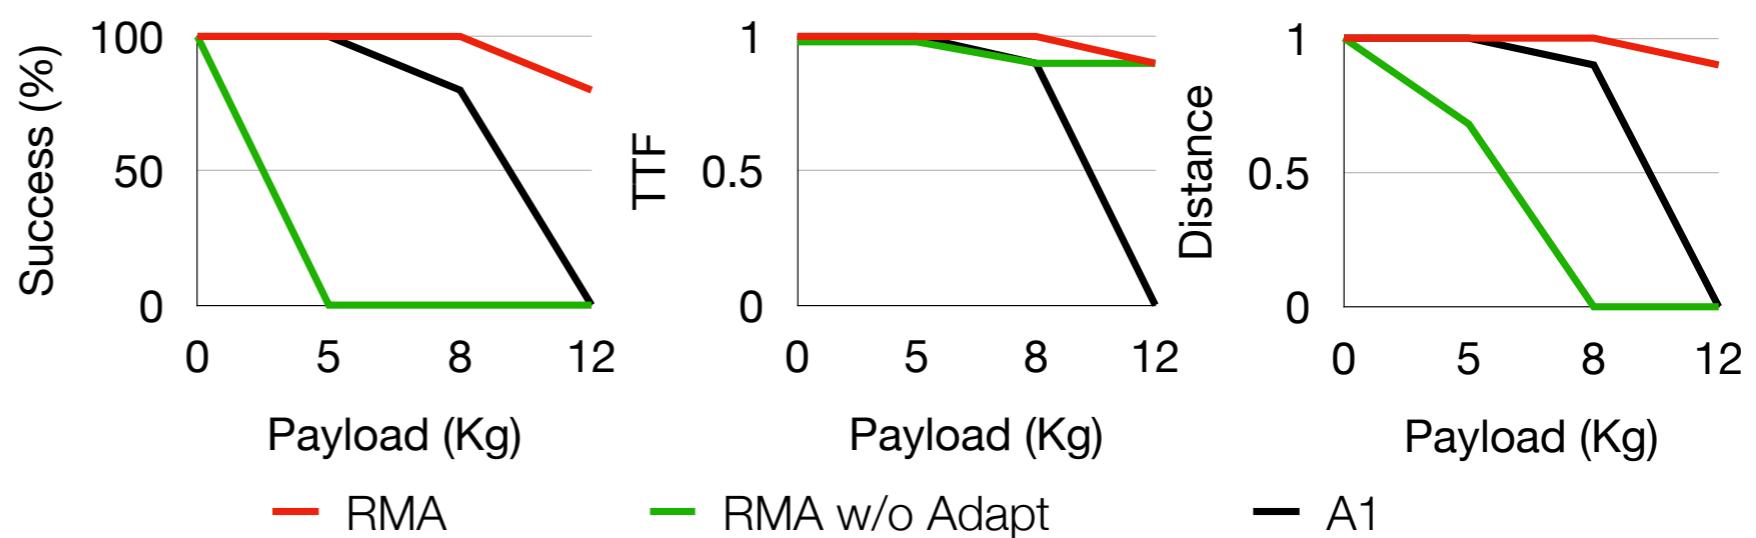
- Consider a 1D system with 1D control. Unknown dynamics, but we know that a positive control u , increases x
- Suppose we want to bring the system to a desired state x_d
- Proportional control: $u_P = -k_p(x - x_d)$
- Just P control has the risk that system has a non-zero steady state error (think a robot arm under gravity)
- Add an integral term: $u_I = -k_I I(t), I(t) = \int_0^t (x(t) - x_d(t))dt$
- Just PI control can cause systems to oscillate
 - Add a derivative term: $u_D = -k_D(\dot{x} - \dot{x}_d)$
- Total control: $u_P + u_I + u_D$
- <https://www.matthewpeterkelly.com/tutorials/pdControl/index.html>

Results

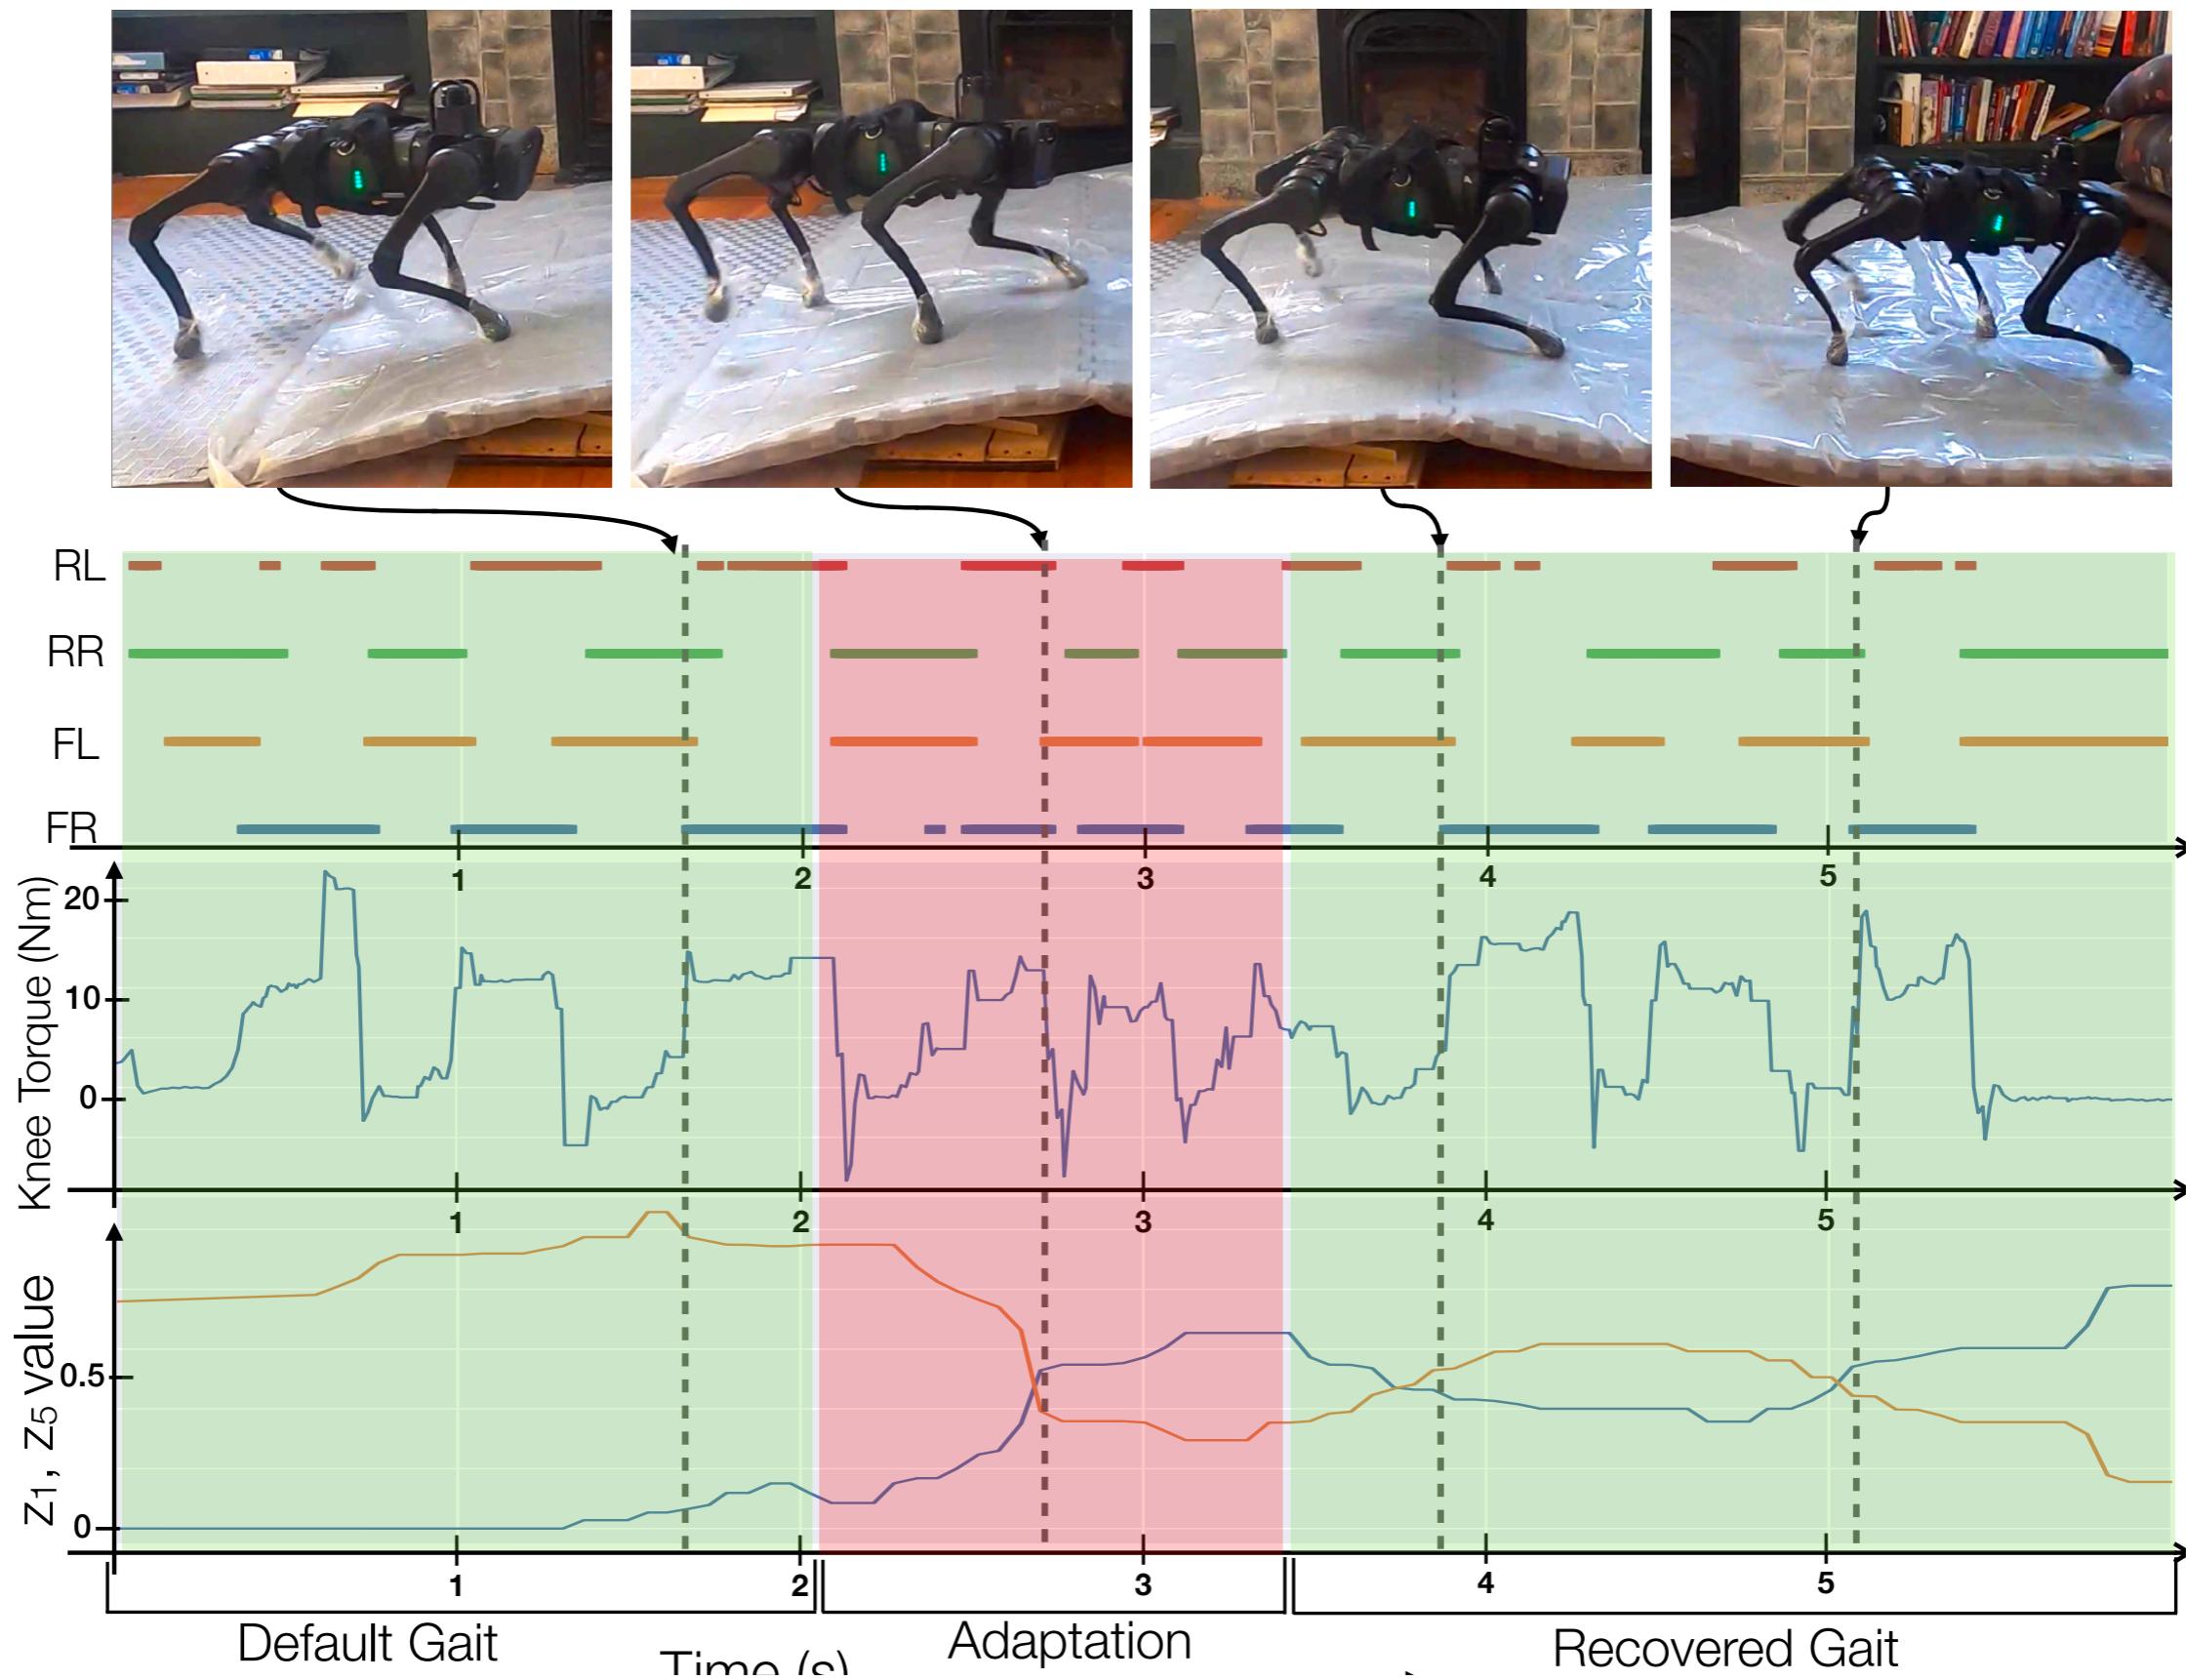


	Success	TTF	Distance
RMA	80	0.9	0.9
RMA w/o Adapt	0	0.1	0.1
A1	20	0.36	0.36

Success
RMA
RMA w/o Adapt
A1



Adaptation



Thank you