

Navigation

Saurabh Gupta

Tasks

- Locomotion
 - Wheeled ground vehicle
 - Drone
 - Quadruped
 - Off-road terrain



Tasks

- Locomotion
 - Platforms:
 - Wheeled ground vehicle, Drone, Quadrupeds, Off-road terrain
 - Move safely, and gracefully. Avoid collisions.
 - Follow high-level commands
 - Visual Servoing, follow an object
- Goal-directed behavior
 - Goals known / unknown
 - Environment seen / novel

Axes of variation

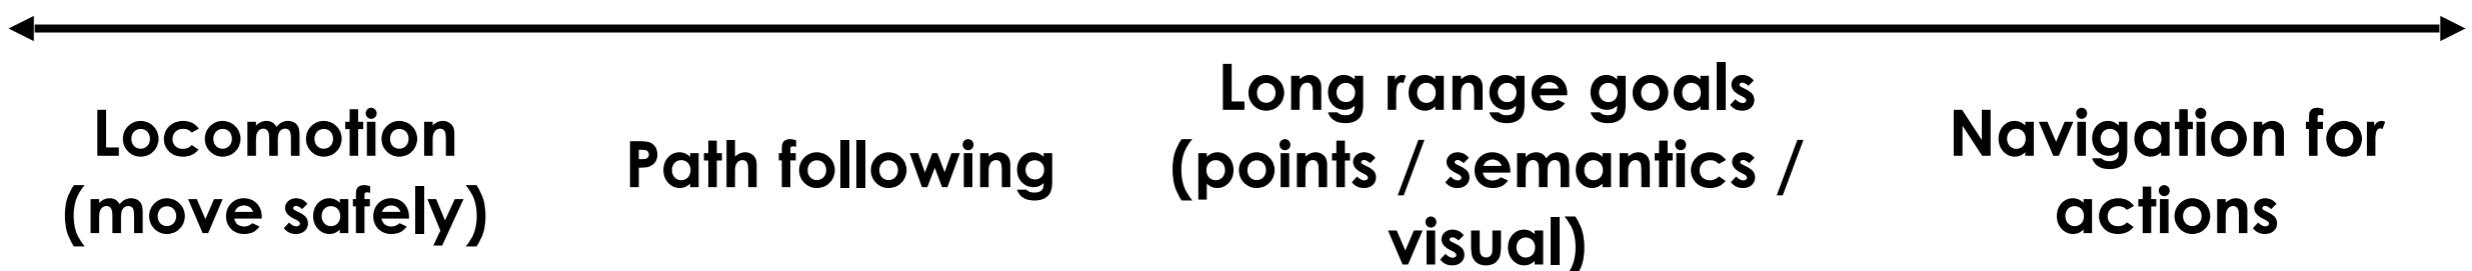
Supervision Signal



Environment Modeling



Task



Generalization



Axes of variation

Abstraction



Environment Properties



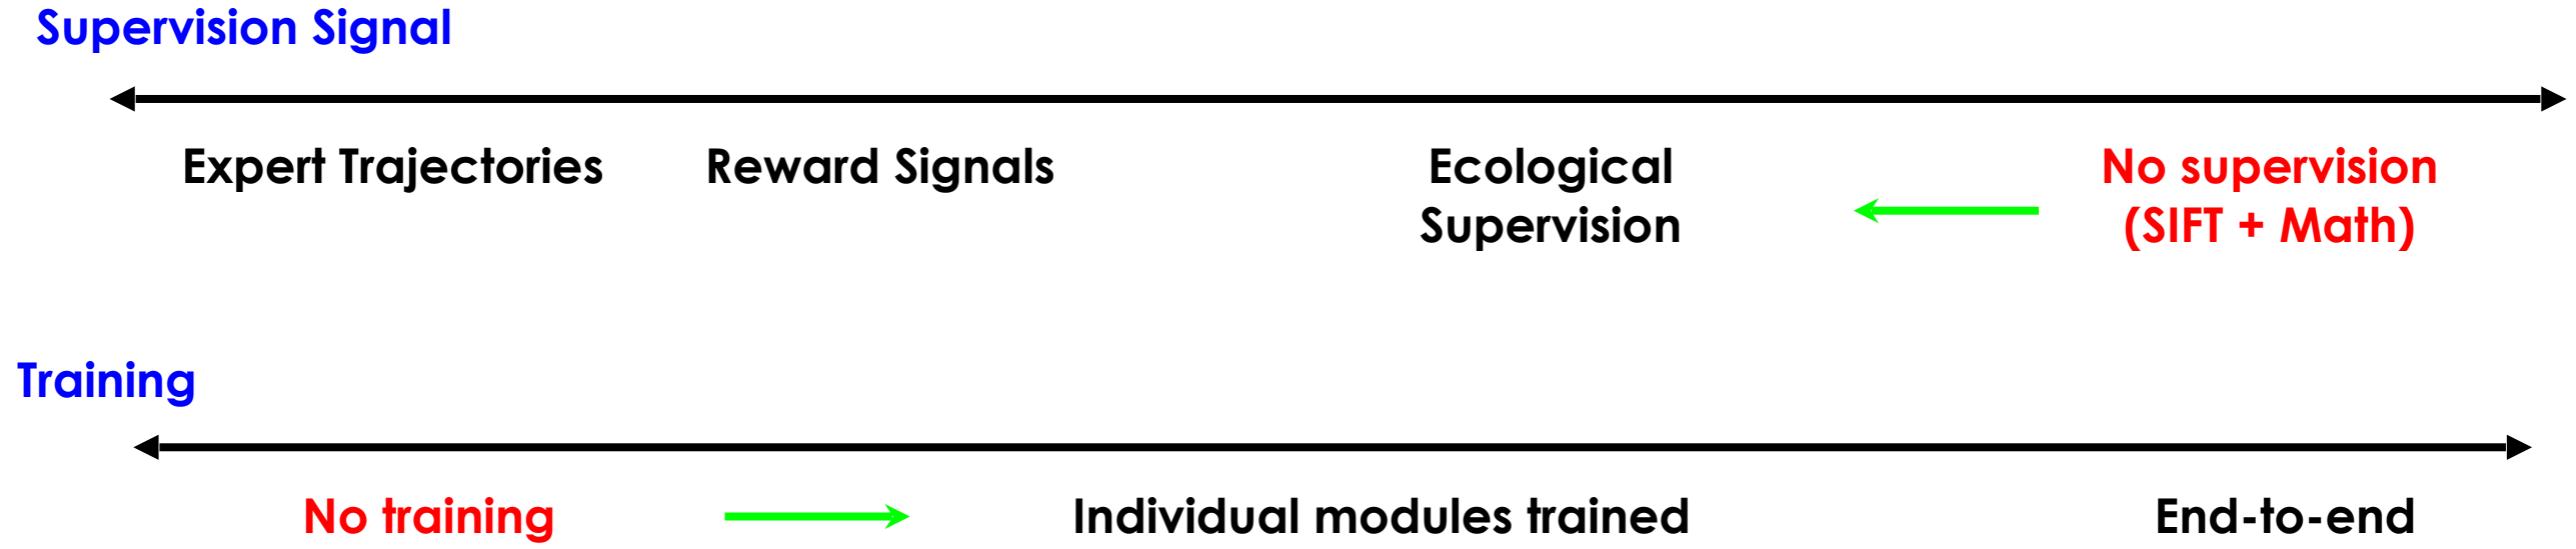
Experimental setup



Training



CNNs for SLAM

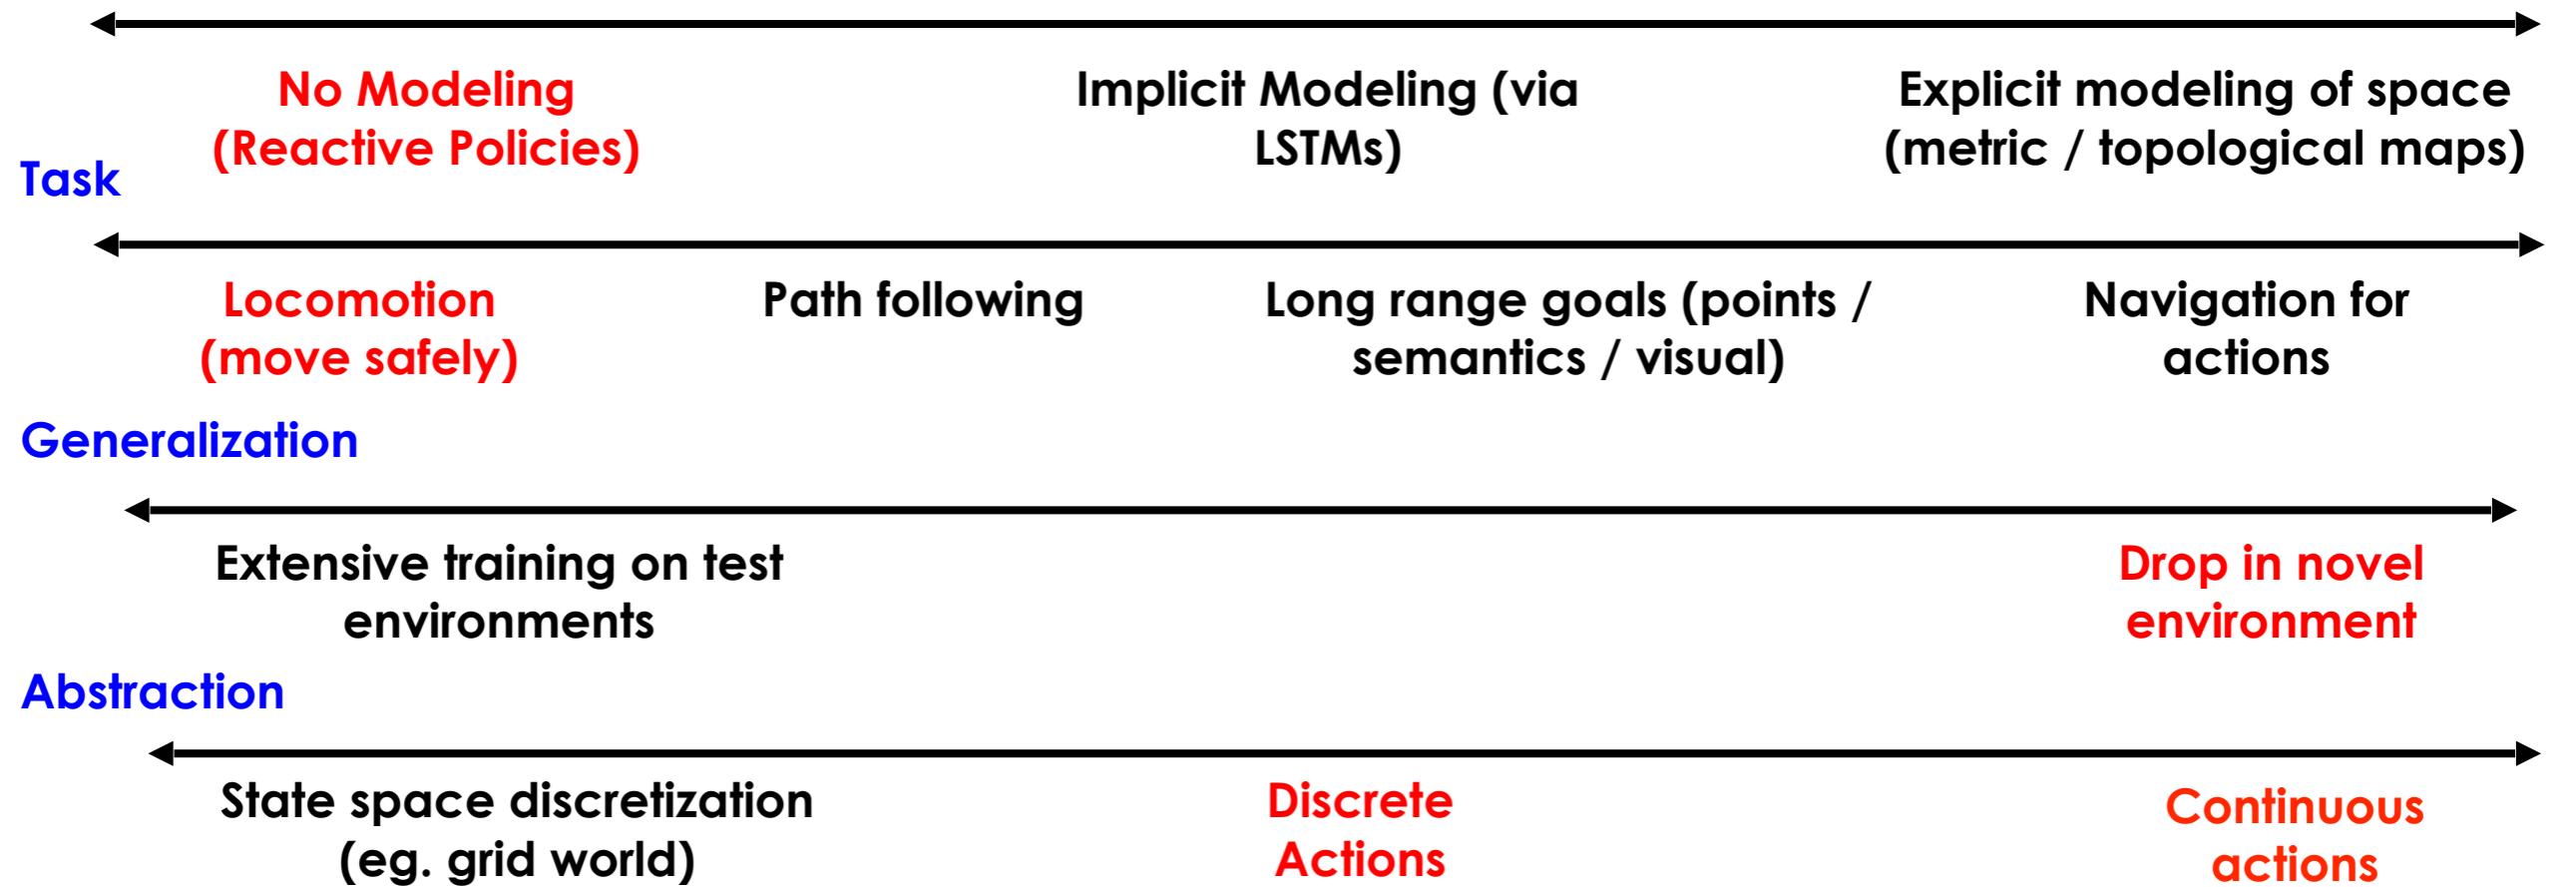


Classical SLAM

- **Localization:**
 - **Semantic Visual Localization** J. Schonberger, M. Pollefeys, A. Geiger, T. Sattler arXiv 17 [pdf](#)
 - **Self-Supervised Place Recognition in Mobile Robots** S. Pillai, J. Leonard. IROS17 [pdf](#)
 - **Active neural localization** D. Chaplot arXiv17 [pdf](#)
- **Visual Odometry**
 - **Towards Visual Ego-motion Learning in Robots.** S. Pillai and J. Leonard IROS17 [pdf](#)
 - **Unsupervised Learning of Depth and Ego-motion from Video** Zhou et al. CVPR17 [pdf](#)
- **Reconstruction**
 - **Learned Stereo Machines.** Kar et al. NIPS17 [pdf](#)
 - **Differentiable ray consistency.** Tulsiani et al. CVPR17 [pdf](#)

Learned Locomotion

Environment Modeling



- **Learning Agile and Dynamic Motor Skills for Legged Robots** J. Hwango et al. Science Robotics 18 [pdf](#), [video](#)

- **Agile Autonomous Driving** Y. Pan et al. R:SS18 [pdf](#), [video](#)

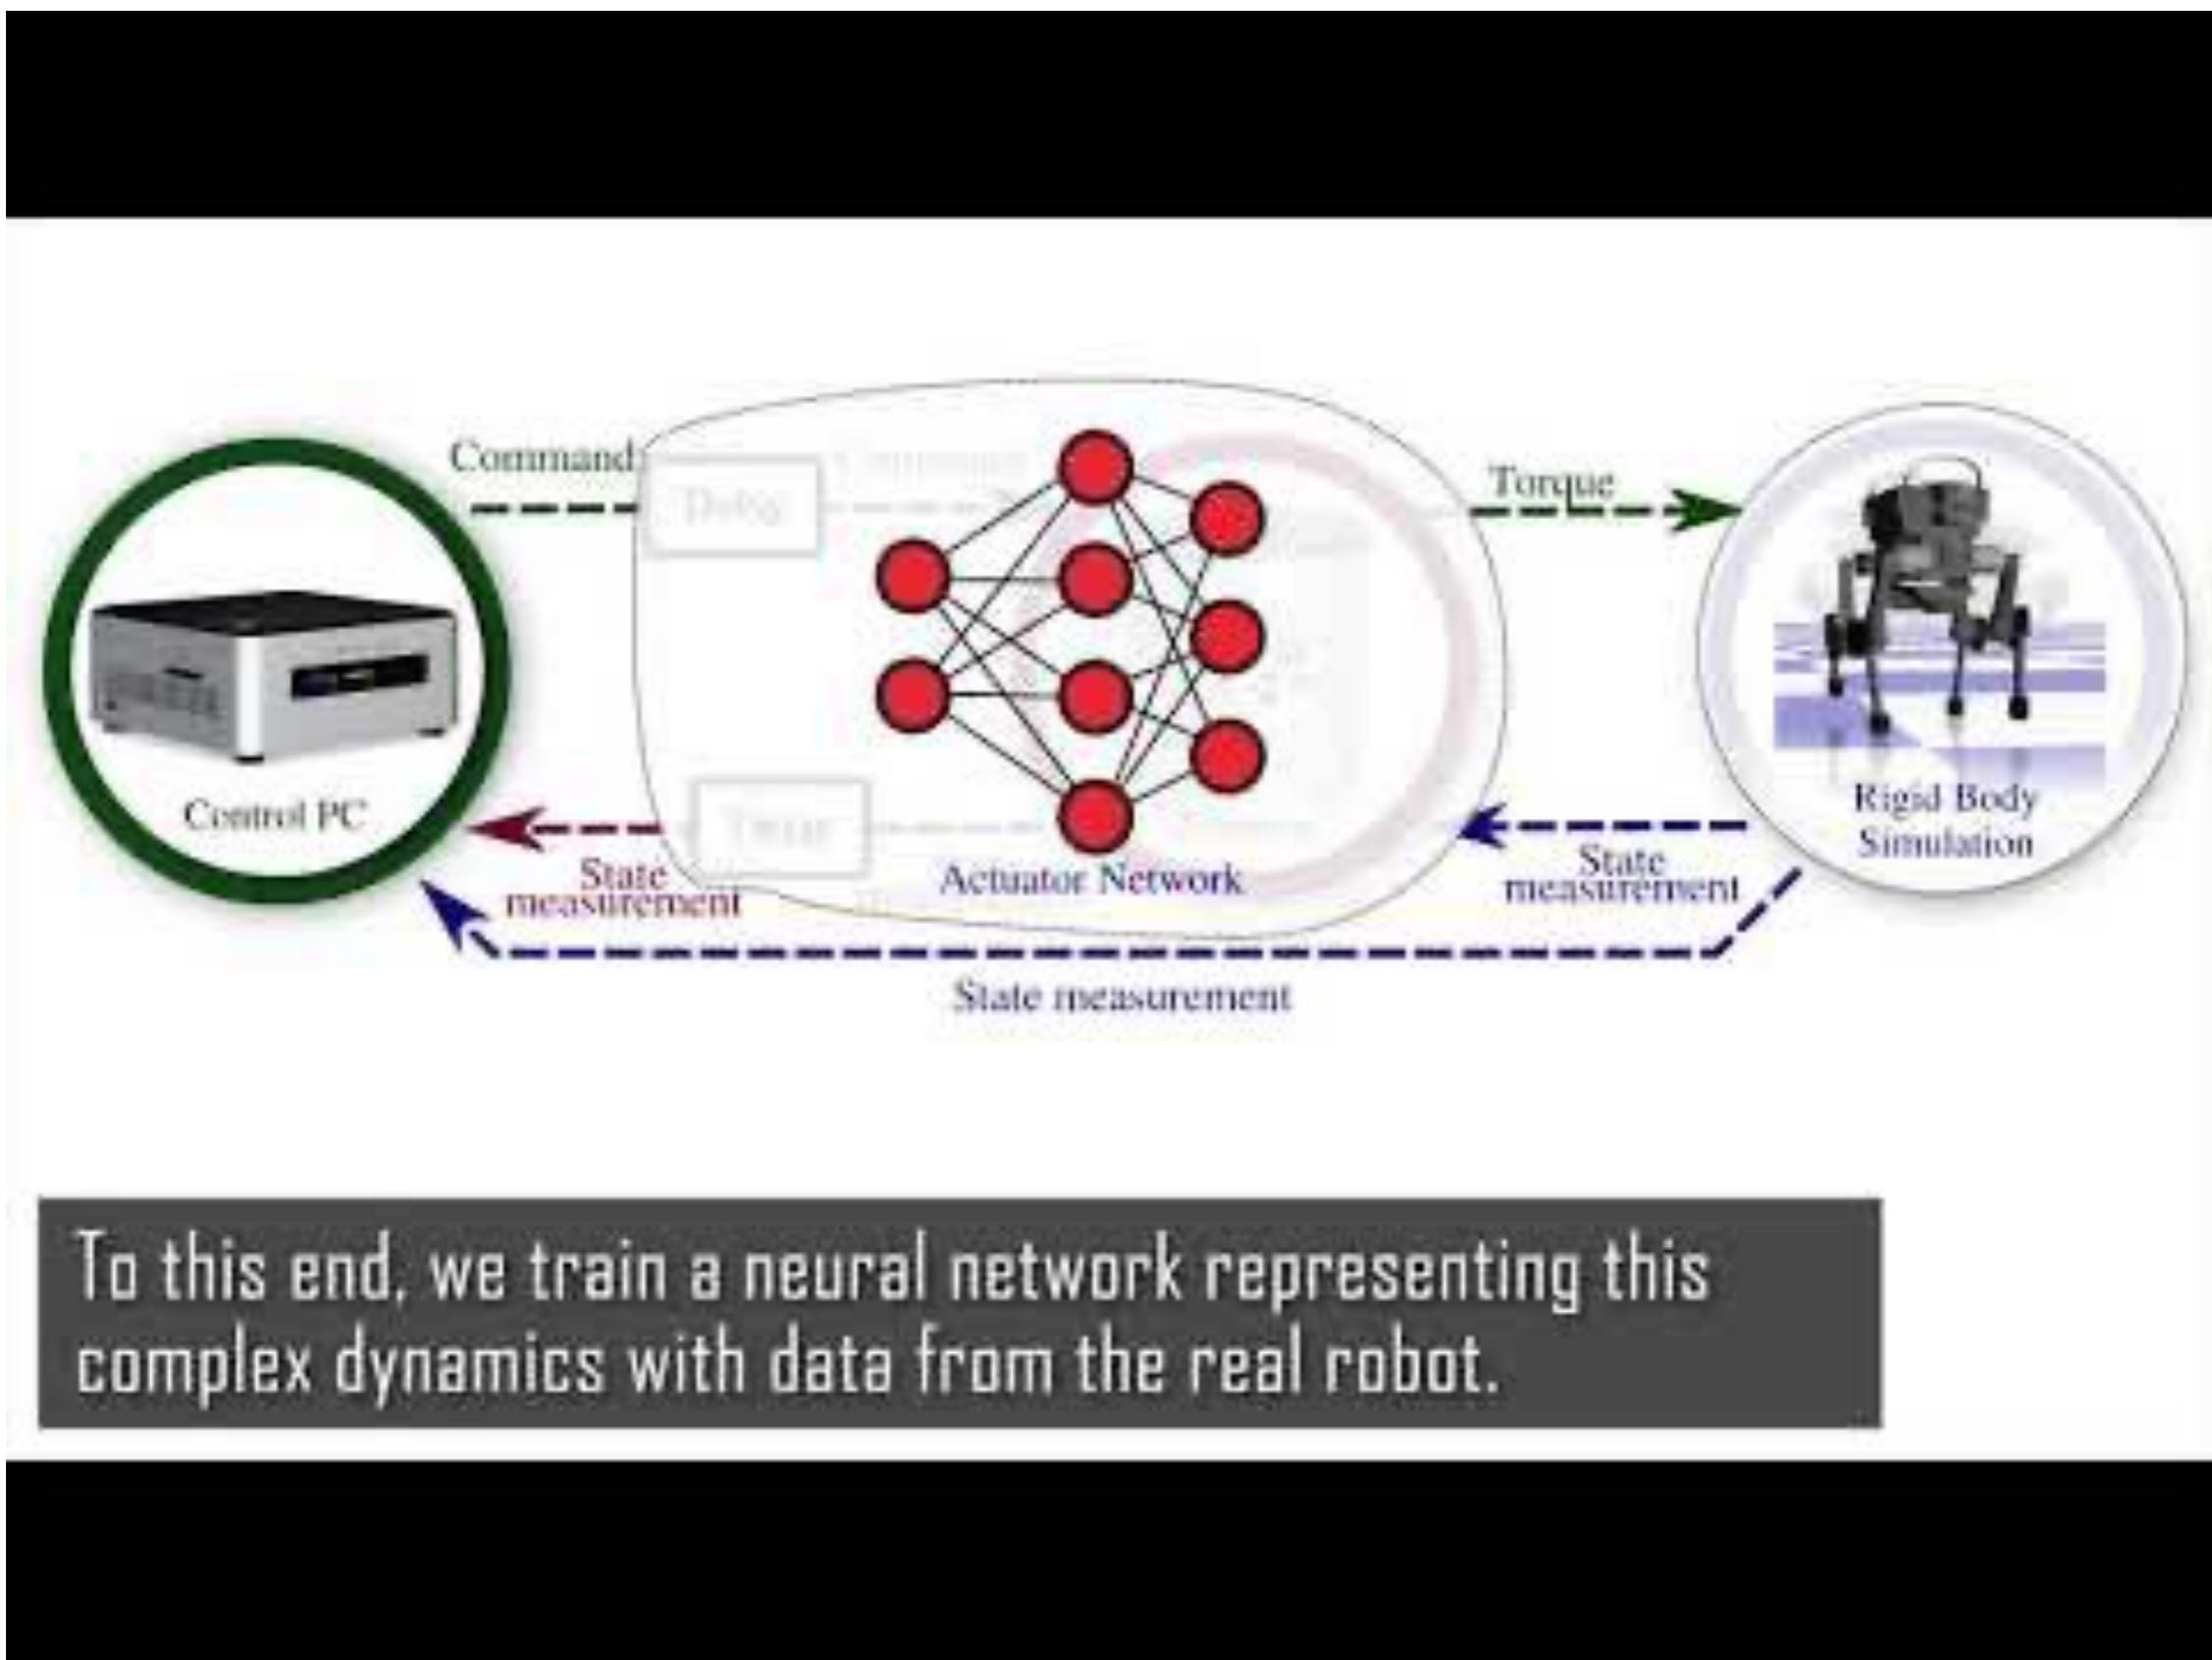
- **Drone Racing** E. Kaufmann et al. ICRA19 [pdf](#), [video](#)

- **Self-supervised Deep Reinforcement Learning with Generalized Computation Graphs for Robot Navigation** G. Kahn, A. Villaflor, B. Ding, P. Abbeel, S. Levine. ICRA18 [pdf](#), [video](#)

- **Learning to Fly by Crashing** D. Gandhi, L. Pinto, A. Gupta. IROS17 [pdf](#), [video](#)

Rewards Signals, Simulation
Imitation Learning
Modularized Learning
Ecological Supervision

- Learning Agile and Dynamic Motor Skills for Legged Robots J. Hwango et al. Science Robotics 18 [pdf](#), [video](#)



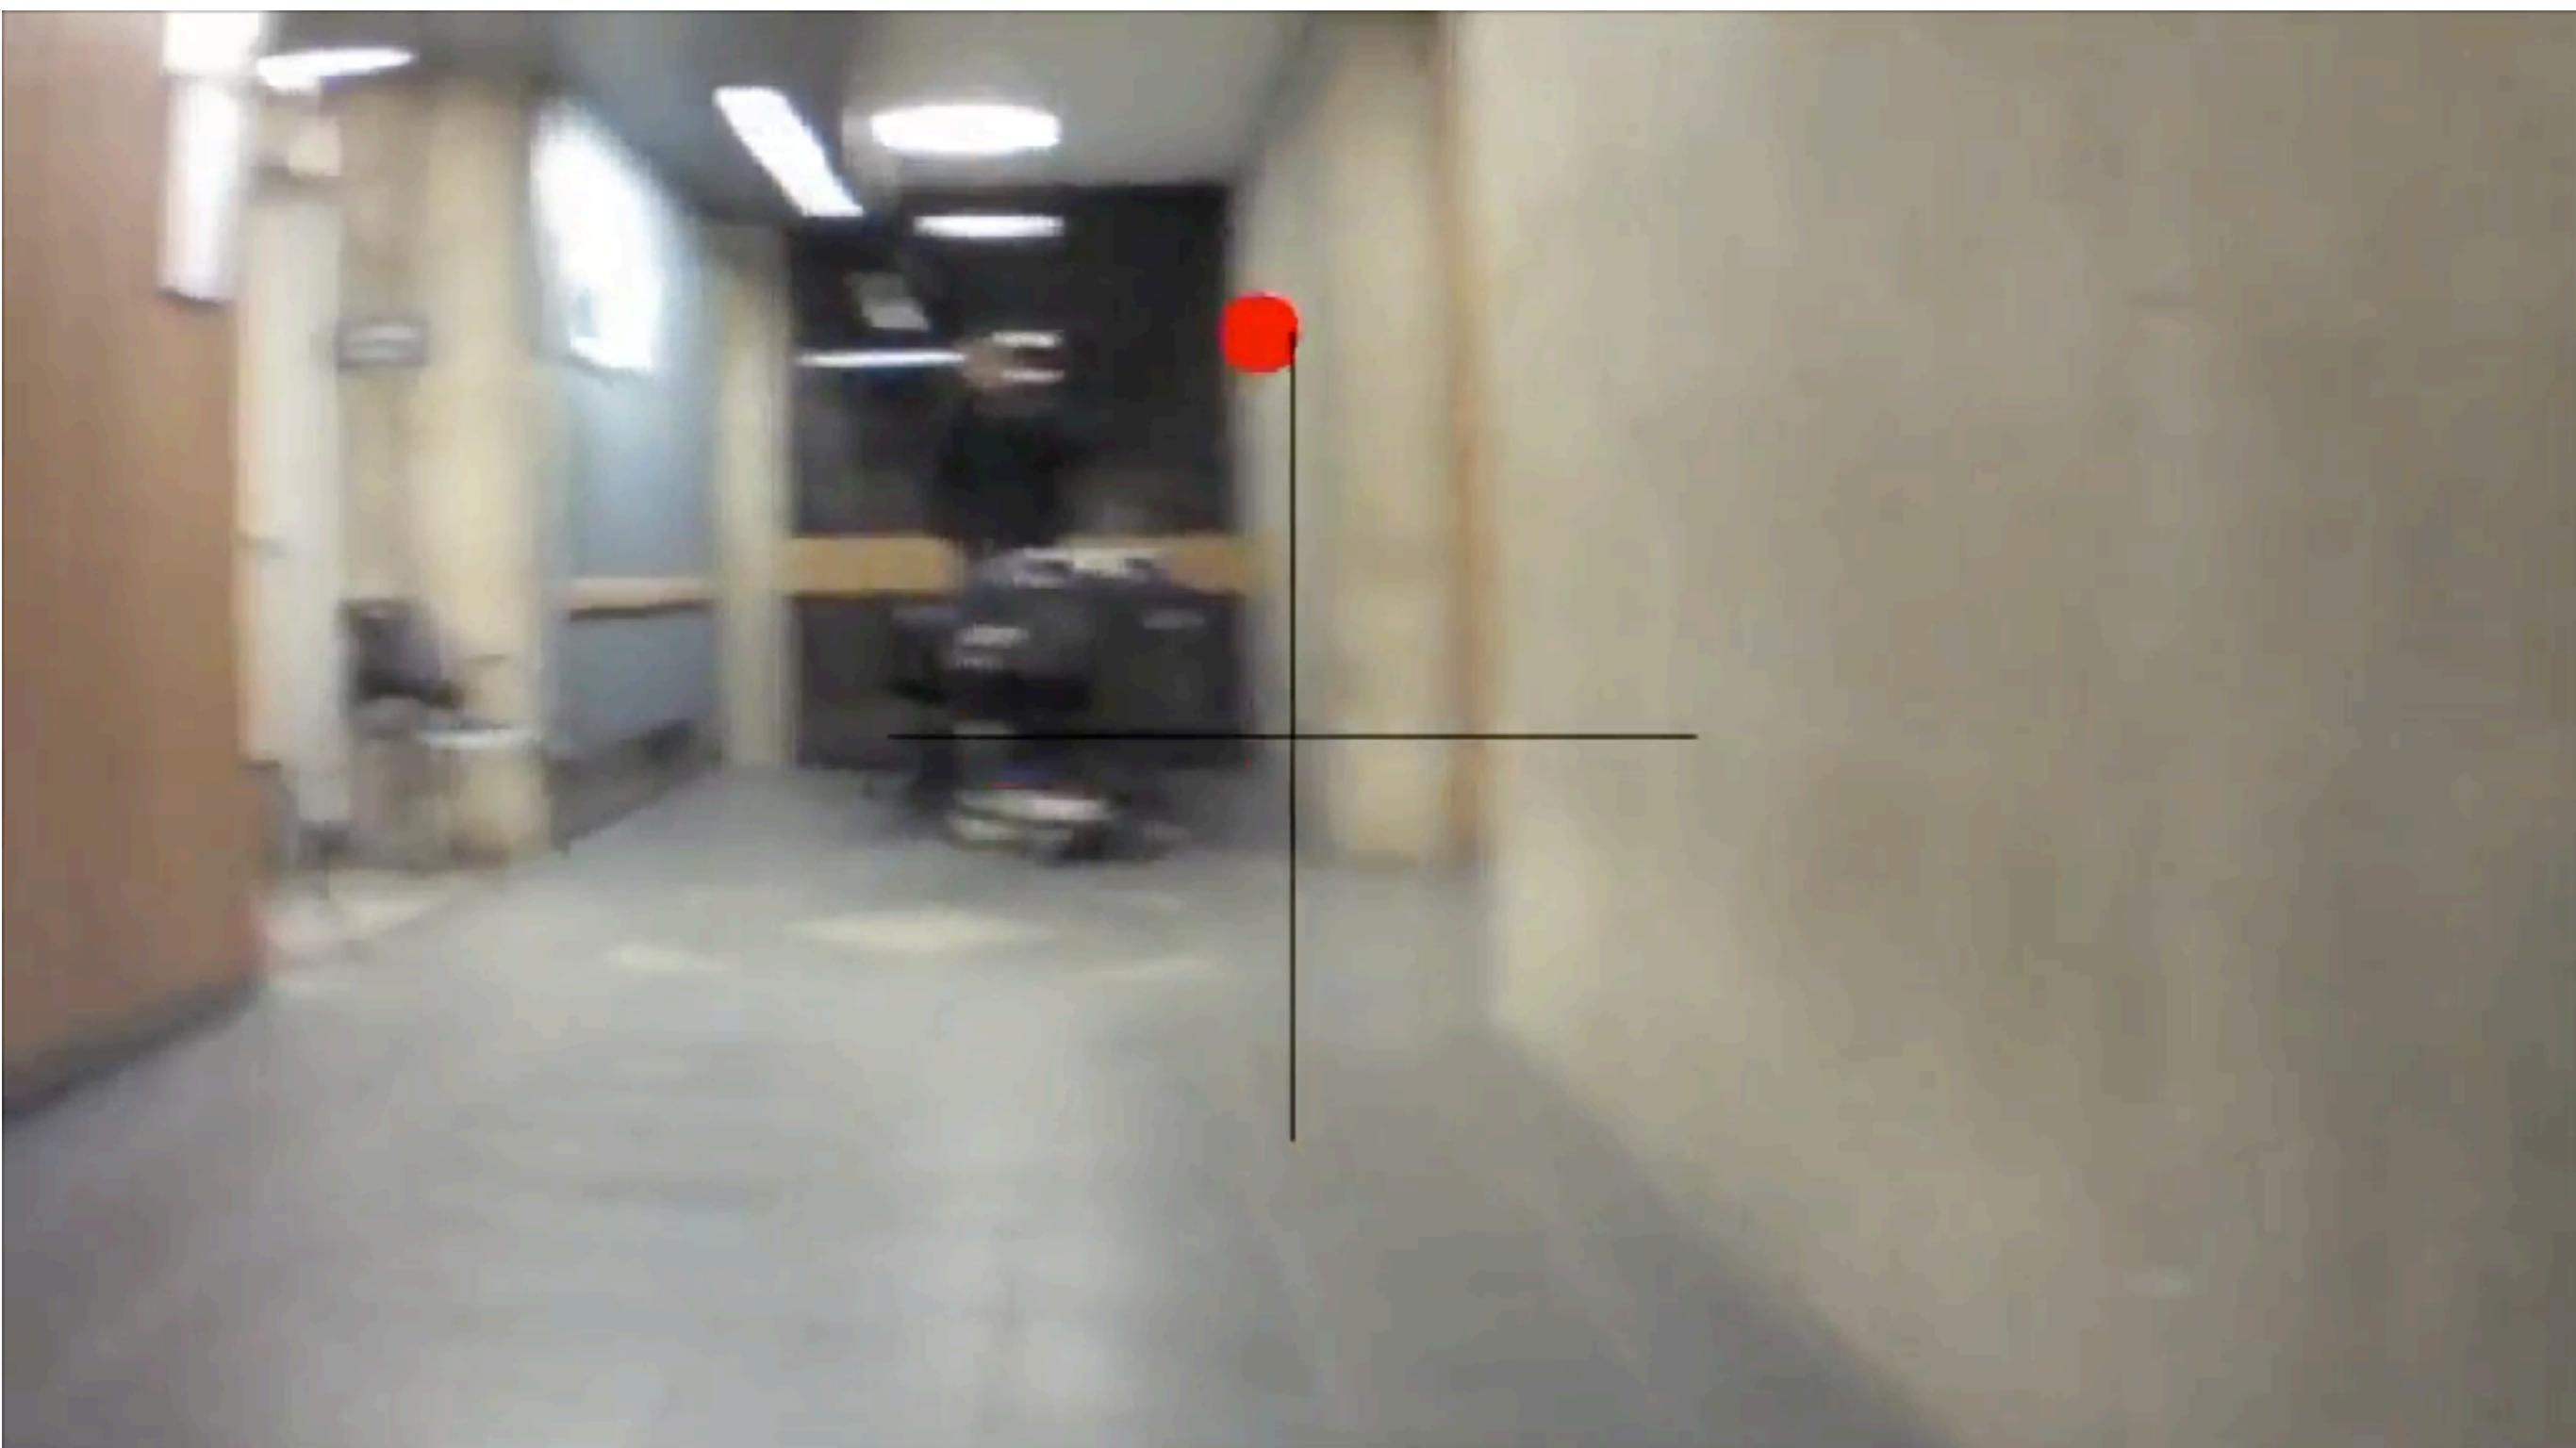
- **Agile Autonomous Driving** Y. Pan et al. R:SS18 [pdf](#), [video](#)



- **Beauty and the Beast: Optimal Methods Meet Learning for Drone Racing** E. Kaufmann et al. ICRA19 [pdf](#), [video](#)



Learning to Fly by Crashing D. Gandhi, L. Pinto, A. Gupta. IROS17. [pdf](#), [video](#)



**Self-supervised Deep Reinforcement Learning with Generalized Computation Graphs
for Robot Navigation** G. Kahn S. Levine. ICRA18. [pdf](#), [video](#)

Self-supervised Deep Reinforcement Learning with Generalized Computation Graphs for Robot Navigation



Gregory Kahn, Adam Villaflor, Bosen Ding, Pieter Abbeel, Sergey Levine



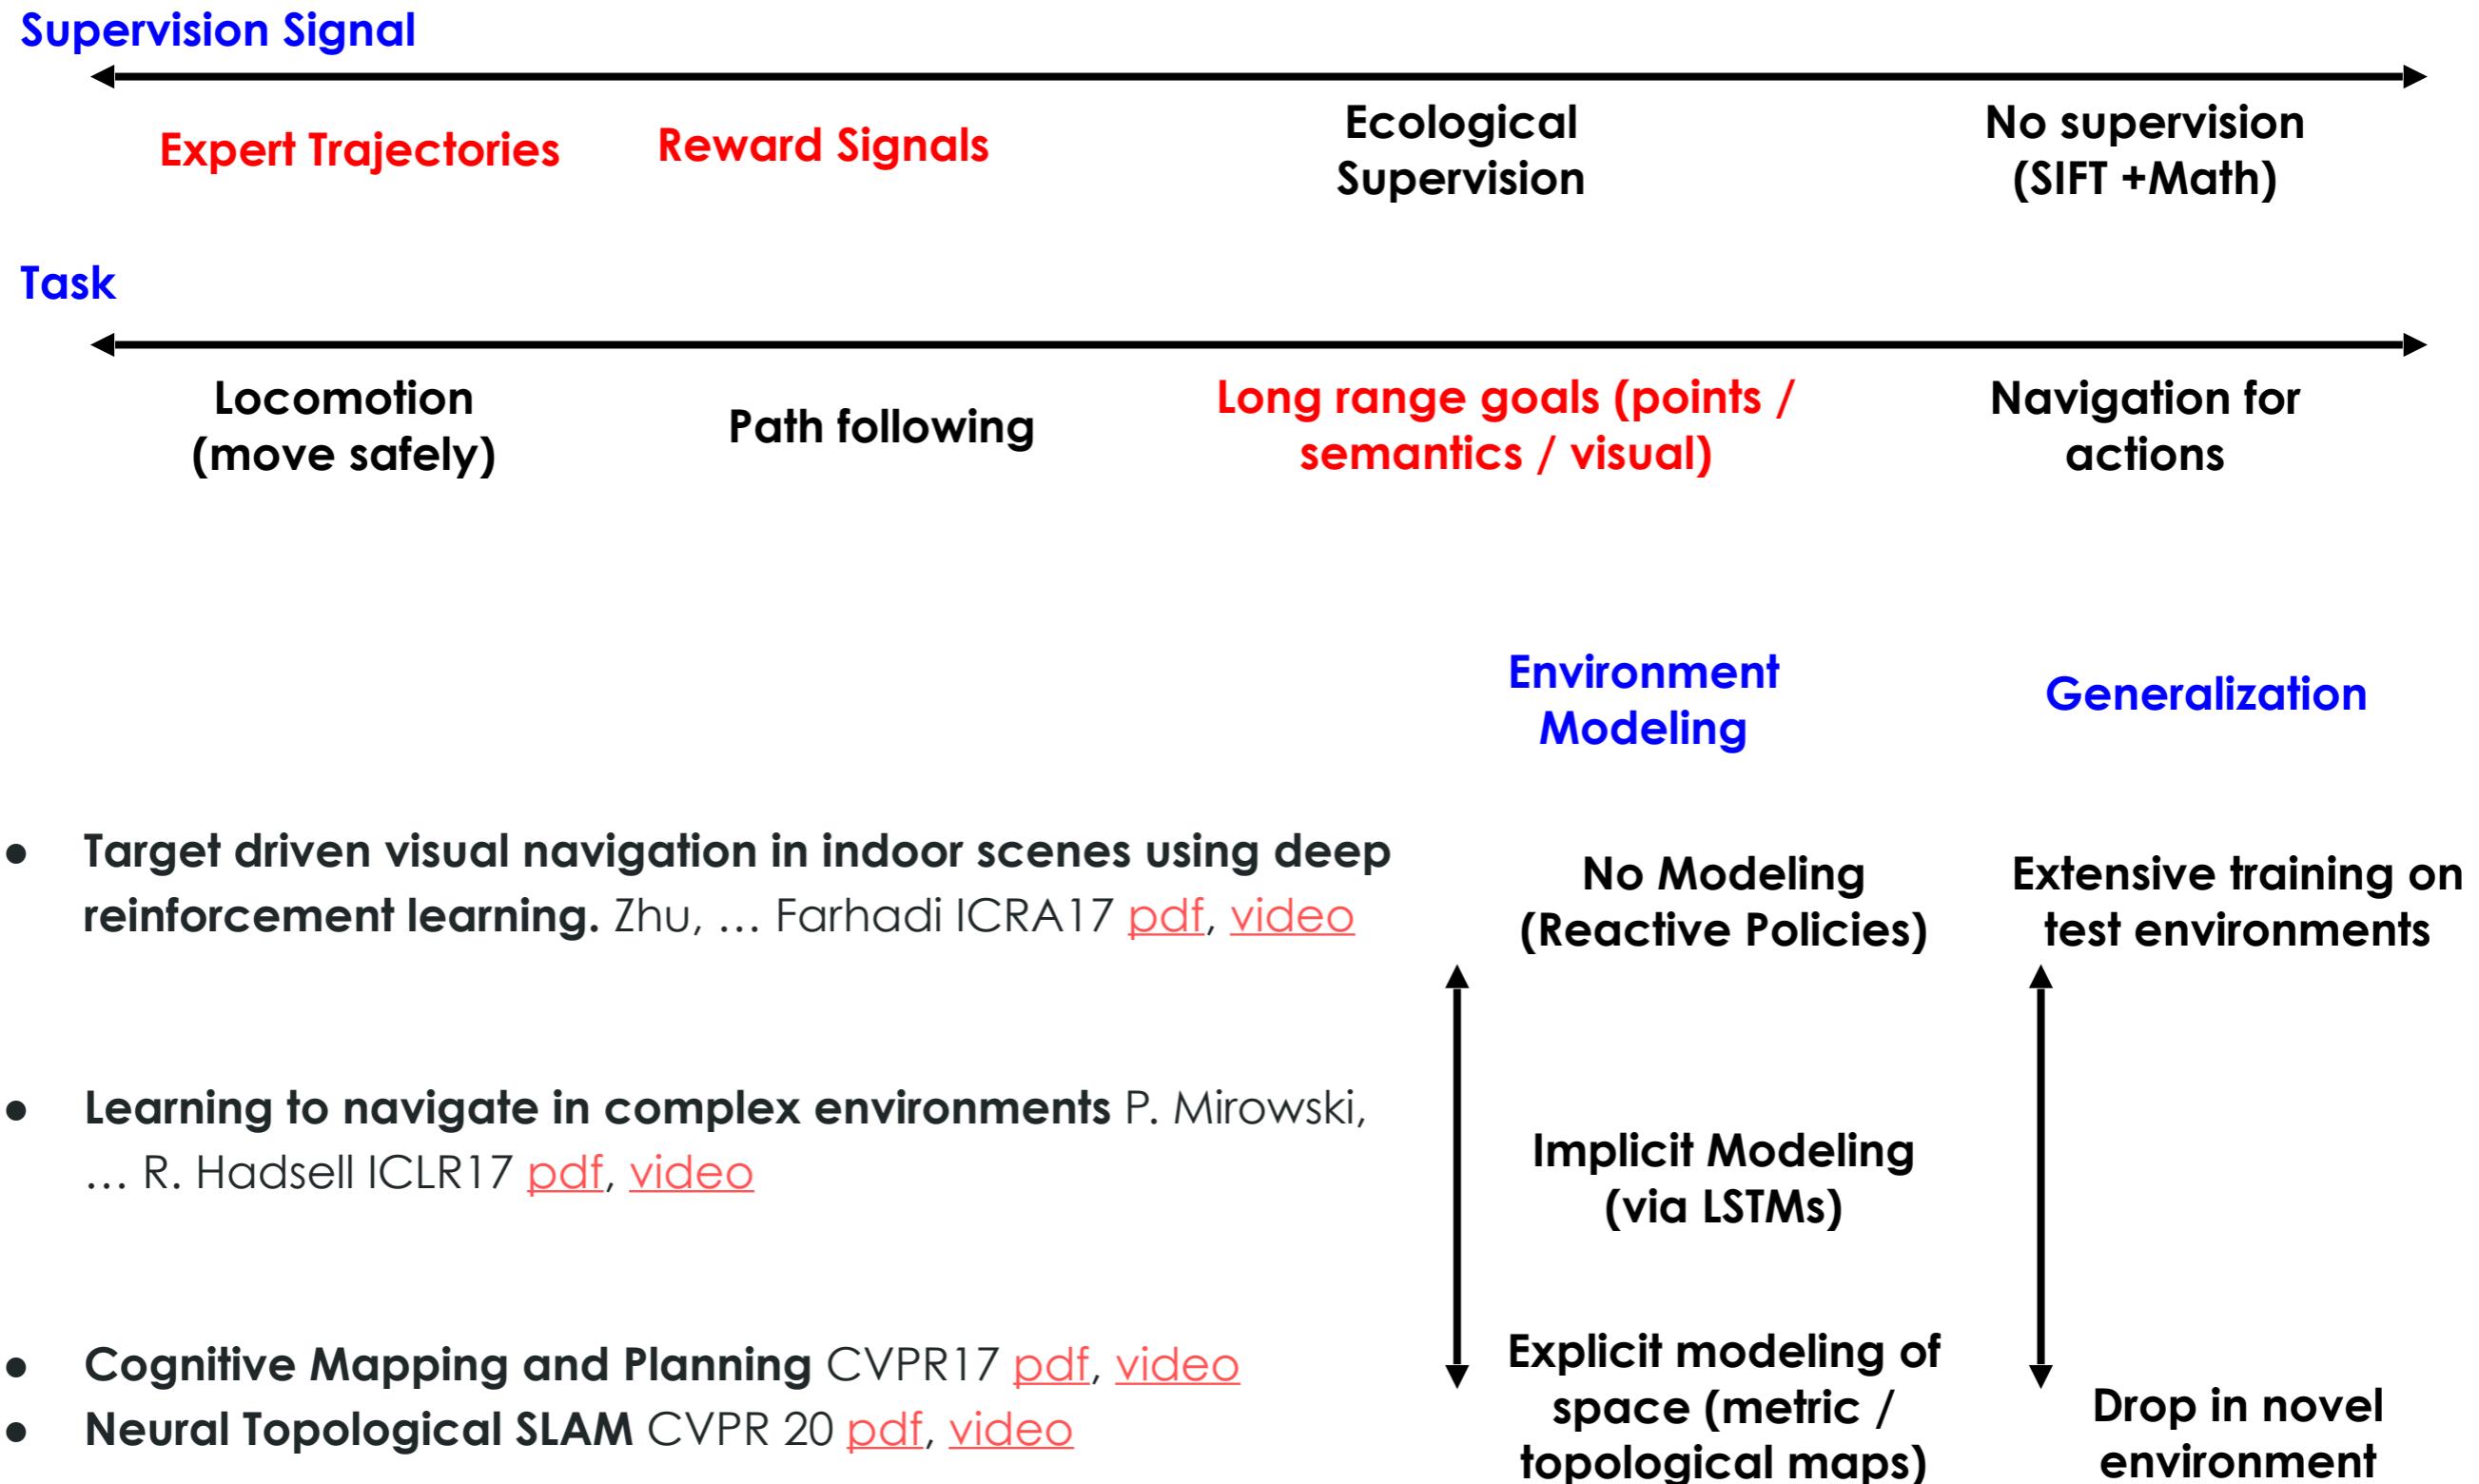
Goal directed navigation

- Goal-directed behavior
 - Goals known / unknown
 - Environment seen / novel
- Central question is that of how to represent space?
 - Reason about what you have seen and where / what you have not seen
 - Spatial reasoning
 - Robustness to actuation noise
- For unknown environments, we may additionally need:
 - Semantic reasoning
 - Easy to build

Representations for Space

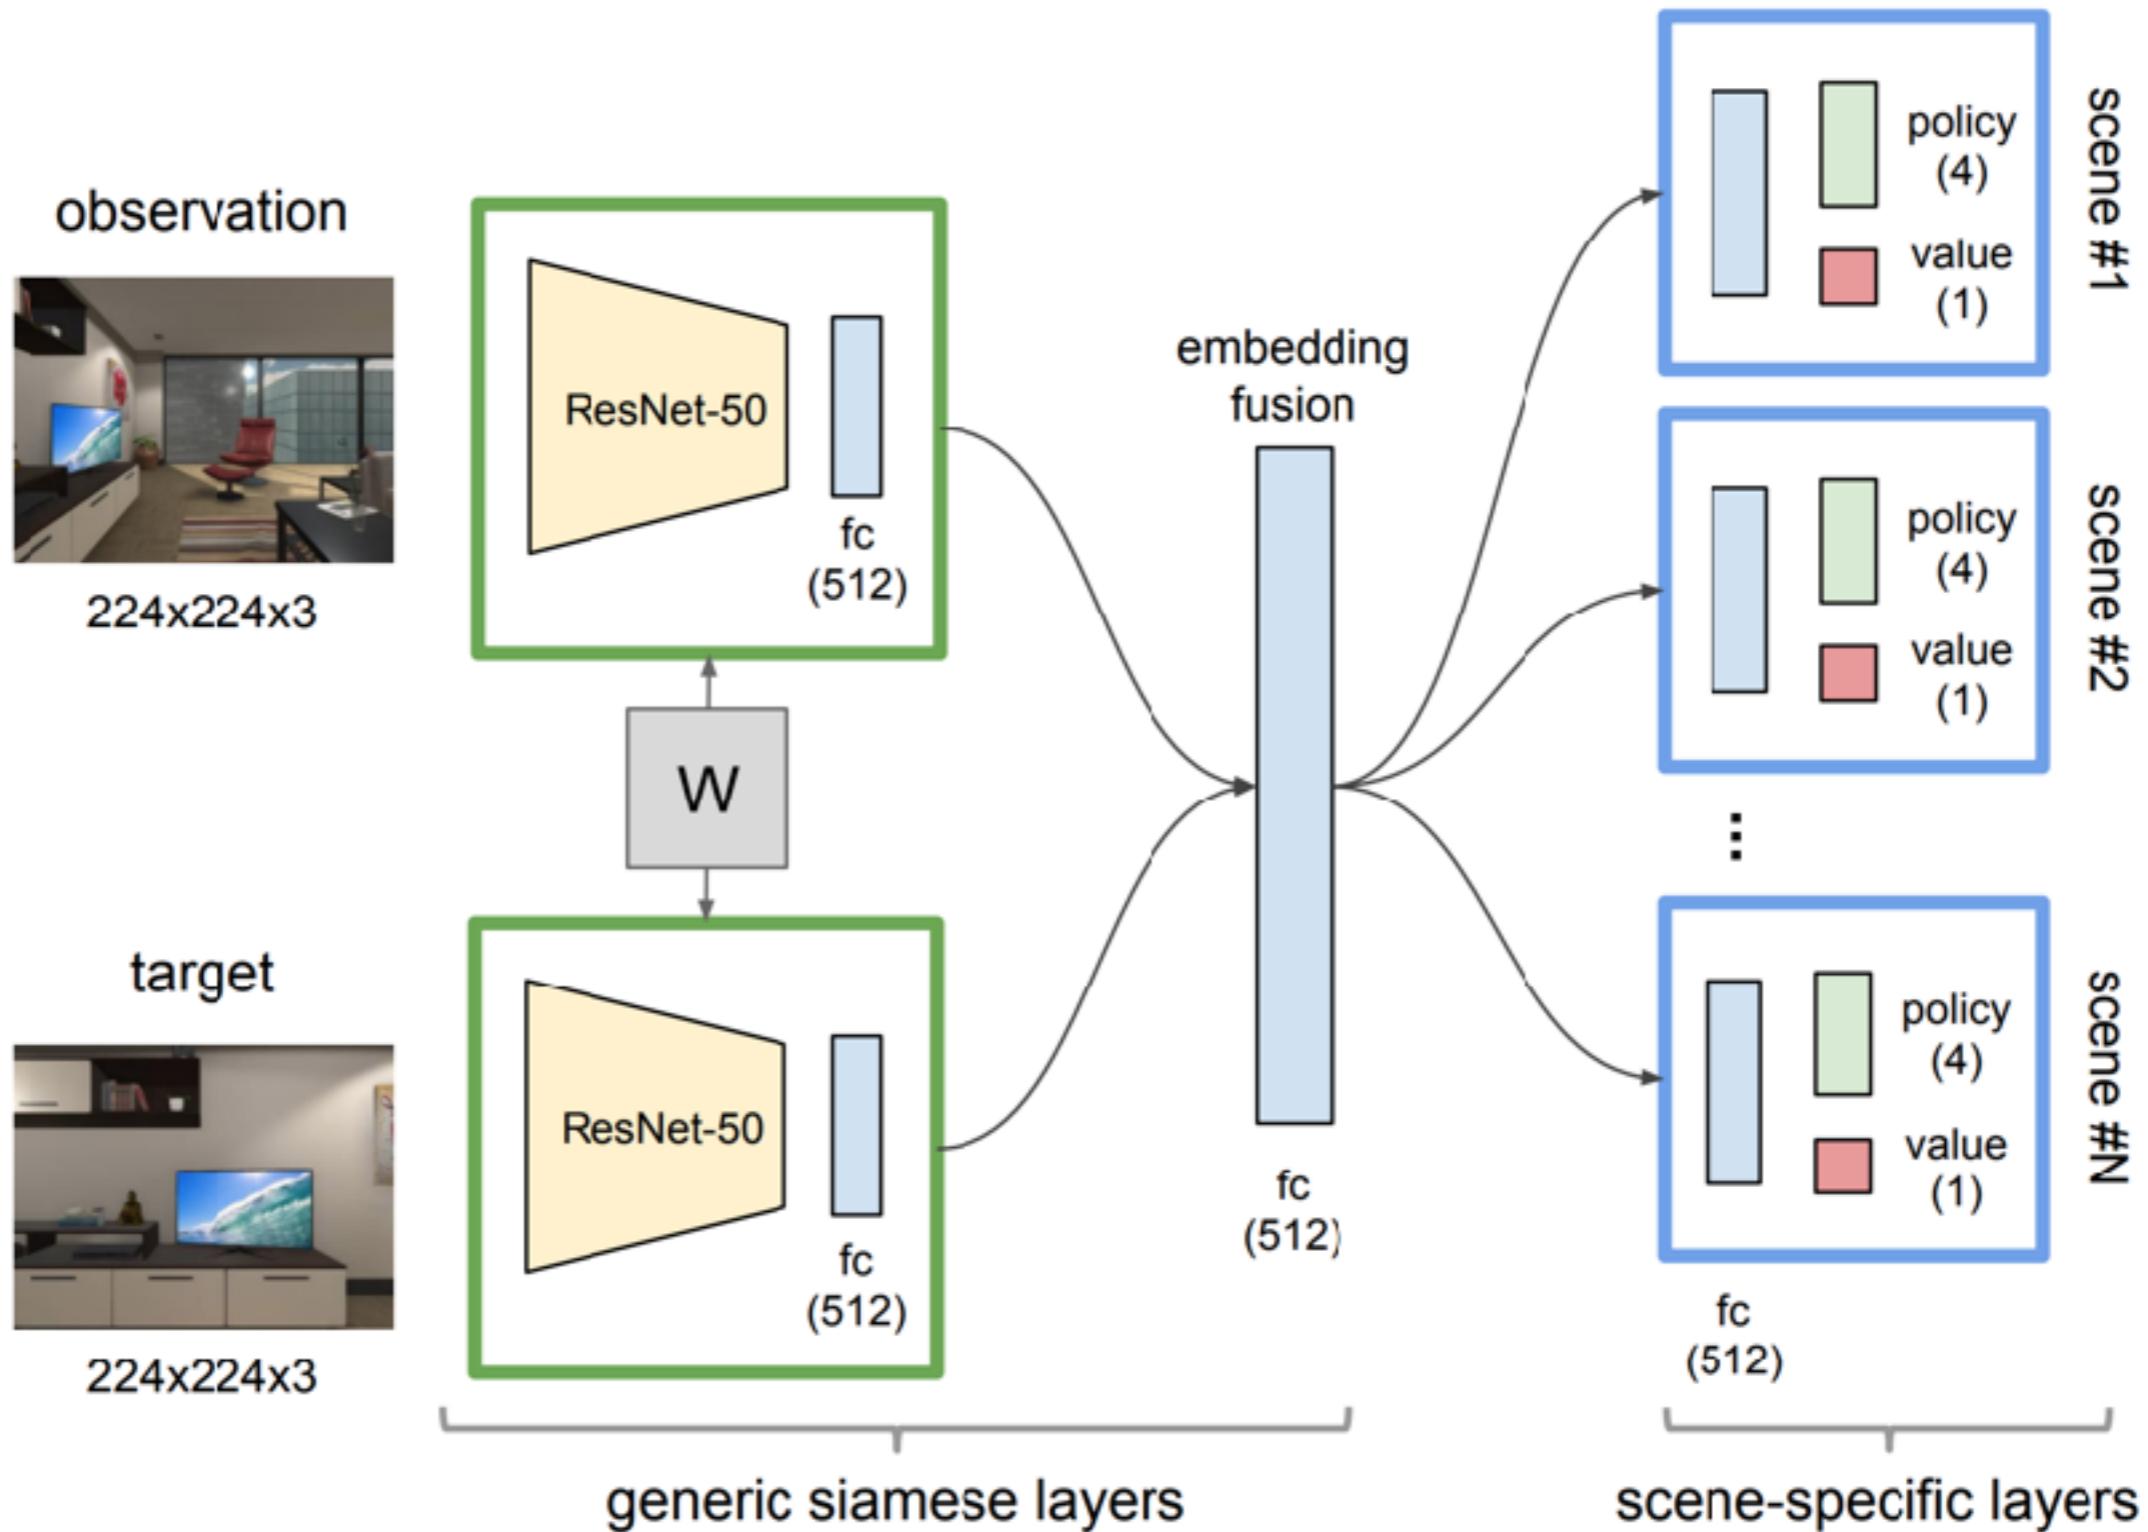
- Metric maps
 - Occupancy maps
 - Occupancy maps flavored with semantic information
- Topological maps
 - Collection of images, and connectivity information

Learned Task driven End-to-end navigation



Target driven visual navigation in indoor scenes using deep reinforcement learning.

Y. Zhu, ..., A. Gupta, ..., A. Farhadi. ICRA17. [pdf](#), [video](#)



Simulation Environments

Abstraction



State space discretization
(eg. grid world)

Discrete Actions

Continuous actions

Environment Properties



Random Mazes (limited semantics)

Real world environments

Experimental setup



Simulation

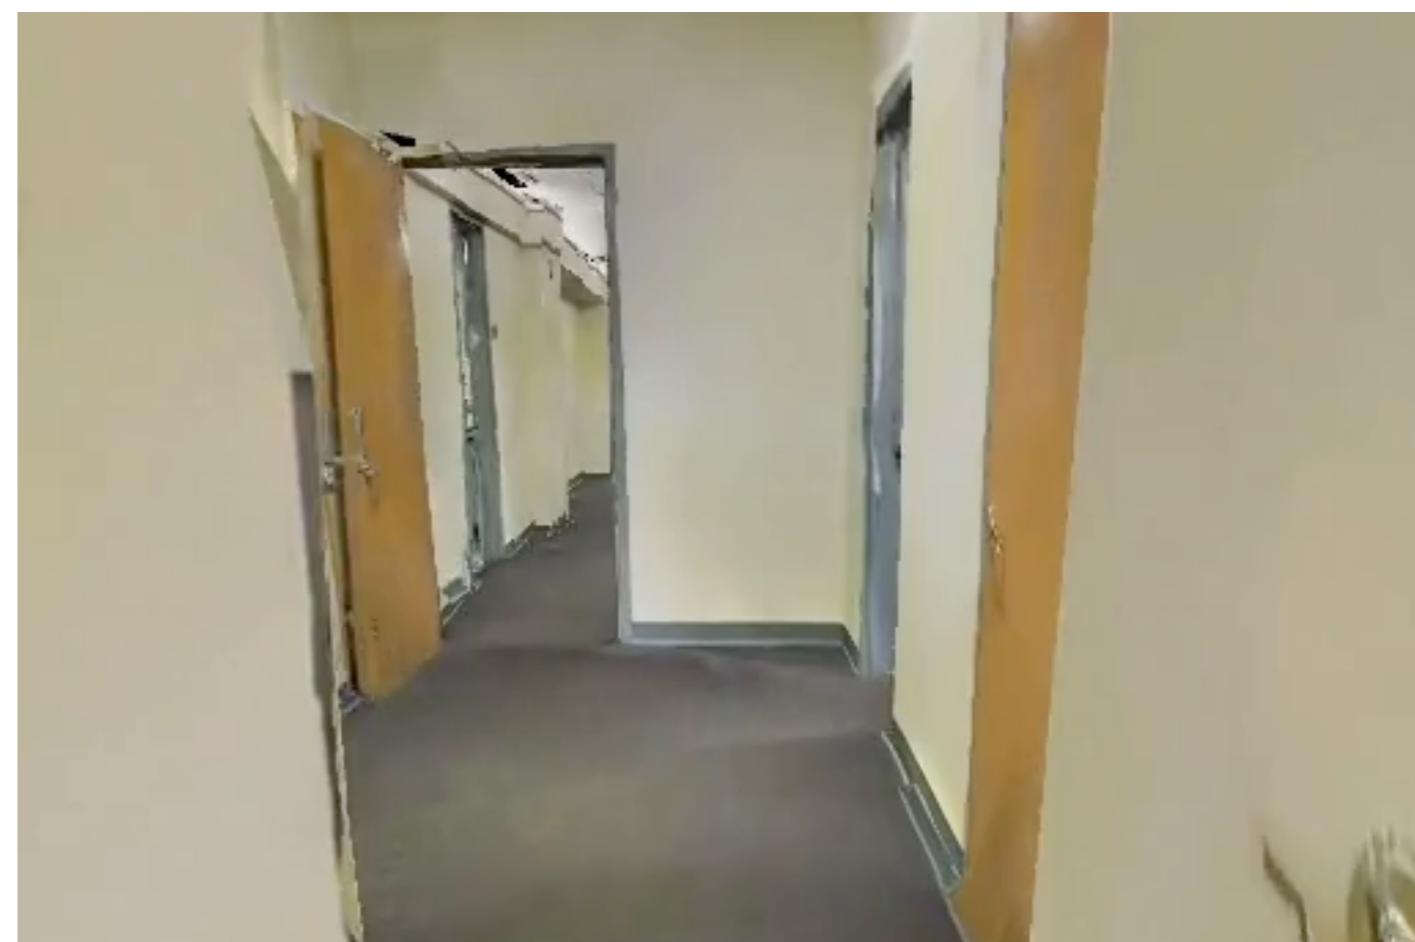
Simulation (based on real world environments)

Real world

Simulators based off games / CG (eg:
AI2Thor, VizDoom, CARLA)

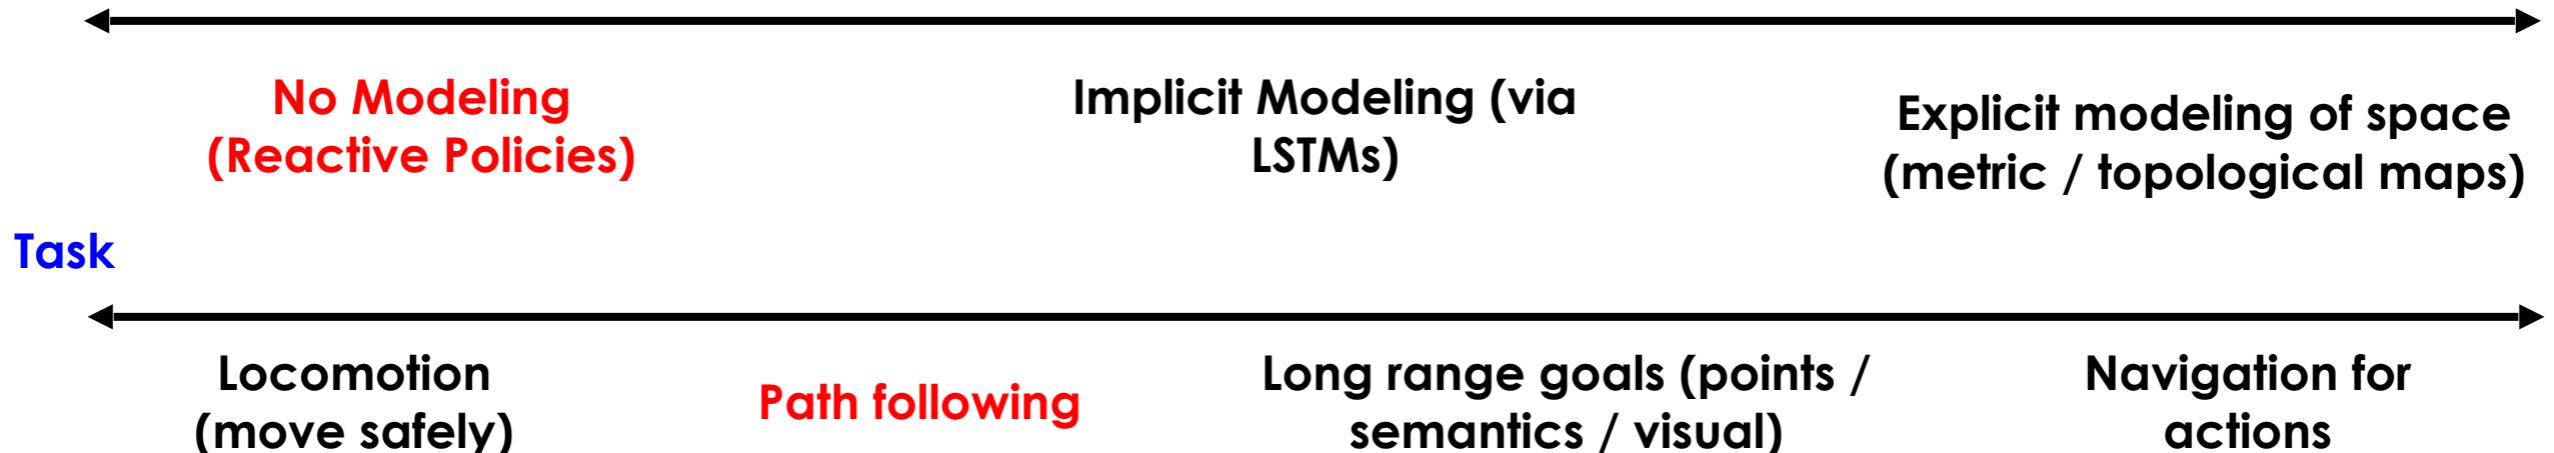
Simulators based off real world scans (eg:
CMPSim, Habitat, Gibson)

Model trained on 10M frames: Go to Sofa



Visual Servoing

Environment Modeling



Supervision Signal

- **Visual Memory for Robust Path Following**
Kumar, Gupta, ... Malik
- **End-to-end Driving via Conditional Imitation Learning** F. Codevilla, ... V. Koltun ICRA18? [pdf](#), [video](#)
- **Learning Visual Servoing with Deep Features and Fitted Q-Iteration**, Alex X. Lee, Sergey Levine, Pieter Abbeel. ICLR17. [pdf](#)
- **Zero-Shot Visual Imitation** D. Pathak, ..., T. Darrell pdf, video
ICLR18? [pdf](#), [video](#)
- **Classical Visual Servoing**

↑
Expert Trajectories
Reward Signals
Ecological Supervision
No supervision (SIFT + Geometry)
↓

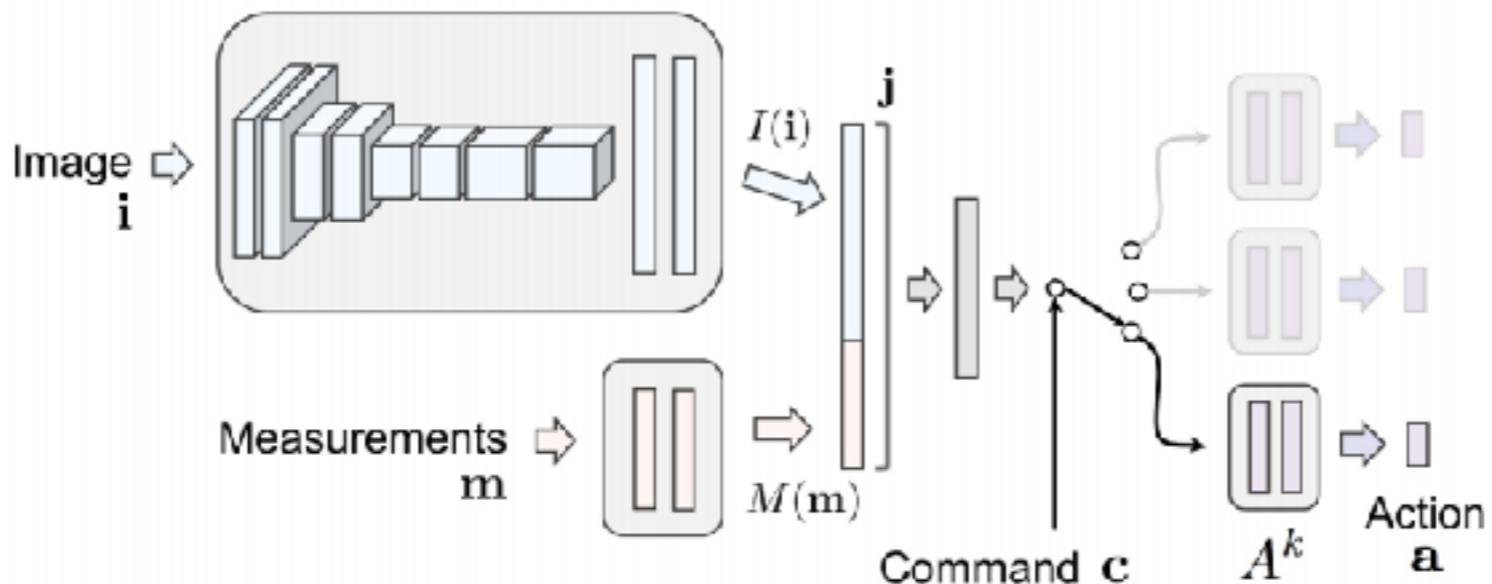
Learning Visual Servoing with Deep Features and Fitted Q-Iteration

Alex X. Lee, Sergey Levine, Pieter Abbeel. ICLR17. [pdf](#)



End-to-end Driving via Conditional Imitation Learning

F. Codevilla, ... V. Koltun ICRA18? [pdf](#), [video](#)

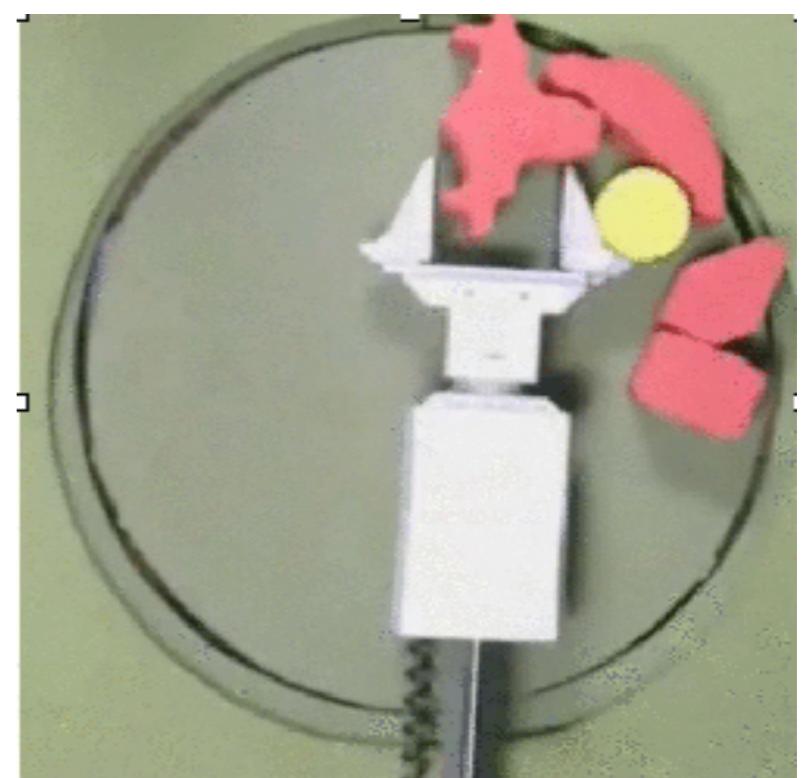
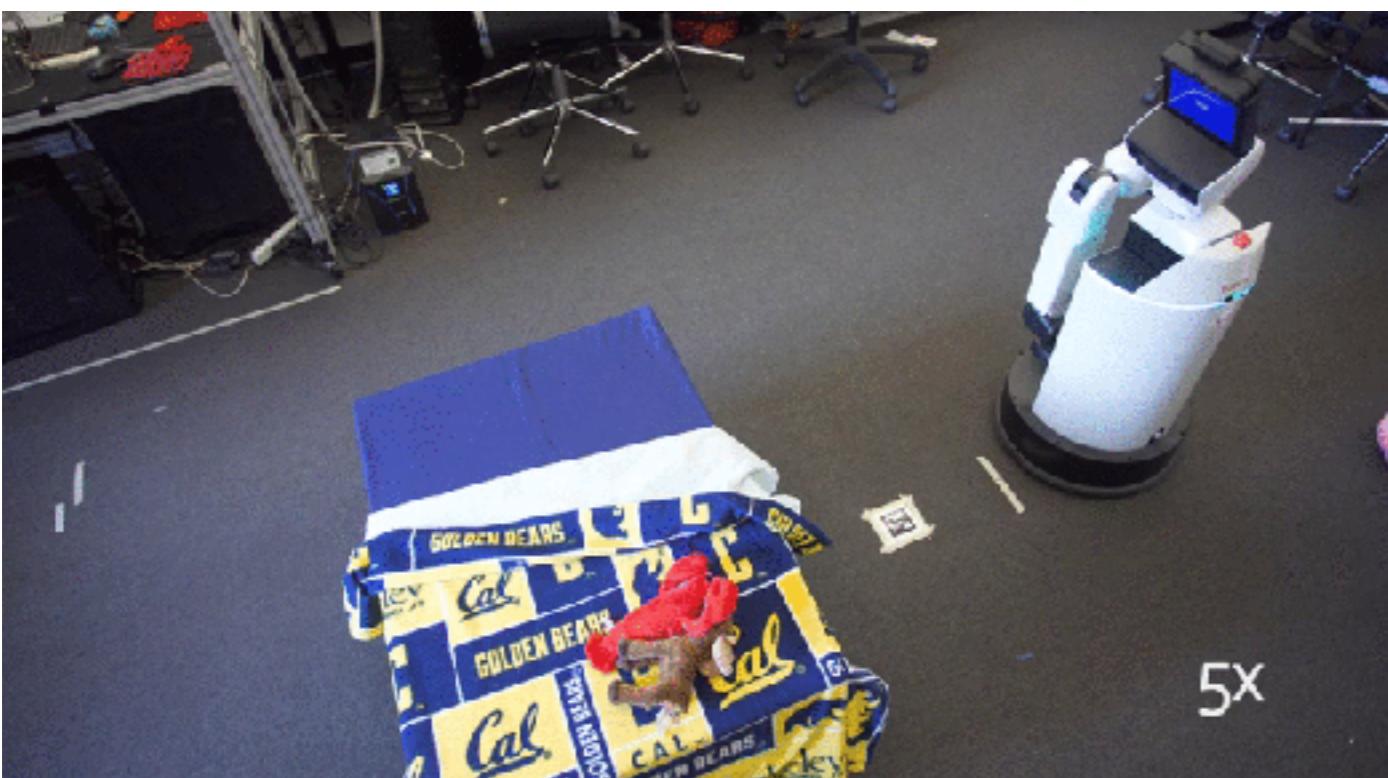
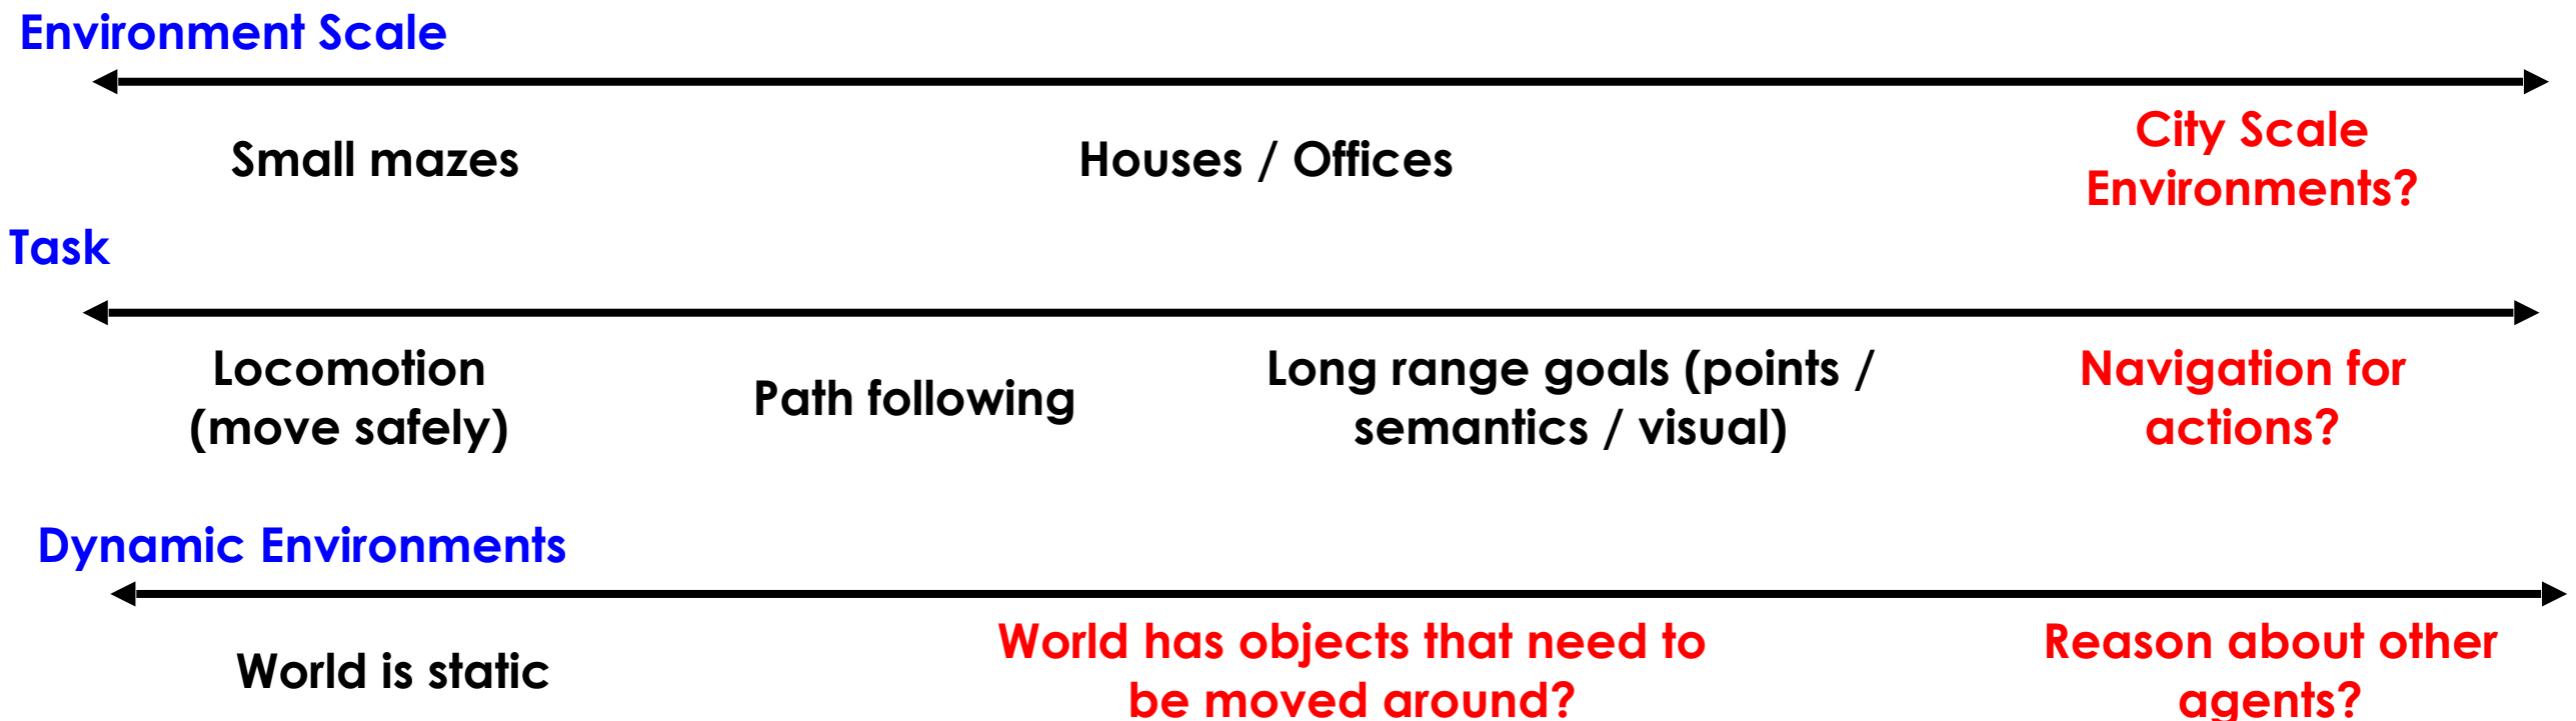


End-to-end Driving via Conditional Imitation Learning

Felipe Codevilla, Matthias Mueller, Alexey Dosovitskiy, Antonio Lopez, Vladlen Koltun

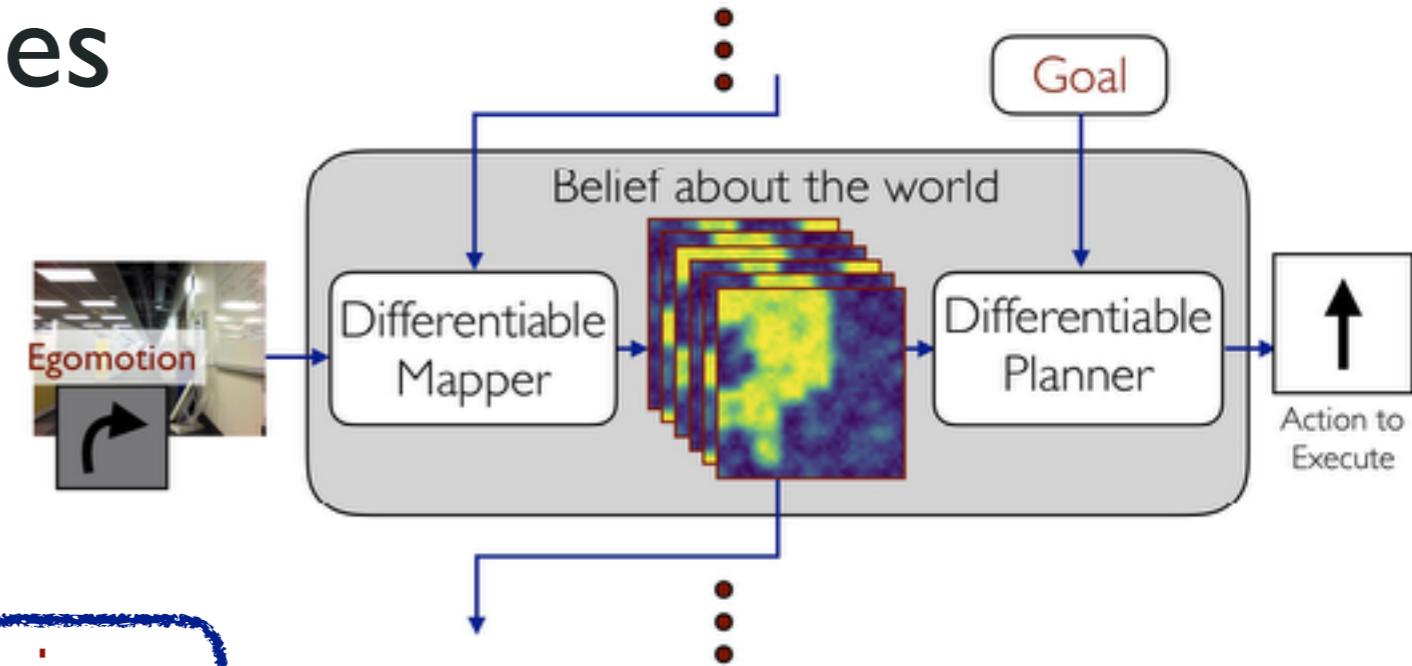
Submitted to ICRA 2018

Other Unexplored Axes?



Representation for Spaces

- Semantic reasoning
- Spatial reasoning
- Robustness to actuation noise
- Easy to acquire



- Topological representations (robustness to noise)
- Ease of training
 - Modular approaches
 - Training via supervised learning

Semantic Topological Representations

Panorama

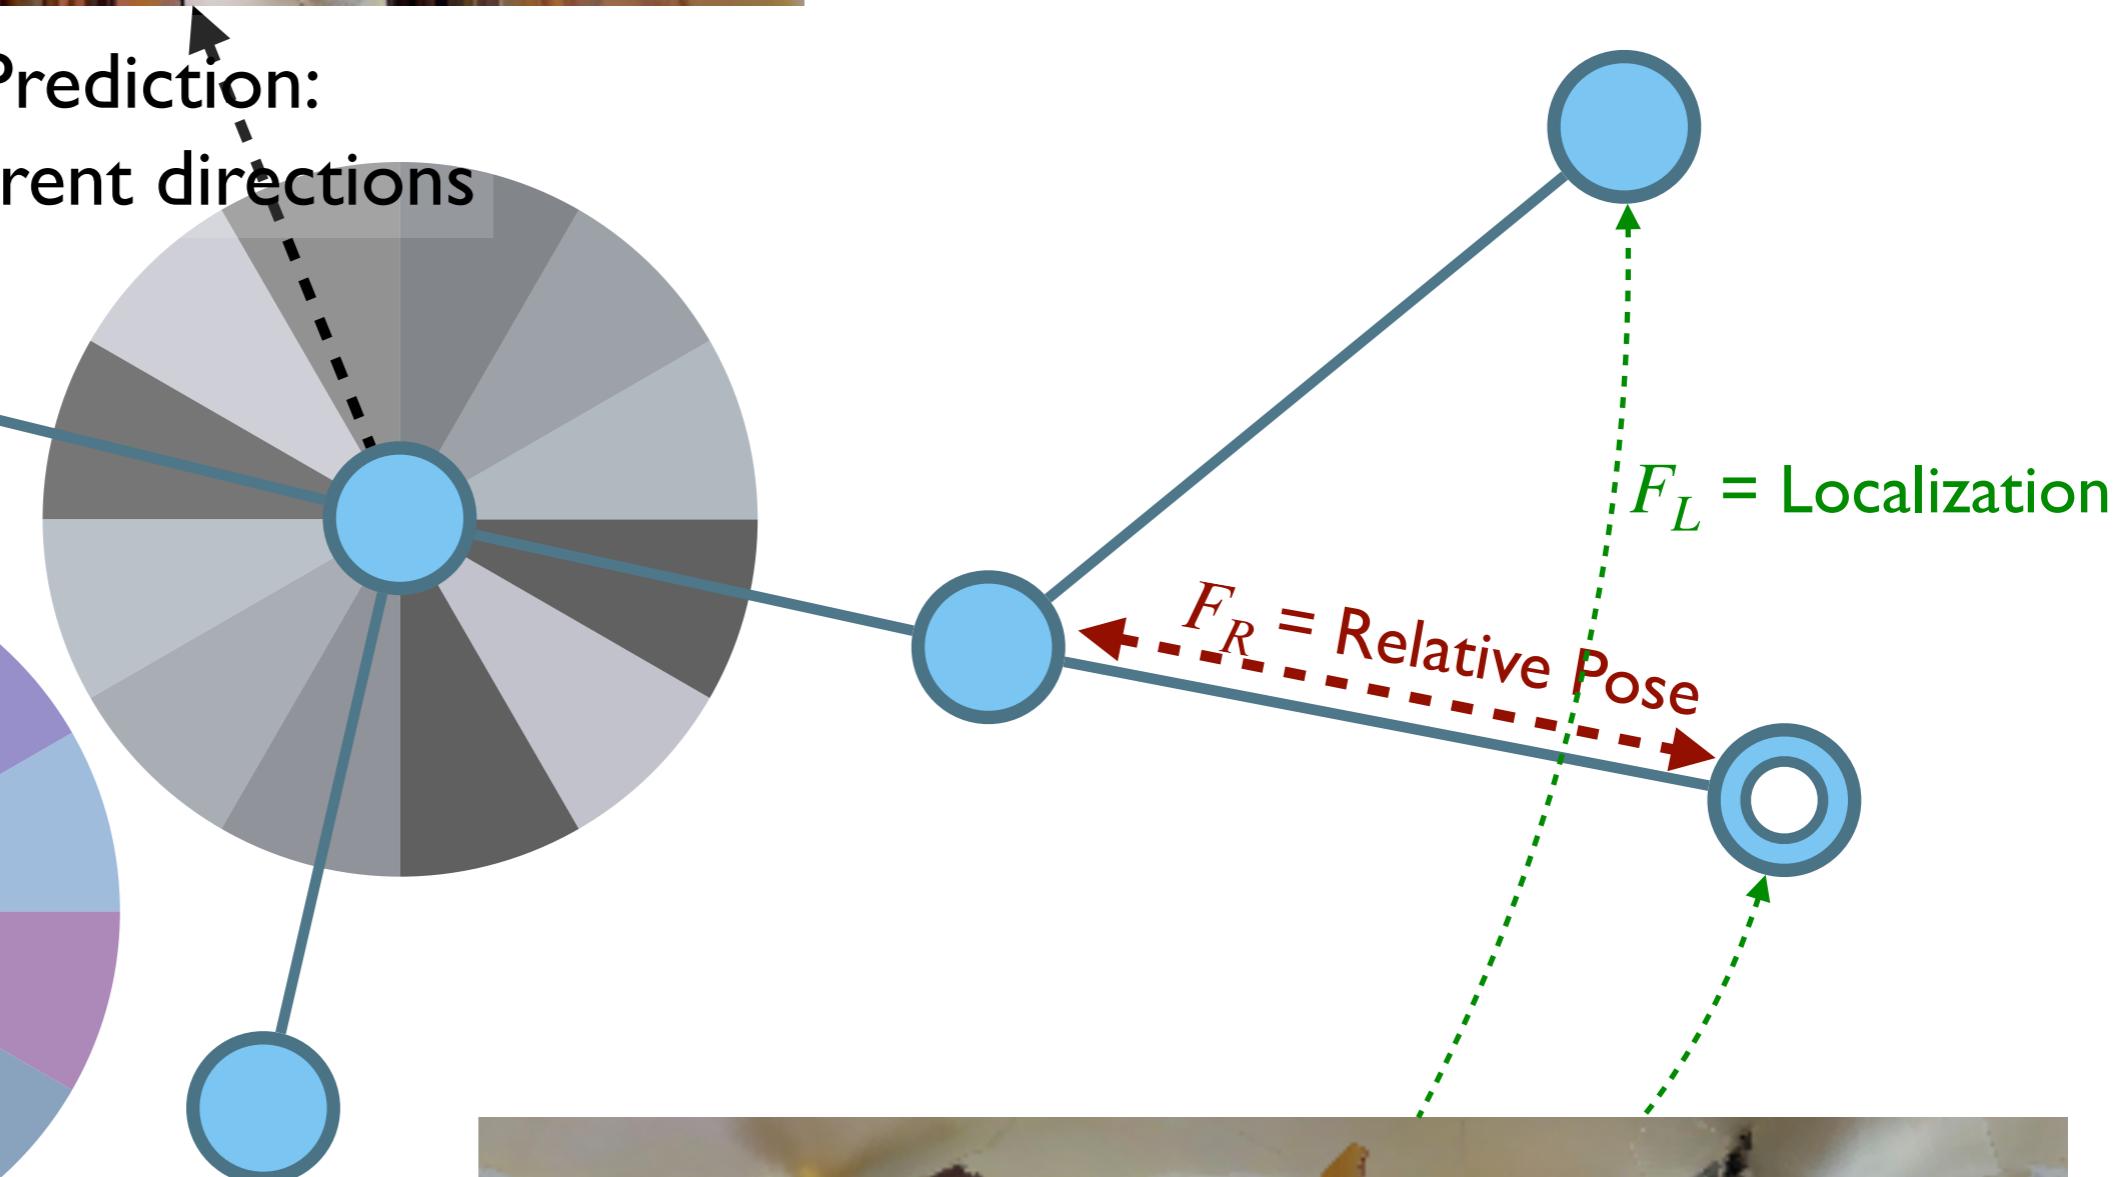
F_G = Geometric Prediction:
Free space in different directions

“Ghost
Nodes” 0.8



F_S = Semantic Prediction:
Closeness to target

D. S. Chaplot et al. Neural topological slam for visual navigation. In CVPR, 2020.



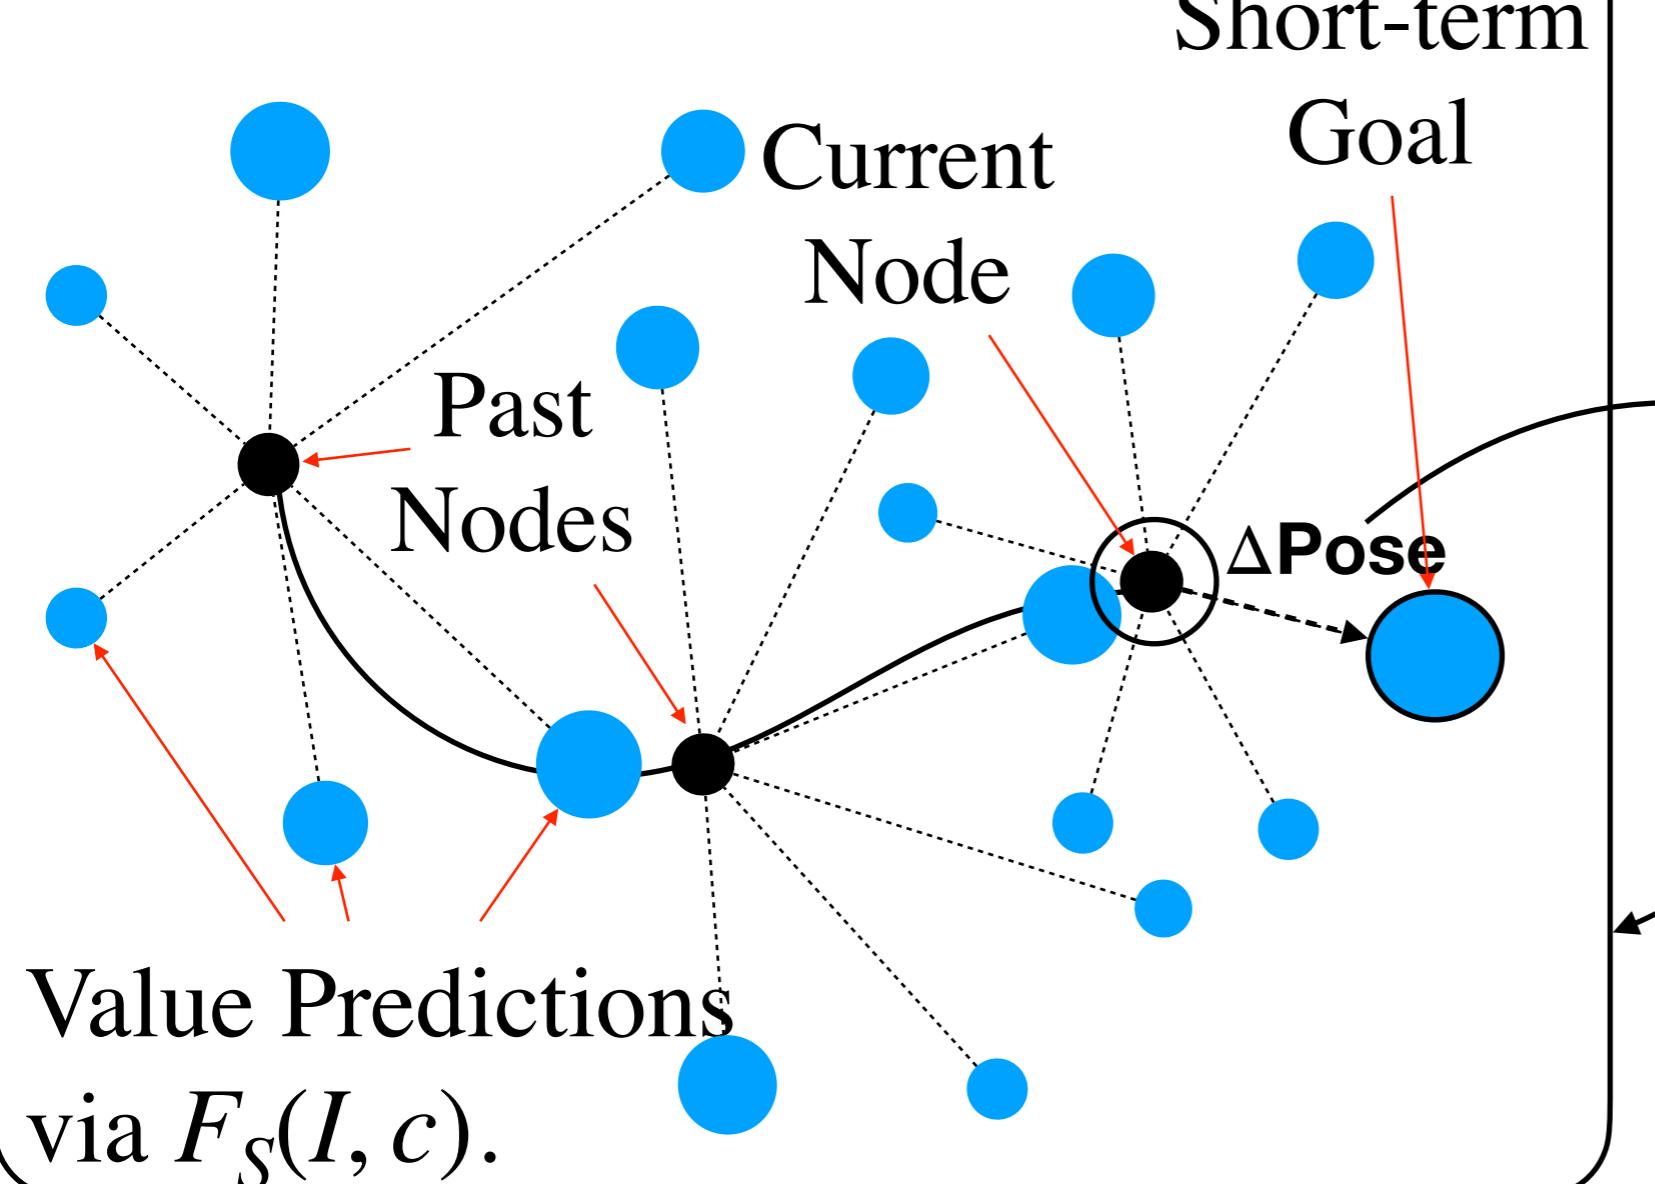
4 Functions

- $F_S(I_1, I_2)$ = Semantic Prediction: Closeness to target
- $F_G(I_1)$ = Geometric Prediction: Free directions
- $F_R(I_1, I_2)$ = Relative Pose
- $F_L(I_1, I_2)$ = Localization

Goal-driven Navigation

Hierarchical Policy

High-Level Policy



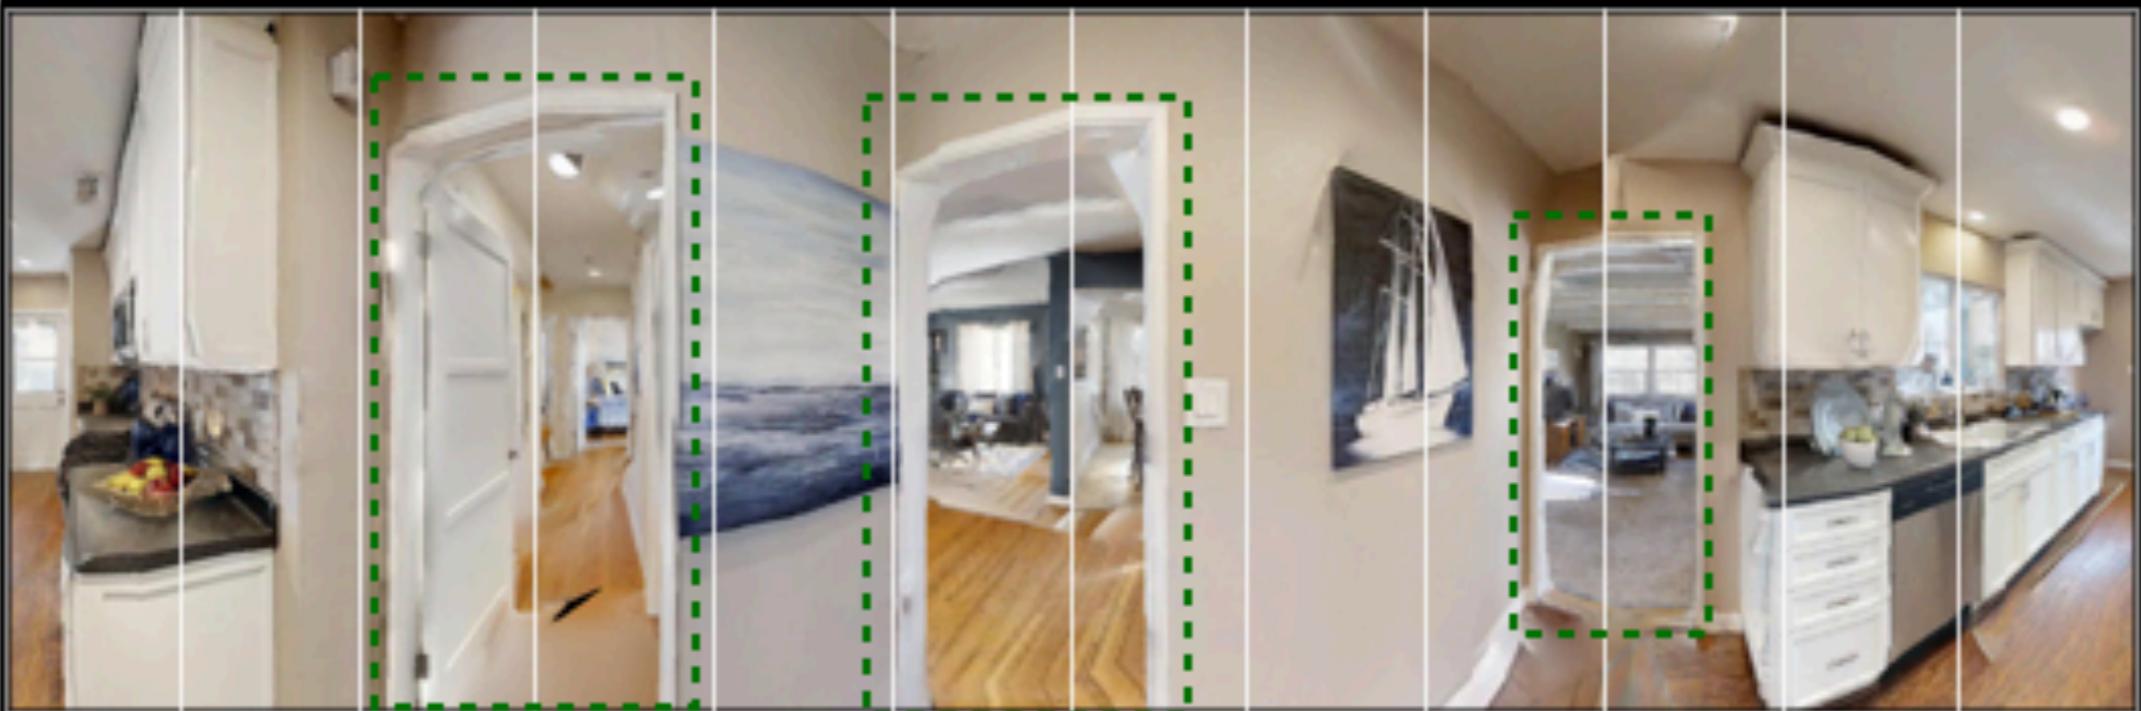
Low-Level Policy



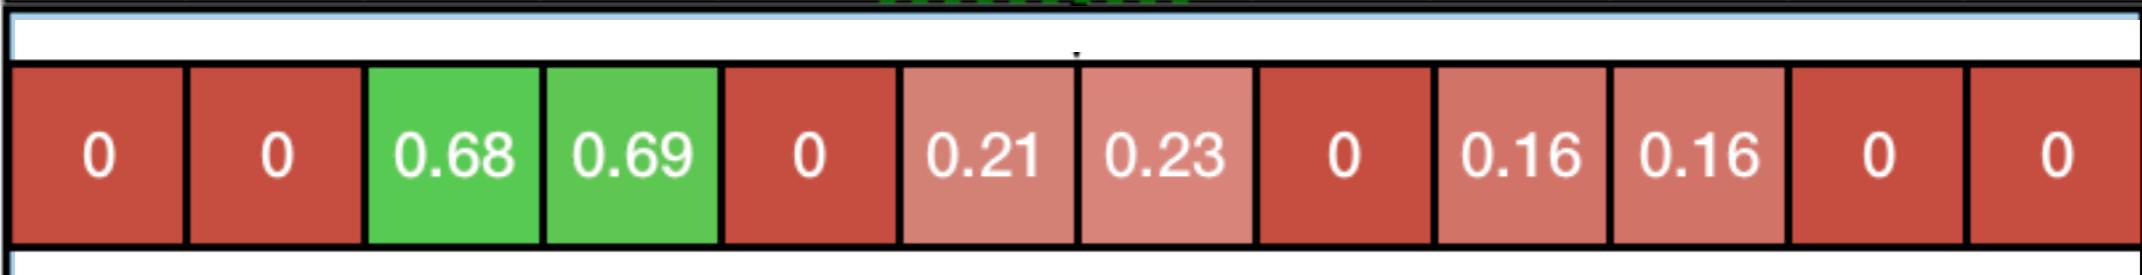
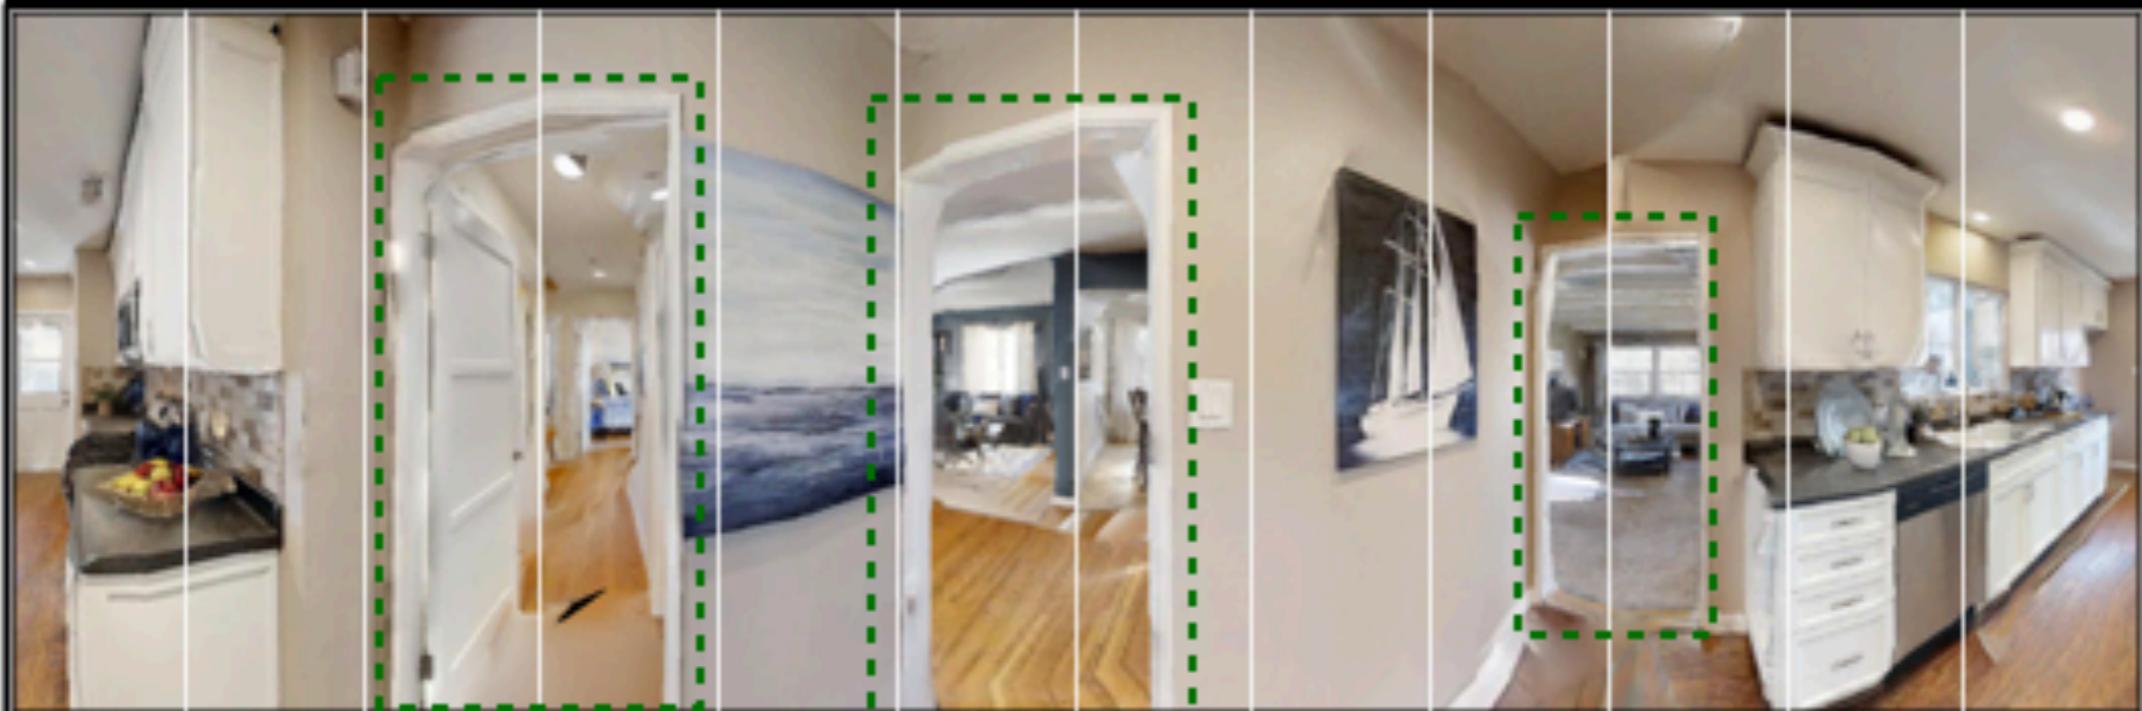
Execution (Image Goal)



Source Image (I_S)



Source Image (I_S)



Results

Map based methods are better than vanilla learning methods even in presence of noise.

	RGB	RGBD	RGBD (No Noise)	RGBD (No Stop)
End-to-end Learning	0.10	0.14	0.15	0.18
LSTM + Imitation	0.10	0.13	0.14	0.17
LSTM + RL	0.10	0.13	0.14	0.17
Occupancy Maps + FBE + RL		0.26	0.31	0.24
ANS	0.23	0.29	0.35	0.39
NTS (Our)	0.38	0.43	0.45	0.60

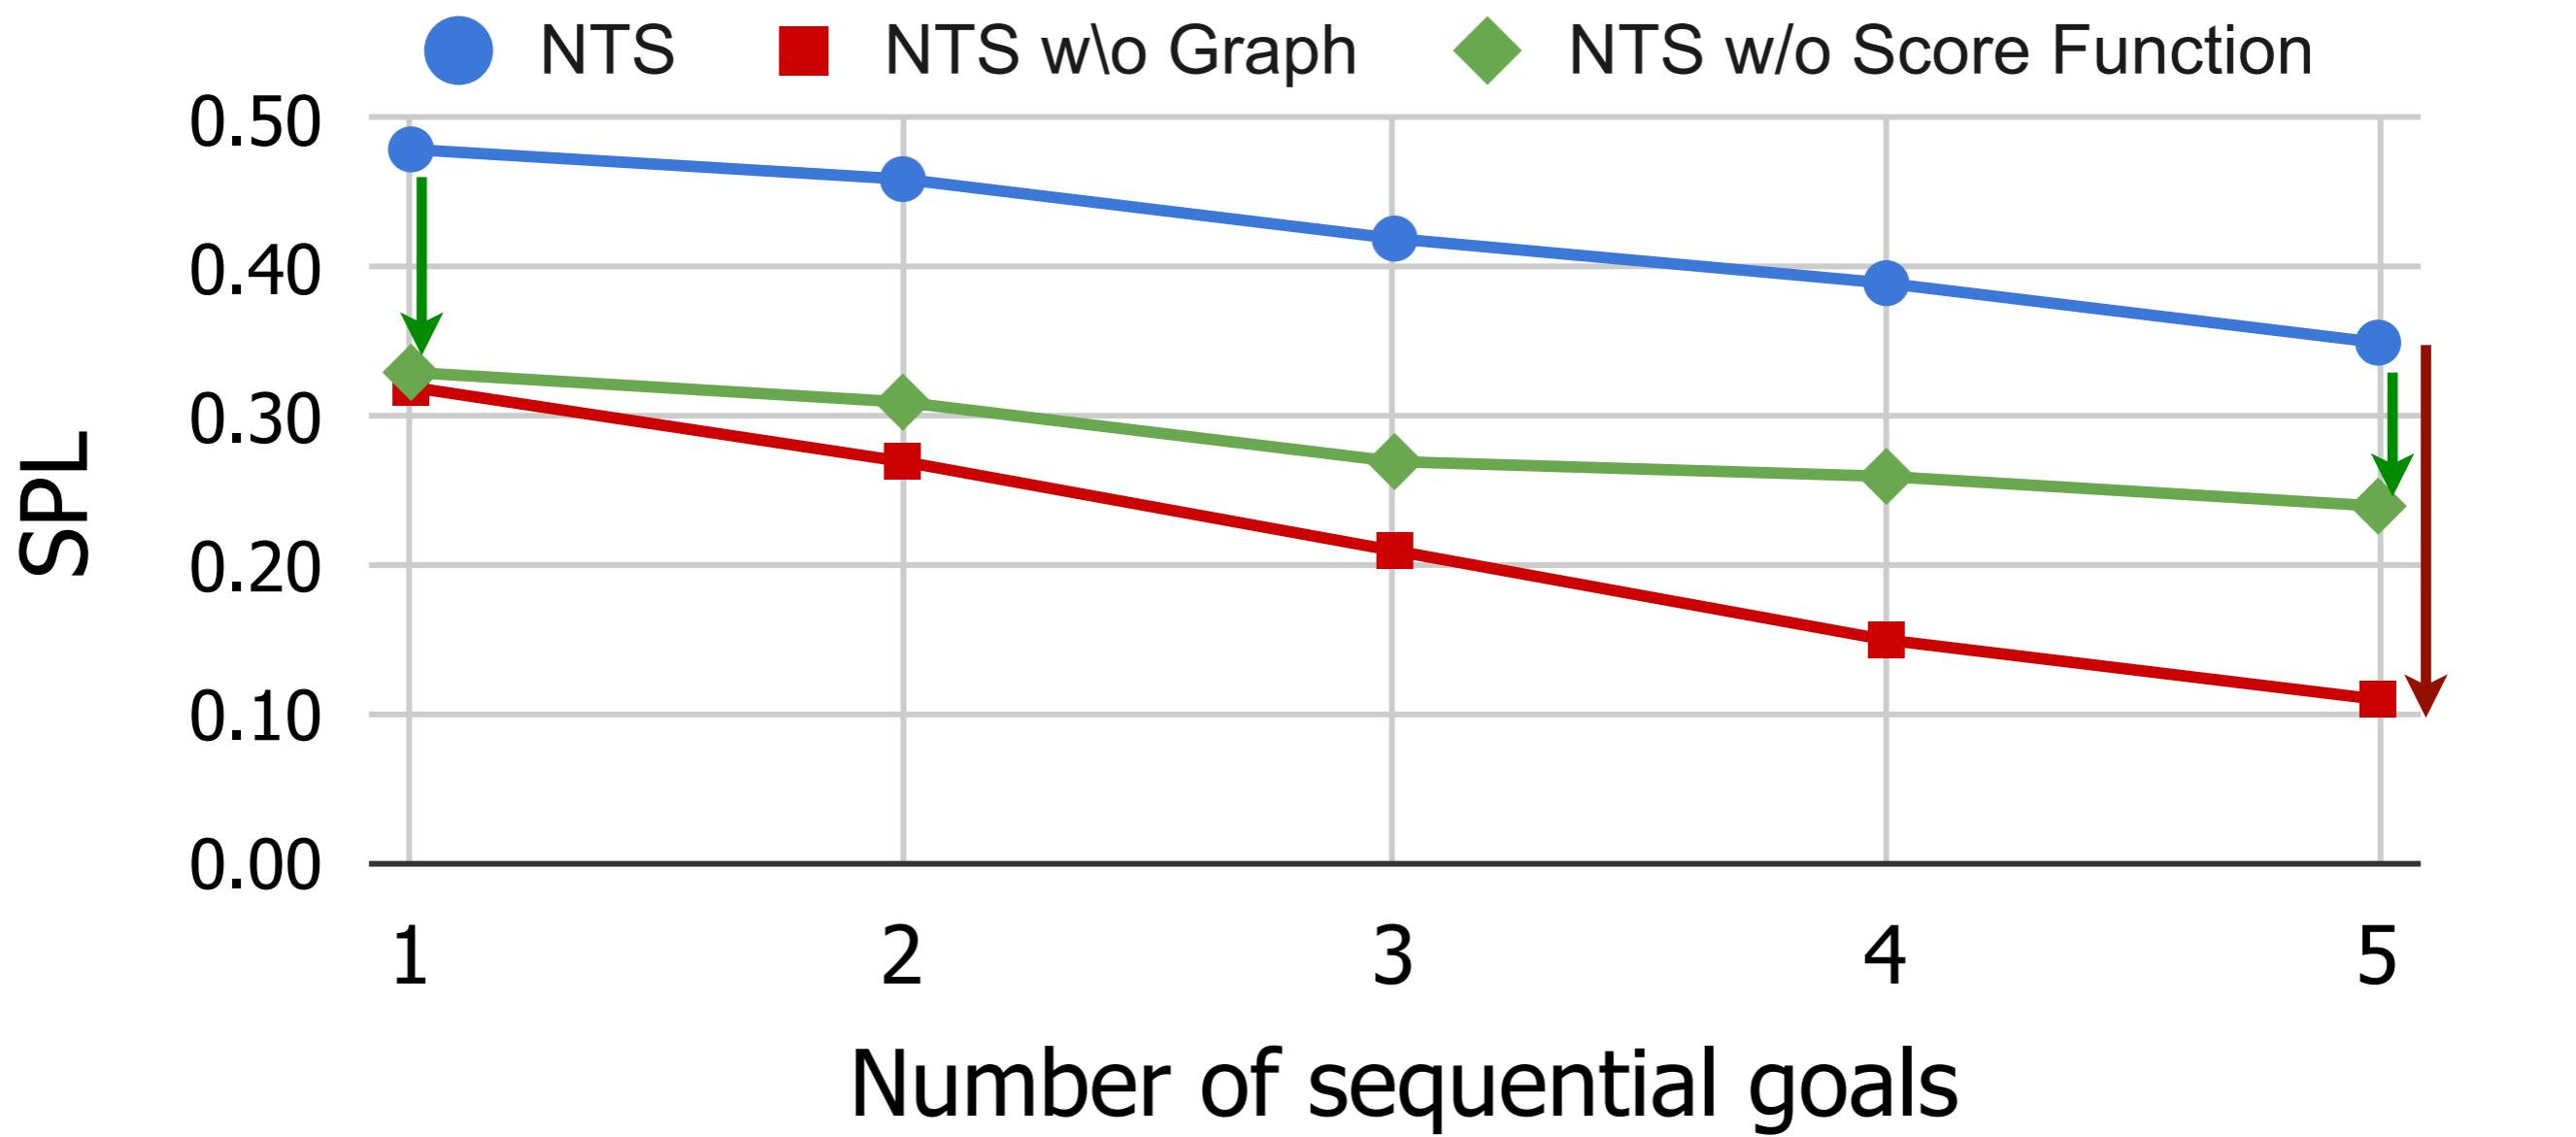
Robust to Actuation Noise

NTS is better than occupancy map models, captures and uses semantic priors.

Results

Semantic score function improves efficiency when no prior experience with environment is available.

As experience in environment increases, utility of semantic function decreases



*But, at the same time,
importance of the topological
representation increases*

Thanks!